

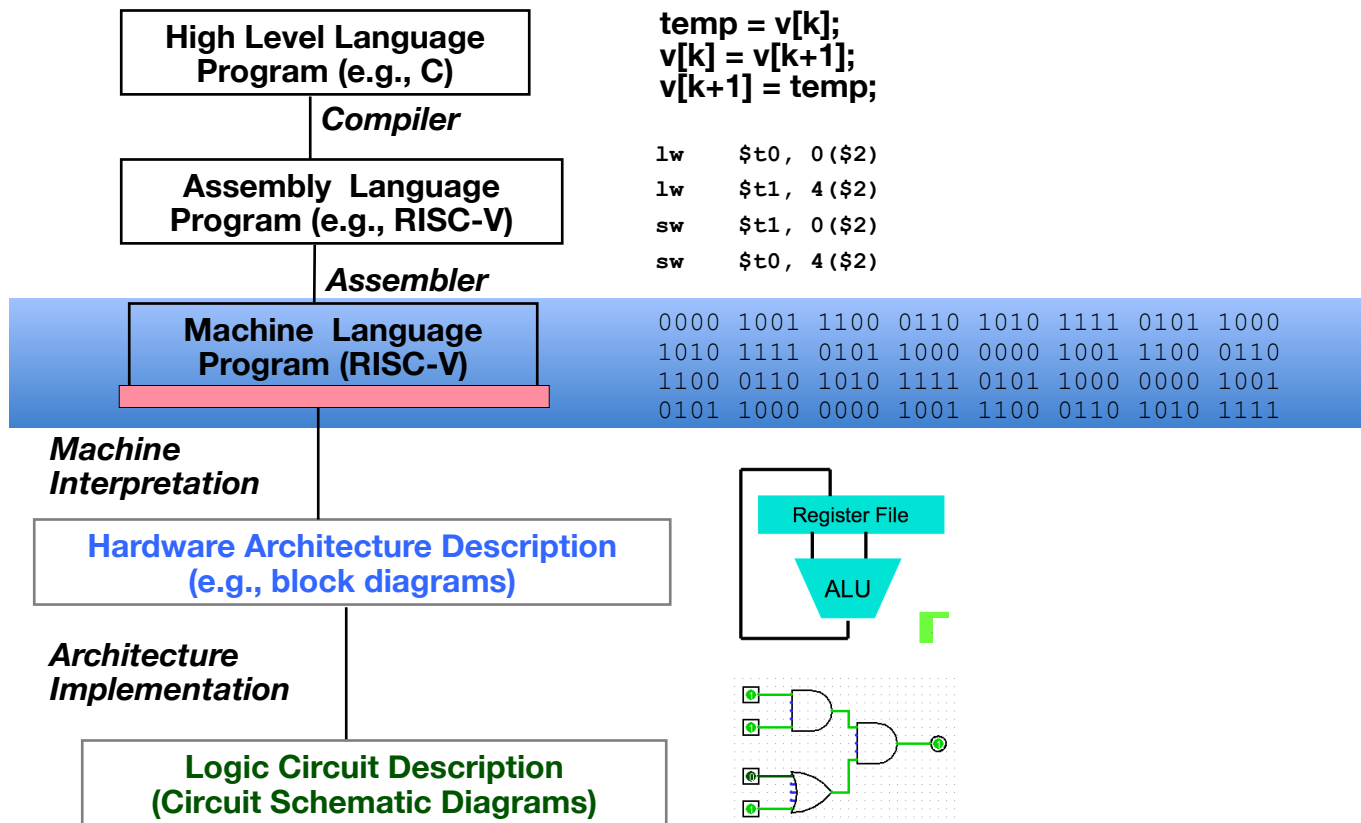
CS 61C:

Great Ideas in Computer Architecture

RISC-V Instruction Formats

Great Idea #1: Abstraction

Levels of Representation/Interpretation

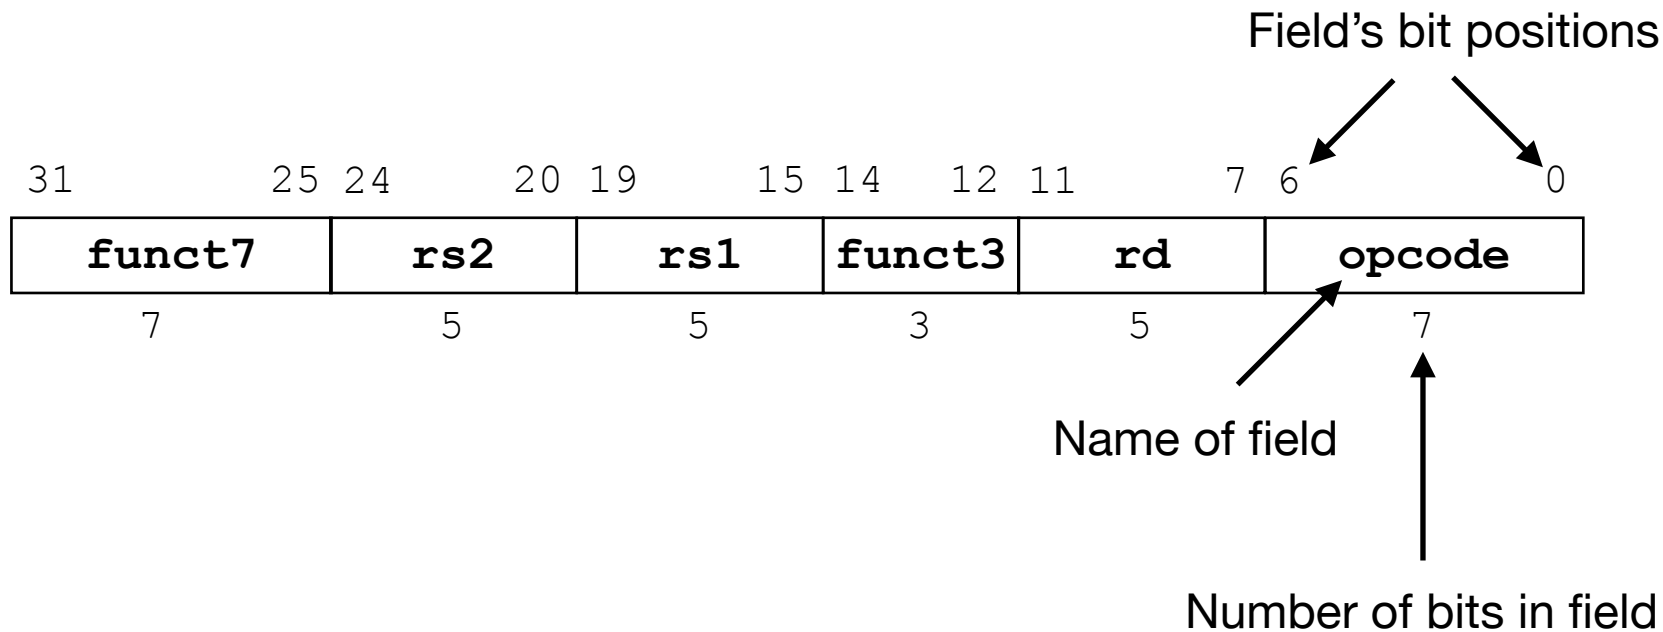


Instruction Formats


- Each instruction is 32 bits wide
- The instruction is broken down into different fields
- There are several ways that instructions are broken up

R-Format

R-Format Layout




R-Format Instructions opcode/funct fields



- **opcode:** partially specifies which instruction it is
 - opcode = 0b0110011 for all 32-bit R-Format arithmetic/logical instructions
- **funct7+funct3:** combined with opcode, these two fields describe what operation to perform

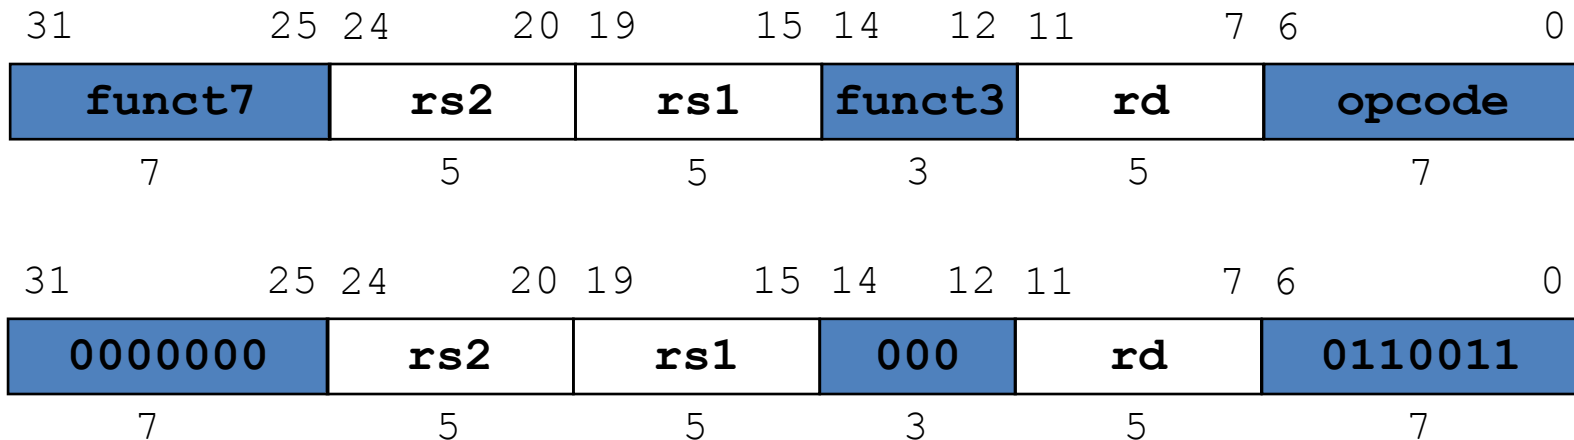
R-Format Instructions Registers



- Each register field (rs1, rs2, rd) holds a 5-bit unsigned integer [0-31] corresponding to a register number (x0-x31)
- rs1 = source register #1
- rs2 = source register #2
- rd = destination register

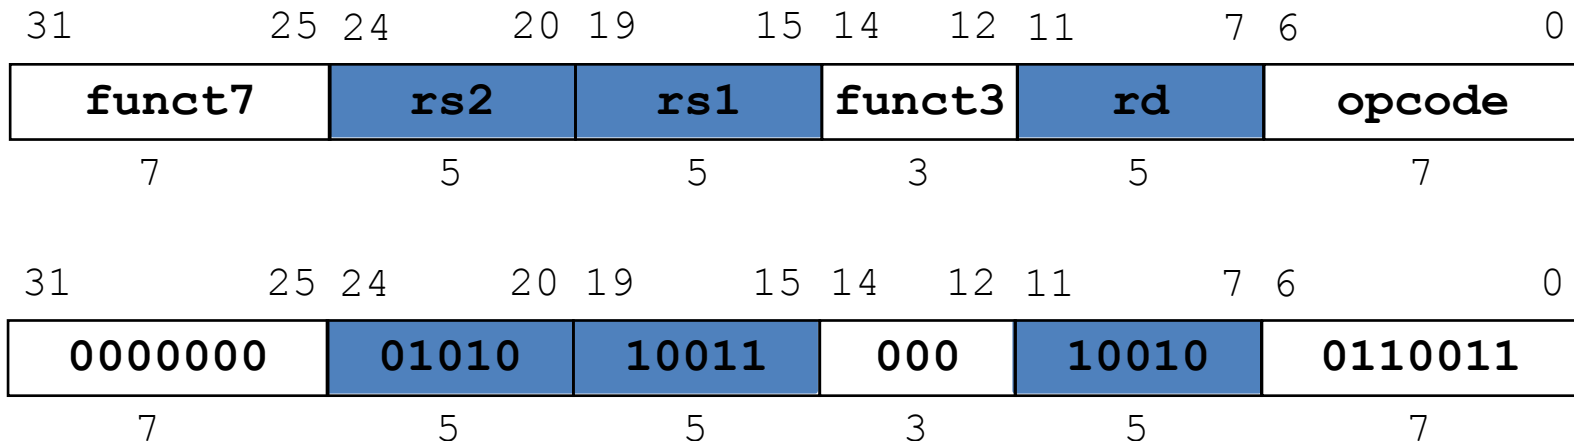
R-Format Example

add x18, x19, x10



R-Format Example

add x18, x19, x10



All RV32 R-format instructions

funct7			funct3		opcode	
0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND


funct7 + funct3 selects particular operation

I-Format

I-Format Layout


`add x18, x19, x10`

`addi x18, x19, 2`



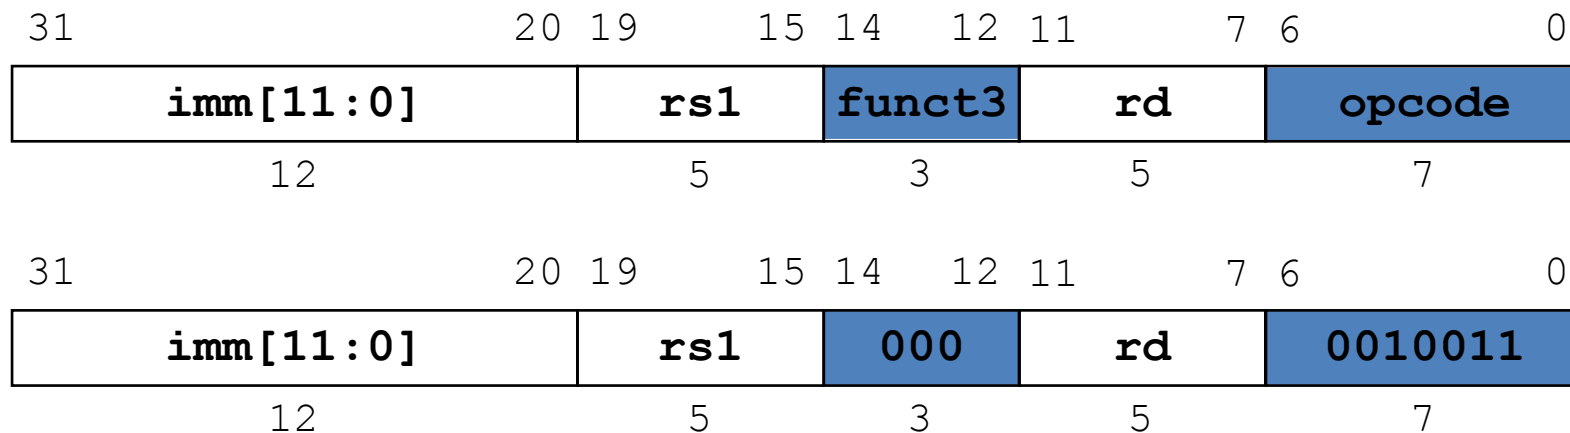
I-Format Instructions

- **imm[11:0]** can hold values in range $[-2048_{10}, +2047_{10}]$
- Immediate is always sign-extended to 32-bits before use in an arithmetic operation
- We'll later see how to handle immediates > 12 bits



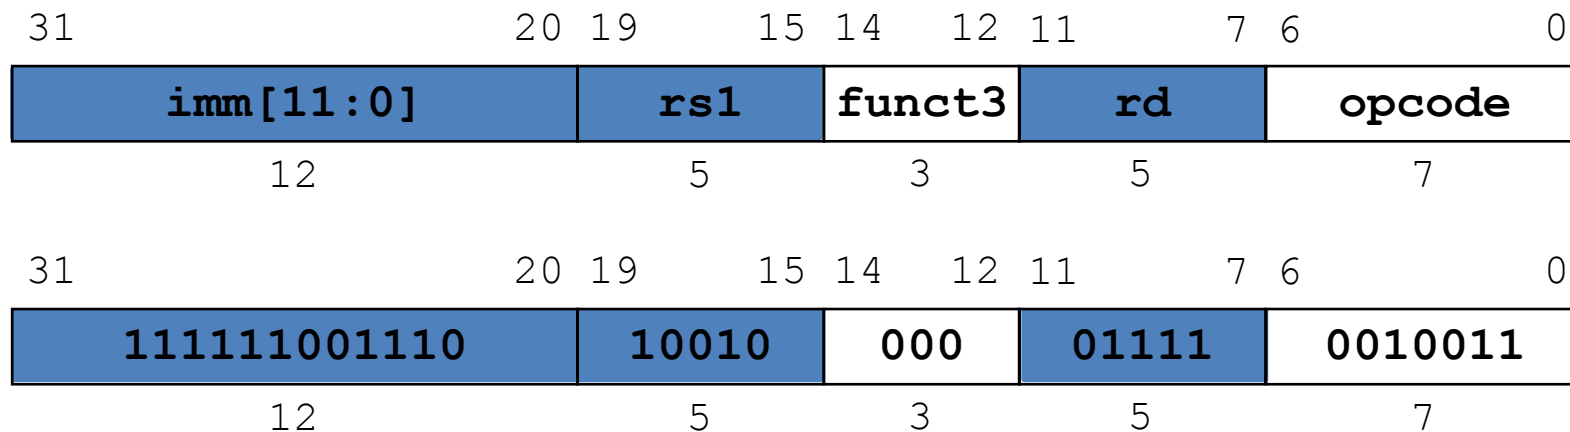
I-Format Example

addi x15, x18, -50



I-Format Example

addi x15, x18, -50



All RV32 I-format Arithmetic/Logical Instructions

Immediate		funct3			opcode	
imm[11:0]		rs1	000	rd	0010011	ADDI
imm[11:0]		rs1	010	rd	0010011	SLTI
imm[11:0]		rs1	011	rd	0010011	SLTIU
imm[11:0]		rs1	100	rd	0010011	XORI
imm[11:0]		rs1	110	rd	0010011	ORI
imm[11:0]		rs1	111	rd	0010011	ANDI
0000000	shamt	rs1	001	rd	0010011	SLLI
0000000	shamt	rs1	101	rd	0010011	SRLI
0100000	shamt	rs1	101	rd	0010011	SRAI

↑
Distinguishes
logical and
arithmetic right

↑
The max number of bits we need
to shift by is 31, so we only need
5 bits to encode the shift amount

↑
These funct3 fields are the
same as the corresponding
r-type operation

↑
Opcode is the same for
all arithmetic and
logical immediate
instructions


Loads

Review: Load Instructions


- Load word
 - Loads one word from memory into the destination register
- Load Byte
 - Loads one byte from memory into the lowest byte of the destination register
 - Value is sign extended
- Load Byte Unsigned
 - Loads one byte from memory into the lowest byte of the destination register
 - Value is not sign extended

Review: Load Instructions

Load Byte



Load Byte Unsigned



Load Halfword

lh x10, 0(x15)


Destination Register Offset Base register



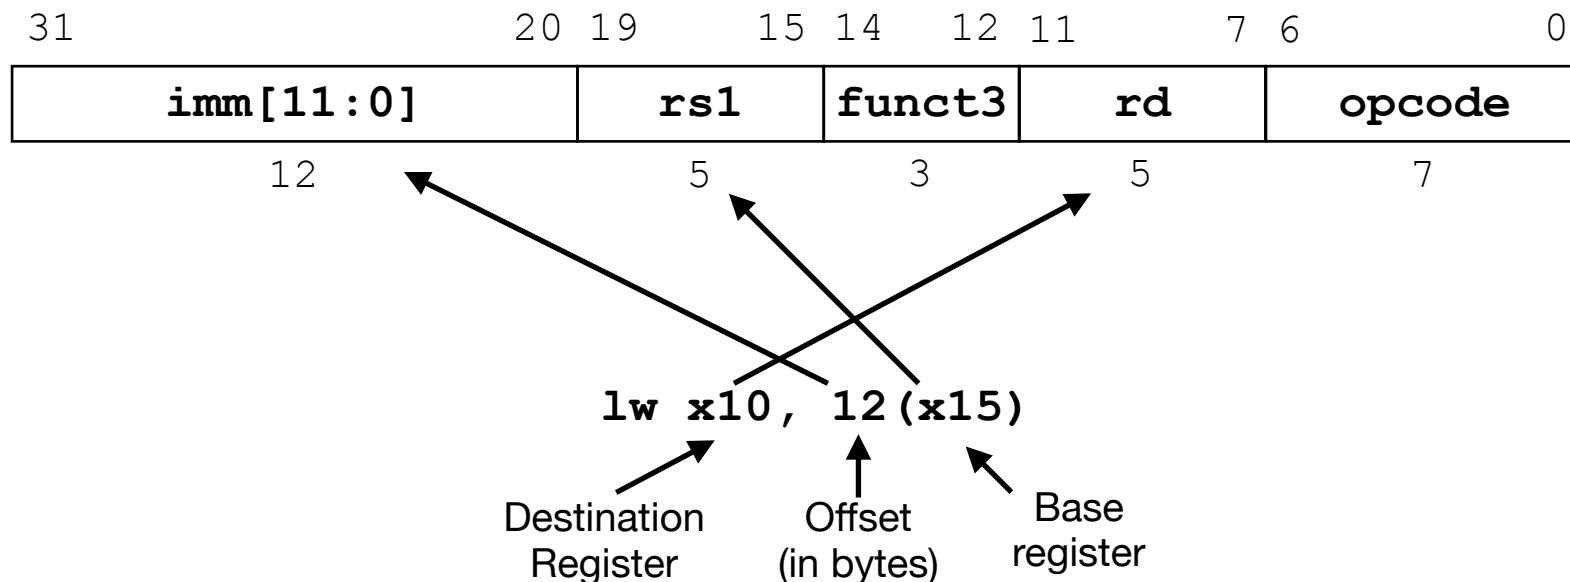
Load Halfword Unsigned

`lhu x10, 0(x15)`

Destination Register Offset Base register



Load Instructions are also I-Type



All RV32 Load Instructions

funct7		funct3		opcode	
imm[11:0]	rs1	000	rd	0000011	LB
imm[11:0]	rs1	001	rd	0000011	LH
imm[11:0]	rs1	010	rd	0000011	LW
imm[11:0]	rs1	100	rd	0000011	LBU
imm[11:0]	rs1	101	rd	0000011	LHU

There is no lwu because a register is 32 bits, so its never sign extended

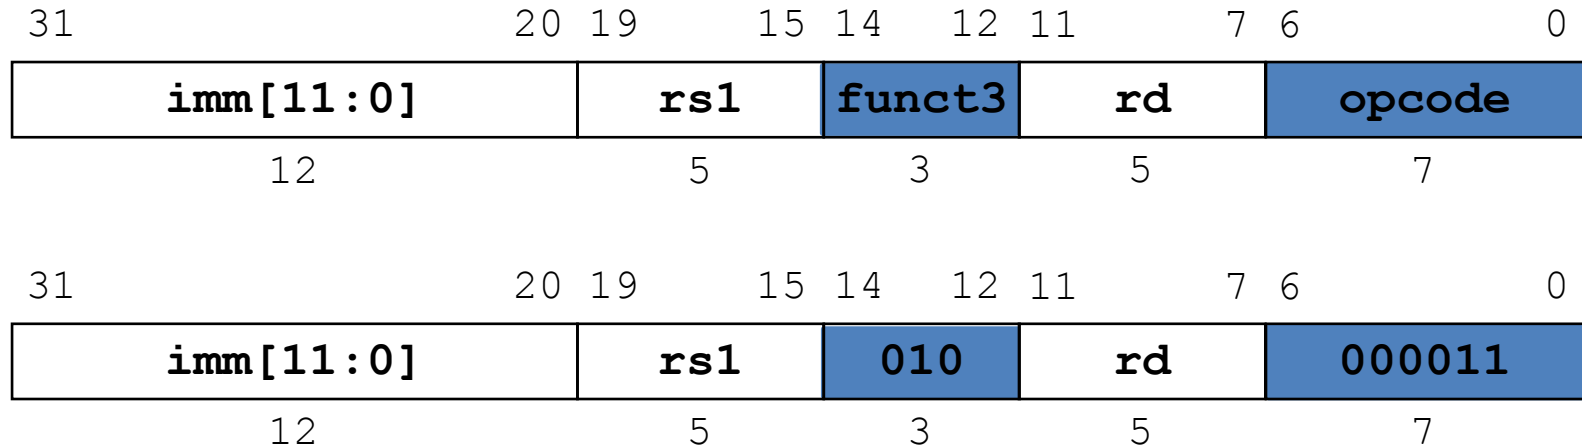
MSB tells us if it's unsigned
0 = signed
1 = unsigned

Bottom 2 bits tell us how much to load
00: 1 byte
01: 2 bytes
10: 4 bytes

Same opcode

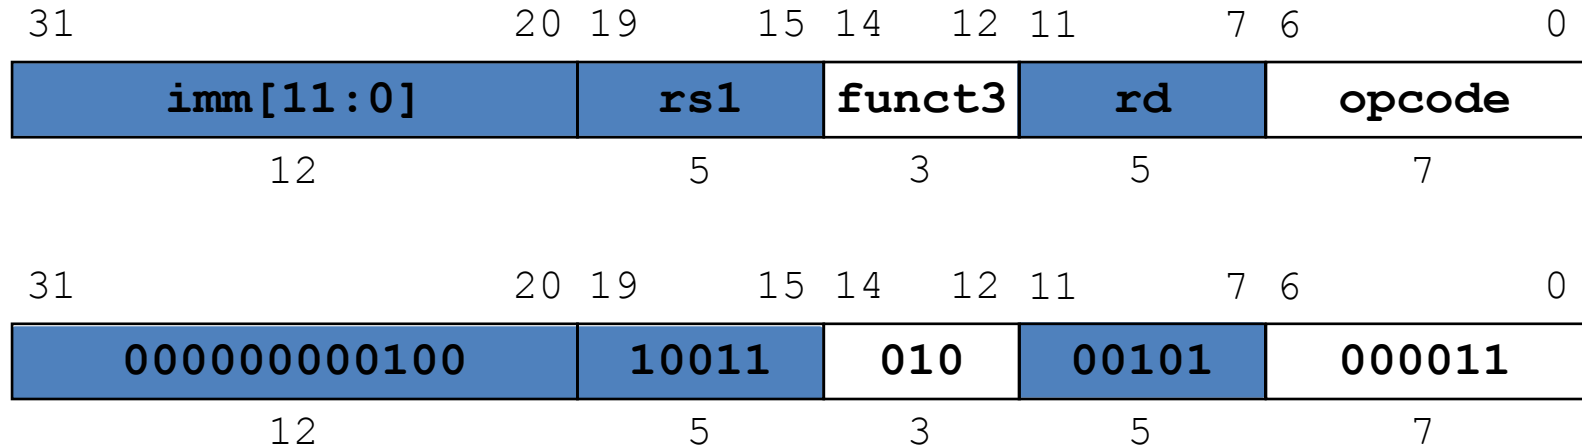
Load Example

lw x18, 4(x19)




Load Example

lw x5, 4(x19)



S-Format

Store Instructions

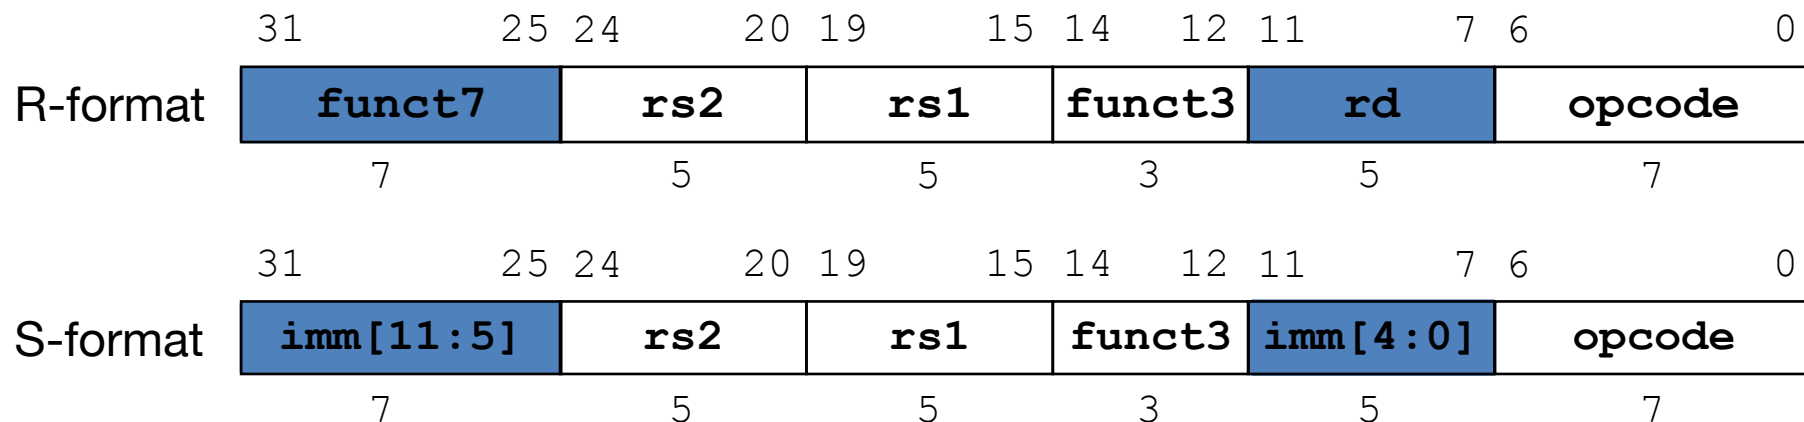


- Stores have 2 source registers, no destination
- I-type instructions have 1 source and 1 destination
- We want to prioritize keeping registers in the same place

Store Instructions

sw x7, 12 (x5)


Source Register Offset (in bytes) Base register



Review: Store Instructions

- Store Word
 - Stores the entire contents of the source register into memory
- Store Byte
 - Stores the least significant byte of the source register into memory
- No Sign extension

Store Halfword



- Stores the lower 16 bits of register x14 into memory at address $[x5] + 36$
- No sign extension

All RV32 Store Instructions

imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	SW

Bottom 2 bits tell
us how much to
store

00: 1 byte

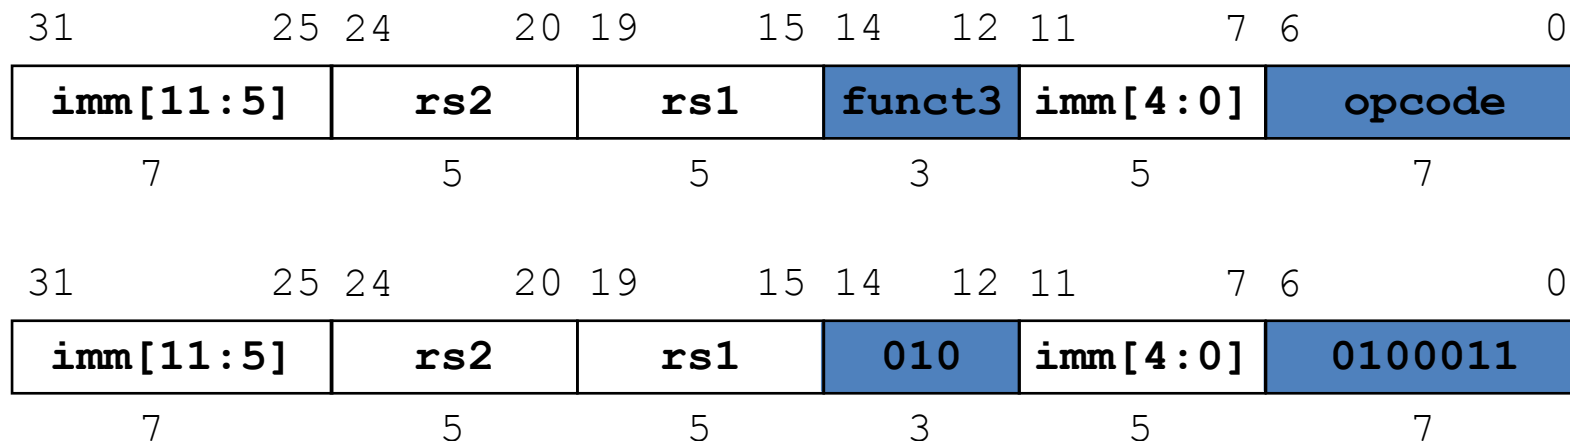
01: 2 bytes

10: 4 bytes

Same
opcode

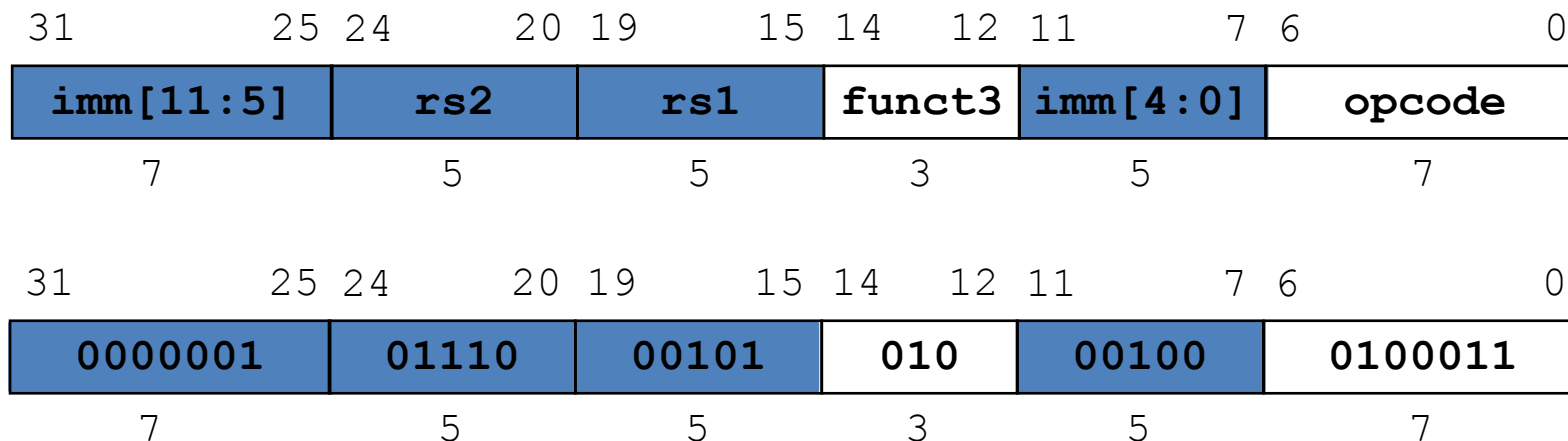
Store Format Example

sw x14, 36(x5)



Store Format Example

sw x14, 36(x5)



36 = 0b

01	00100
----	-------

Branches

Recall: Conditional Branches

- Used for ifs, loops, etc...
- Format: {comparison} {reg1} {reg2} {label}
 - beq
 - bne
 - blt, bltu
 - bge, bgeu

Recall: Incrementing PC

- RV32 instructions are 32 bits = 4 bytes
- When we want to move to the next instruction, the processor increments PC by 4 bytes

```
if (a == b)      x10 = a
    e = c + d;    x11 = b
else             x12 = c
    e = c - d;    x13 = d
                 x14 = e
```

```
PC → bne x10,x11,else
      add x14,x12,x13
      j  done
else:  sub x14,x12,x13
done:
```

Recall: Conditional Branches

- Format: {comparison} {reg1} {reg2} {label}
- If we don't branch
 - $PC = PC + 4$
- If we do branch
 - $PC = PC + \text{immediate}$

What range of instructions can we branch to?

- 2's complement range: $[-2^{n-1}, 2^{n-1}-1]$
- With 12 bits: $\pm 2^{11}$ bytes away from the PC
- Instructions are 4-bytes, so we can jump $\pm 2^9$ instructions away from the current instruction

Instruction Addressing

0x10000000	Instruction 0
0x10000004	Instruction 1
0x10000008	Instruction 2
0x1000000C	Instruction 3
0x10000010	Instruction 4
0x10000014	Instruction 5
0x10000018	Instruction 6
0x1000001C	Instruction 7

The last nibble is always 0x (0, 4, 8, or C)

0x0 = 0b0000
0x4 = 0b0100
0x8 = 0b1000
0xC = 0b1100



The last two bits
are always 0b00

Branch Instruction Addressing

- The last two bits are always 0, so we can increase our range by not storing those bits

encoded immediate = 0b0000 0000 0011

actual immediate offset = 0b00 0000 0000 1100

- (PC + immediate) will go 3 instructions (or 12 bytes) away
- Now, we can jump $\pm 2^{11}$ instructions (or $\pm 2^{13}$ bytes) away from the current PC

RISC-V Feature, $n \times 16$ -bit instructions

- Extensions to RISC-V base ISA support 16-bit compressed instructions and also variable-length instructions that are multiples of 16 bits in length

Instruction Addressing

0x10000000	Instruction 0
0x10000002	Instruction 1
0x10000004	Instruction 2
0x10000006	Instruction 3
0x10000008	Instruction 4
0x1000000A	Instruction 5
0x1000000C	Instruction 6
0x1000000E	Instruction 7

The last nibble always ends in an even number

0x0	=	0b0000
0x2	=	0b0010
0x4	=	0b0100
0x6	=	0b0110
0x8	=	0b1000
0xA	=	0b1010
0xC	=	0b1100
0xE	=	0b1110




The last bit is
always 0b0

RISC-V Feature, nx16-bit instructions

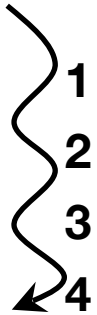
- With the 16-bit encoding, we can only save room by not storing the last bit
- Don't want to have two different branch encodings, so we will choose to only discard 1 bit instead of discarding 2 bits
- Range of bytes we can jump to
 - $\pm 2^{12}$
- Range of 32-bit instructions we can jump to
 - $\pm 2^{10}$

Branch Format



Branch Example

```
Loop: beq  x19,x10,End
      add  x18,x18,x10
      addi x19,x19,-1
      j    Loop
End:   # target instruction
```

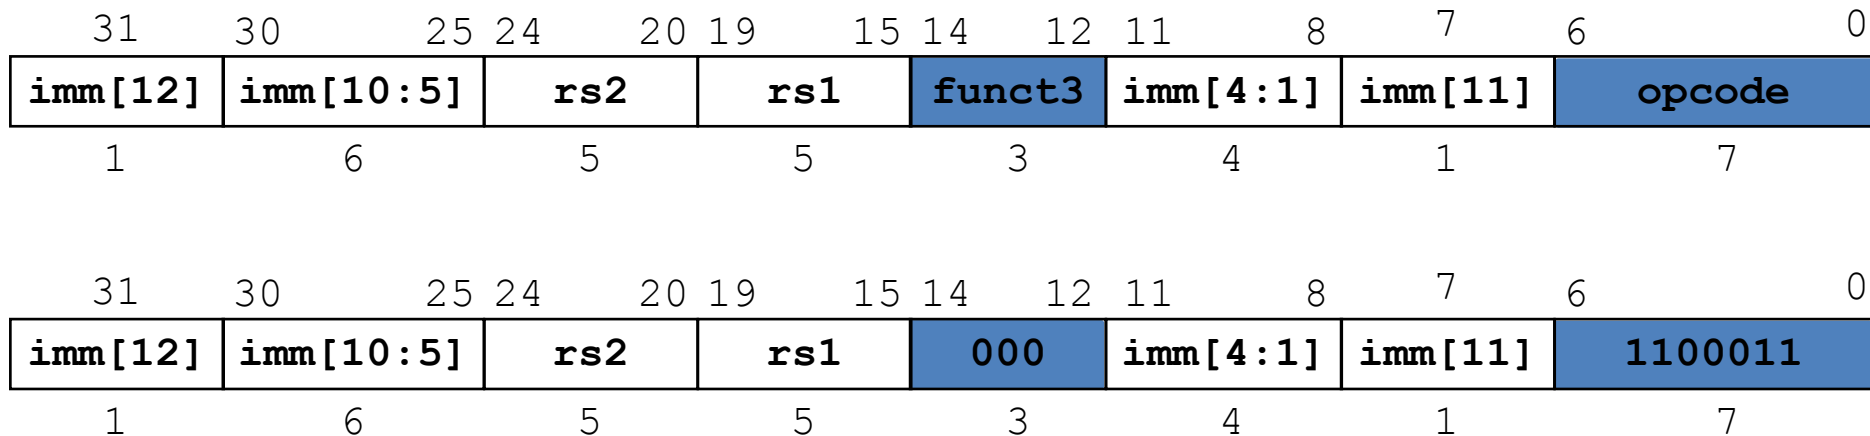


1 Count
2 instructions
3 from branch
4

- Branch offset = 4 instructions
 - (4 instructions x 4 bytes) = 16 bytes
 - 0b 0000 0001 0000

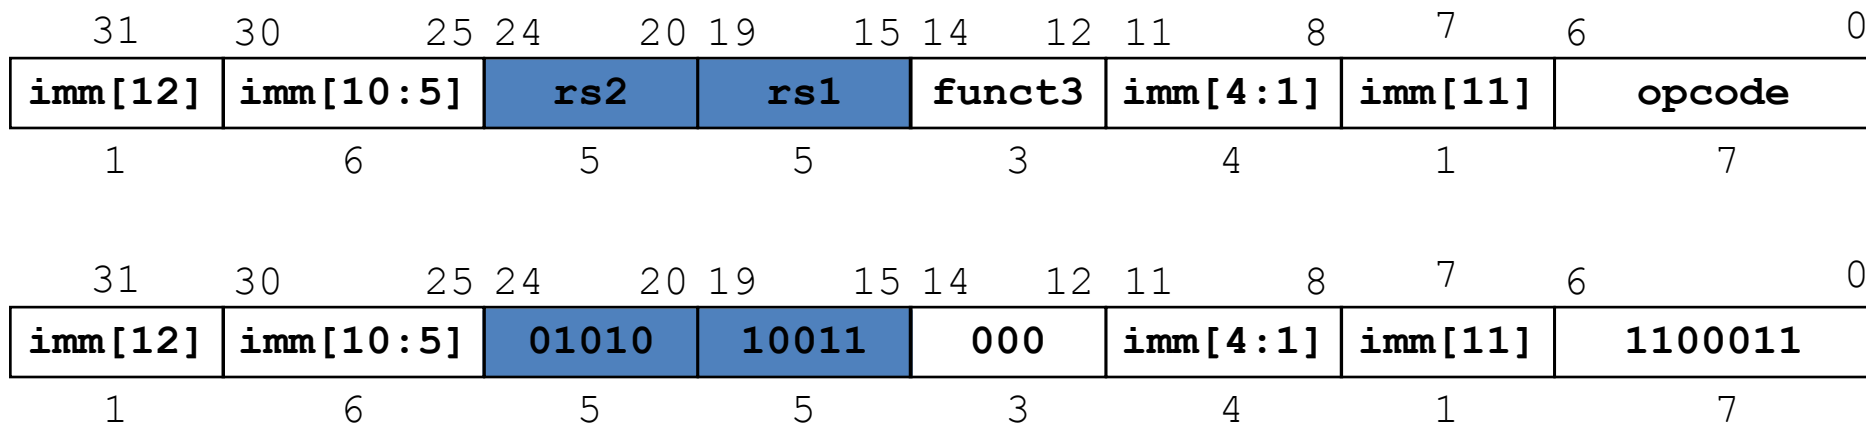
Branch Layout Example

beq x19, x10, End



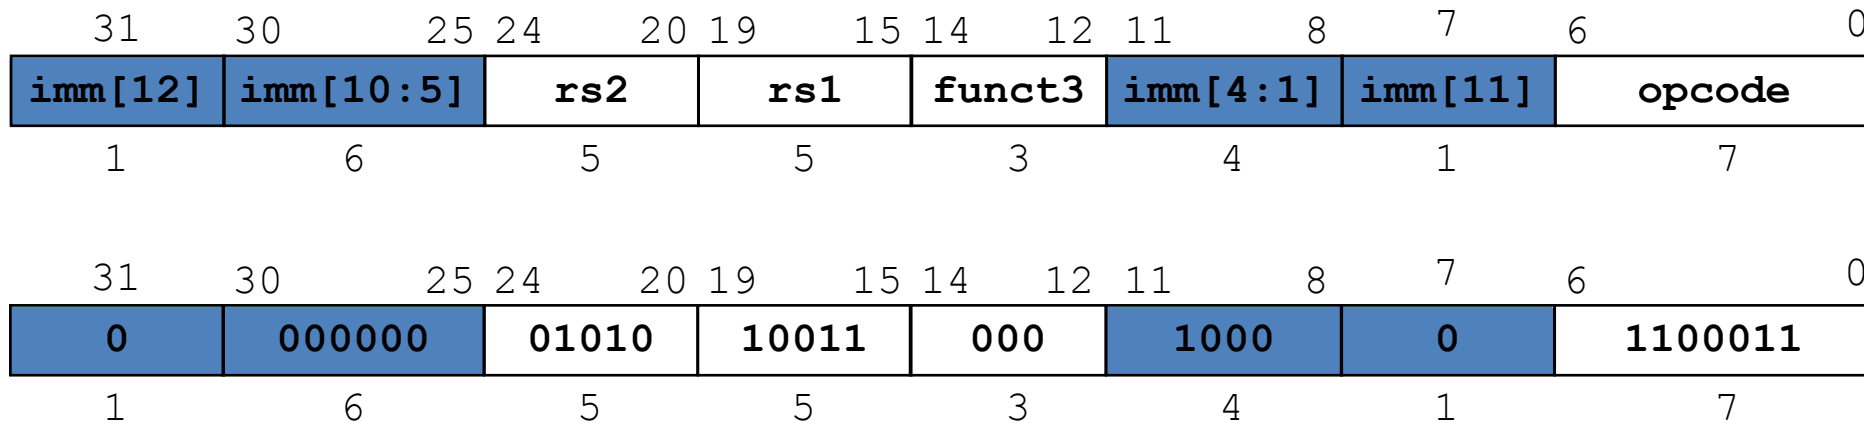
Branch Layout Example

beq x19, x10, End




Branch Layout Example

beq x19, x10, End



Branch offset = 16 bytes = 0b 0000 0001 0000



All RISC-V Branch Instructions

imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011	BEQ
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011	BNE
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	1100011	BLT
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	1100011	BGE
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	1100011	BLTU
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	1100011	BGEU

↑
funct3 field
specifies with
type

↑
Same
opcode

Conditional Branches

- Used for if-else statements and loops
 - These are usually pretty small (< 50 lines of code)
- Branch instructions have a limited range
 - $\pm 2^{10}$ 32-bit instructions
- What do we do if the location where we want to jump is farther away?

```
beq x10,x0,far  
# next instr
```

```
bne x10,x0,next  
j    far  
next: # next instr
```

J-Format

Recall: JAL

`jal rd, Label`



rd = register where the
return address will be
stored

← The label that we
want to jump to

rd = return address

$PC = PC + \text{offset}$

- The label that we want to jump to gets translated by the assembler to a 20-bit offset


J-Format

- 20 bits for immediate field
- Just like with branch instructions, we can leave off the last bit
- So we can jump by
 - $\pm 2^{20}$ bytes
 - $\pm 2^{18}$ 32-bit instructions

J-Format

`jal rd, Label` ← The label that we want to jump to

↑
rd = register where the return address will be stored

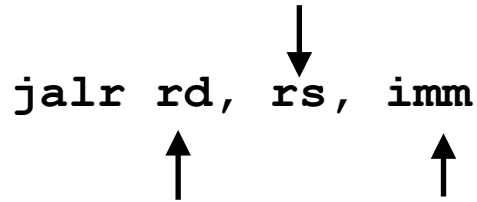


JALR

Recall: JALR

Register containing
the base address
(source register)

`jalr rd, rs, imm`




rd = register where the
return address will be
stored

Immediate value
to be added to
the base register

rd = return address
 $PC = [rs] + imm$

JALR is an I-type Instruction

- Usually used to return from a function (ret)
- **imm**[11:0] can hold values in range $[-2048_{10}, +2047_{10}]$
- Immediate is always sign-extended to 32-bits before use
- Unlike JAL, we must include the last 0 because we are using the I-format which specifies that we include the 0th bit



PC-Relative Address vs Absolute Address

- PC Relative Address
 - Jump to a location based on the current location of the PC
 - $PC + \text{offset}$
- Absolute Address
 - Jump to a location using that location's full address

Questions on PC-addressing

- Does the value in branch immediate field change if we change the location of the code in memory?
 - If moving all of code, then no (because PC-relative offsets)

U-Format

Load Upper Immediate (LUI)

- destination register = immediate \ll 12

`lui x5, 0xABCDE`

Destination register Immediate value

`x5 = 0xABCDE000`

LUI to create long immediates

- LUI writes the upper 20 bits of the destination with the immediate value, and clears the lower 12 bits.
- Together with an ADDI to set low 12 bits, can create any 32-bit value in a register using two instructions (LUI/ADDI).

```
LUI x10, 0x87654          # x10 = 0x87654000
ADDI x10, x10, 0x321       # x10 = 0x87654321
```

Load Immediate

- Pseudo-instruction that performs lui and addi

```
LUI x10, 0x87654          # x10 = 0x87654000  
ADDI x10, x10, 0x321      # x10 = 0x87654321
```



```
LI x10, 0x87654321        # x10 = 0x87654321
```

One Corner Case

- How to set 0xABCDEEEE?

```
lui x10, ABCDE # x10 = 0xABCDE000
addi x10, 0xEE # x10 = 0xABCDDEEE
```

$$\begin{array}{r} 0xABCDE000 \\ + 0xFFFFFFFF \\ \hline 0xABCDDEEE \end{array}$$

$$\begin{array}{r} 0xABCDE \\ + 0xFFFFF \\ \hline 0xABCDD \end{array} \quad \begin{array}{r} 0x000 \\ + 0xEE \\ \hline 0xEE \end{array}$$

Solution

- If the value being added is negative, add 1 to the upper 20 bits before adding the 12-bit value
- How to set 0xABCDEEEE?

```
lui x10, ABCDF # x10 = 0xABCDF000
addi x10, 0xEE # x10 = 0xABCDEEEE
```

- `li` instruction will automatically handle this corner case

Add Upper Immediate PC (AUIPC)


- $rd = PC + (\text{immediate} \ll 12)$

`auipc x5, 0xABCDE`

Destination register Immediate value

$x5 = PC + 0xABCDE000$

U-Format for Upper Immediate Instructions



lui opcode = 0b0110111
auipc opcode = 0b0010111

LUI and AUIPC with JALR

Call function at any 32-bit absolute address

```
lui    x5, <hi20bits>
```

```
jalr   ra, x5, <lo12bits>
```

Jump PC-relative with 32-bit offset

```
auipc  x5, <hi20bits>
```

```
jalr   ra, x5, <lo12bits>
```

Summary of RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0				
funct7				rs2			rs1		funct3		rd			opcode		R-type		
imm[11:0]						rs1		funct3		rd			opcode		I-type			
imm[11:5]				rs2			rs1		funct3		imm[4:0]			opcode		S-type		
imm[12]		imm[10:5]			rs2			rs1		funct3		imm[4:1]		imm[11]		opcode		B-type
imm[31:12]									rd			opcode			U-type			
imm[20]		imm[10:1]				imm[11]		imm[19:12]			rd			opcode		J-type		

Complete RV32I ISA

imm[31:12]				rd	0110111
imm[31:12]				rd	0010111
imm[20:10:1 11 19:12]				rd	1101111
imm[11:0]				rd	1100111
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	1100011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011

LUI
 AUIPC
 JAL
 JALR
 BEQ
 BNE
 BLT
 BGE
 BLTU
 BGEU
 LB
 LH
 LW
 LBU
 LHU
 SB
 SH
 SW
 ADDI
 SLTI
 SLTIU
 XORI
 ORI
 ANDI

0000000	shamt	rs1	001	rd	0010011
0000000	shamt	rs1	101	rd	0010011
0100000	shamt	rs1	101	rd	0010011
0000000	rs2	rs1	000	rd	0110011
0100000	rs2	rs1	000	rd	0110011
0000000	rs2	rs1	001	rd	0110011
0000000	rs2	rs1	010	rd	0110011
0000000	rs2	rs1	011	rd	0110011
0000000	rs2	rs1	100	rd	0110011
0000000	rs2	rs1	101	rd	0110011
0100000	rs2	rs1	101	rd	0110011
0000000	rs2	rs1	110	rd	0110011
0000000	rs2	rs1	111	rd	0110011

0000	pred	succ	00000	000	00000	0001111
0000	0000	0000	00000	001	00000	0001111
000000000000			00000	000	00000	1110011
0000000000001			00000	000	00000	1110011
csr	rs1	001	rd	1110011		
csr	rs1	010	rd	1110011		
csr	rs1	011	rd	1110011		
csr	zimm	101	rd	1110011		
csr	zimm	110	rd	1110011		
csr	zimm	111	rd	1110011		

SLLI
 SRLI
 SRAI
 ADD
 SUB
 SLL
 SLT
 SLTU
 XOR
 SRL
 SRA
 OR
 AND
 FENCE
 FENCE.I
 ECALL
 EBREAK
 CSRRW
 CSRRS
 CSRRC
 CSRRWI
 CSRRSI
 CSRRCI

Not in CS61C