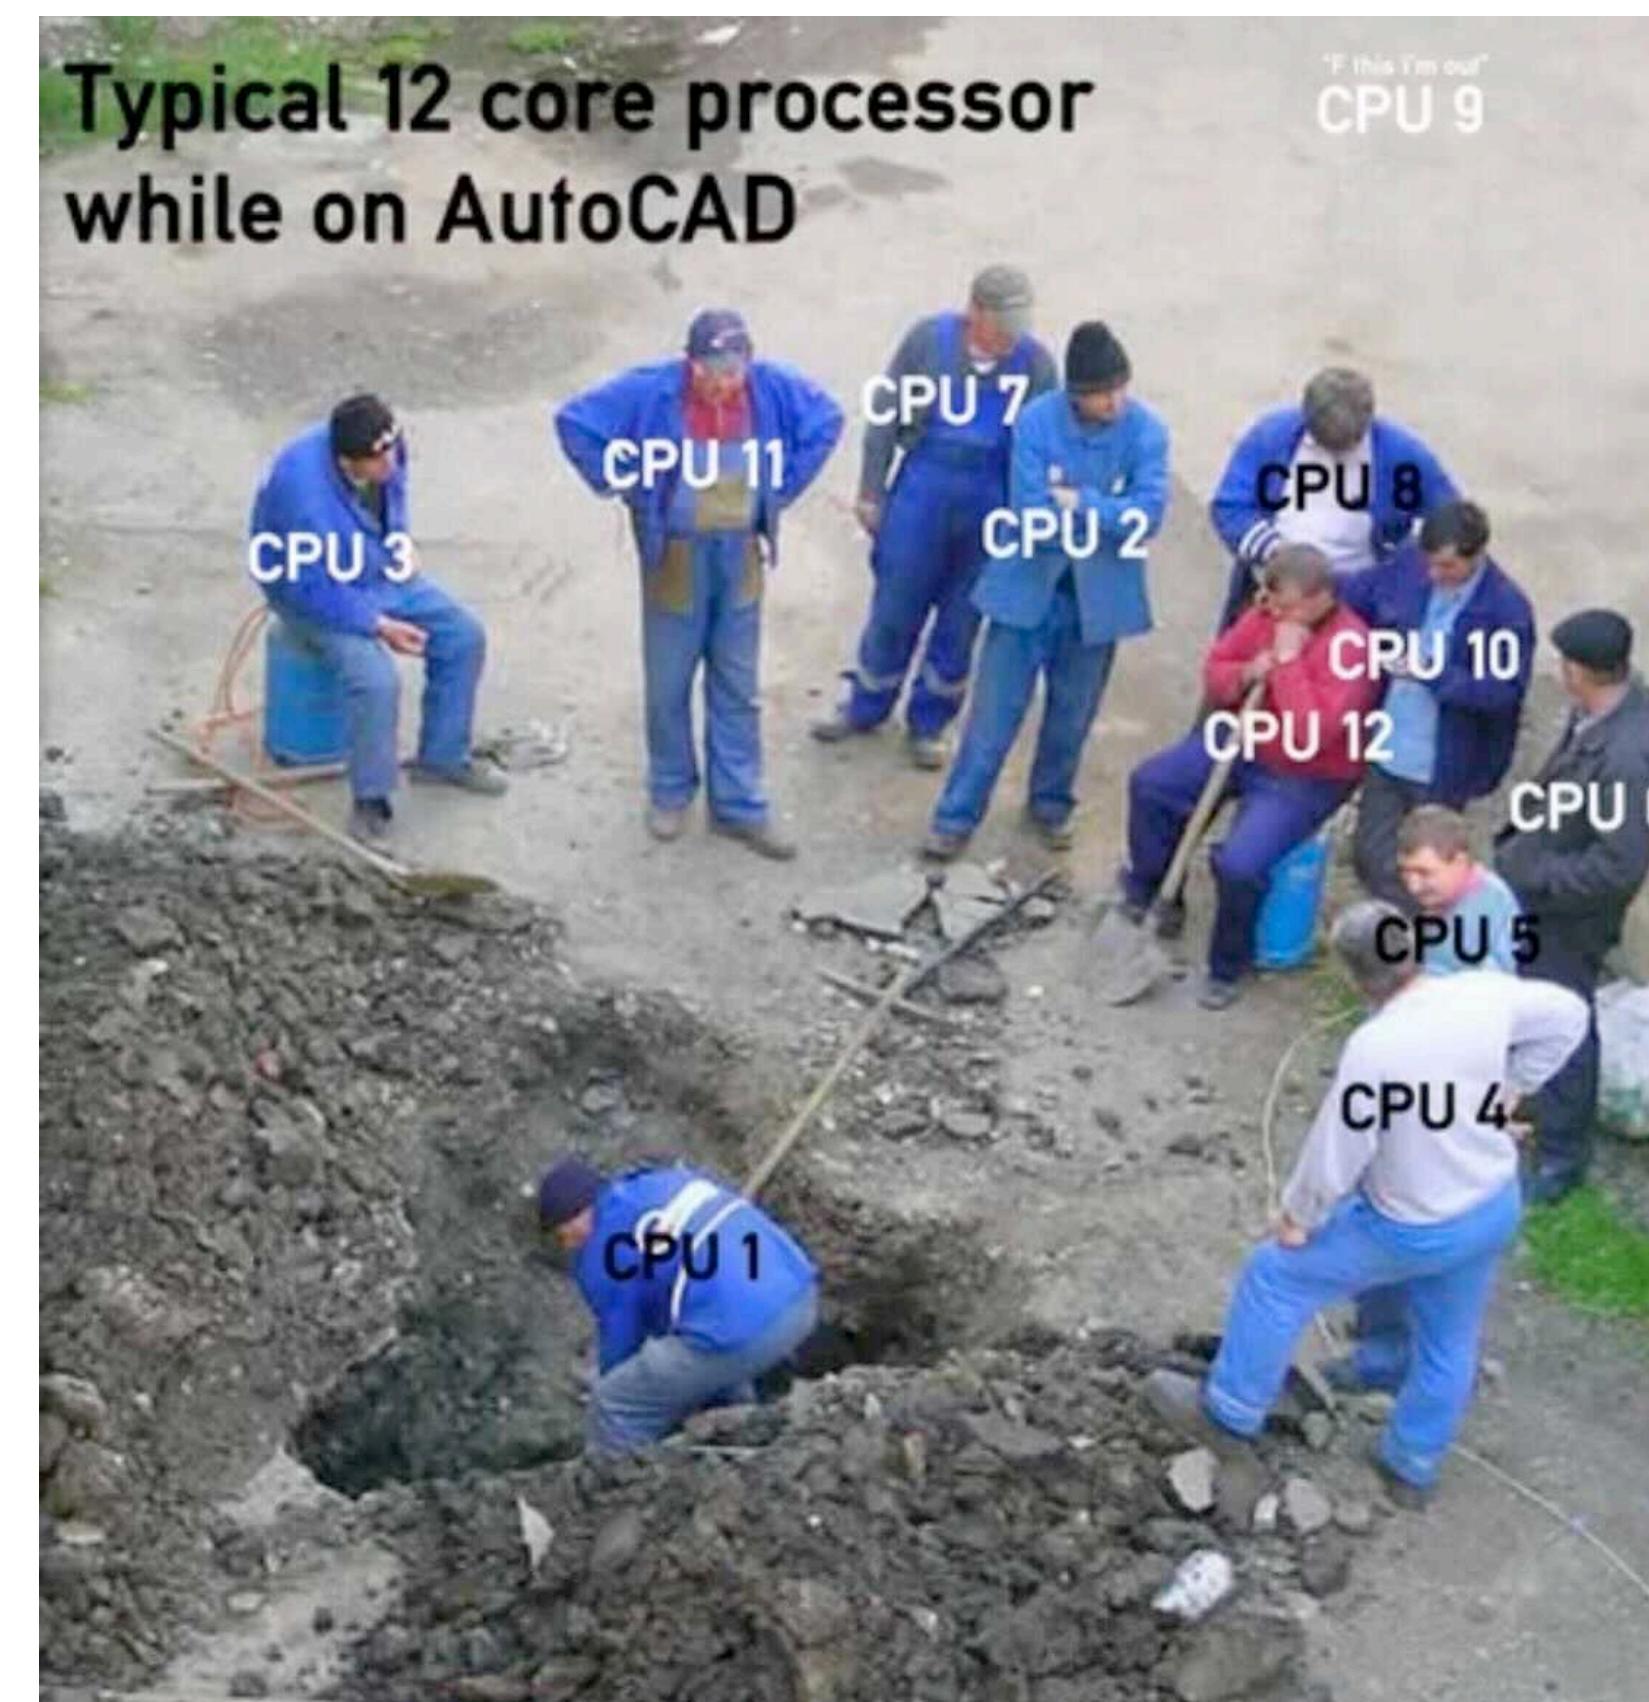


# Parallelism 1



# Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

# 61C Topics so far ...

- What we learned:
  - Binary numbers
  - C
  - Pointers
  - Assembly language
  - Processor micro-architecture
  - Pipelining
  - Caches
  - Floating point
- What does this buy us?
  - Promise: execution speed
  - Let's check!

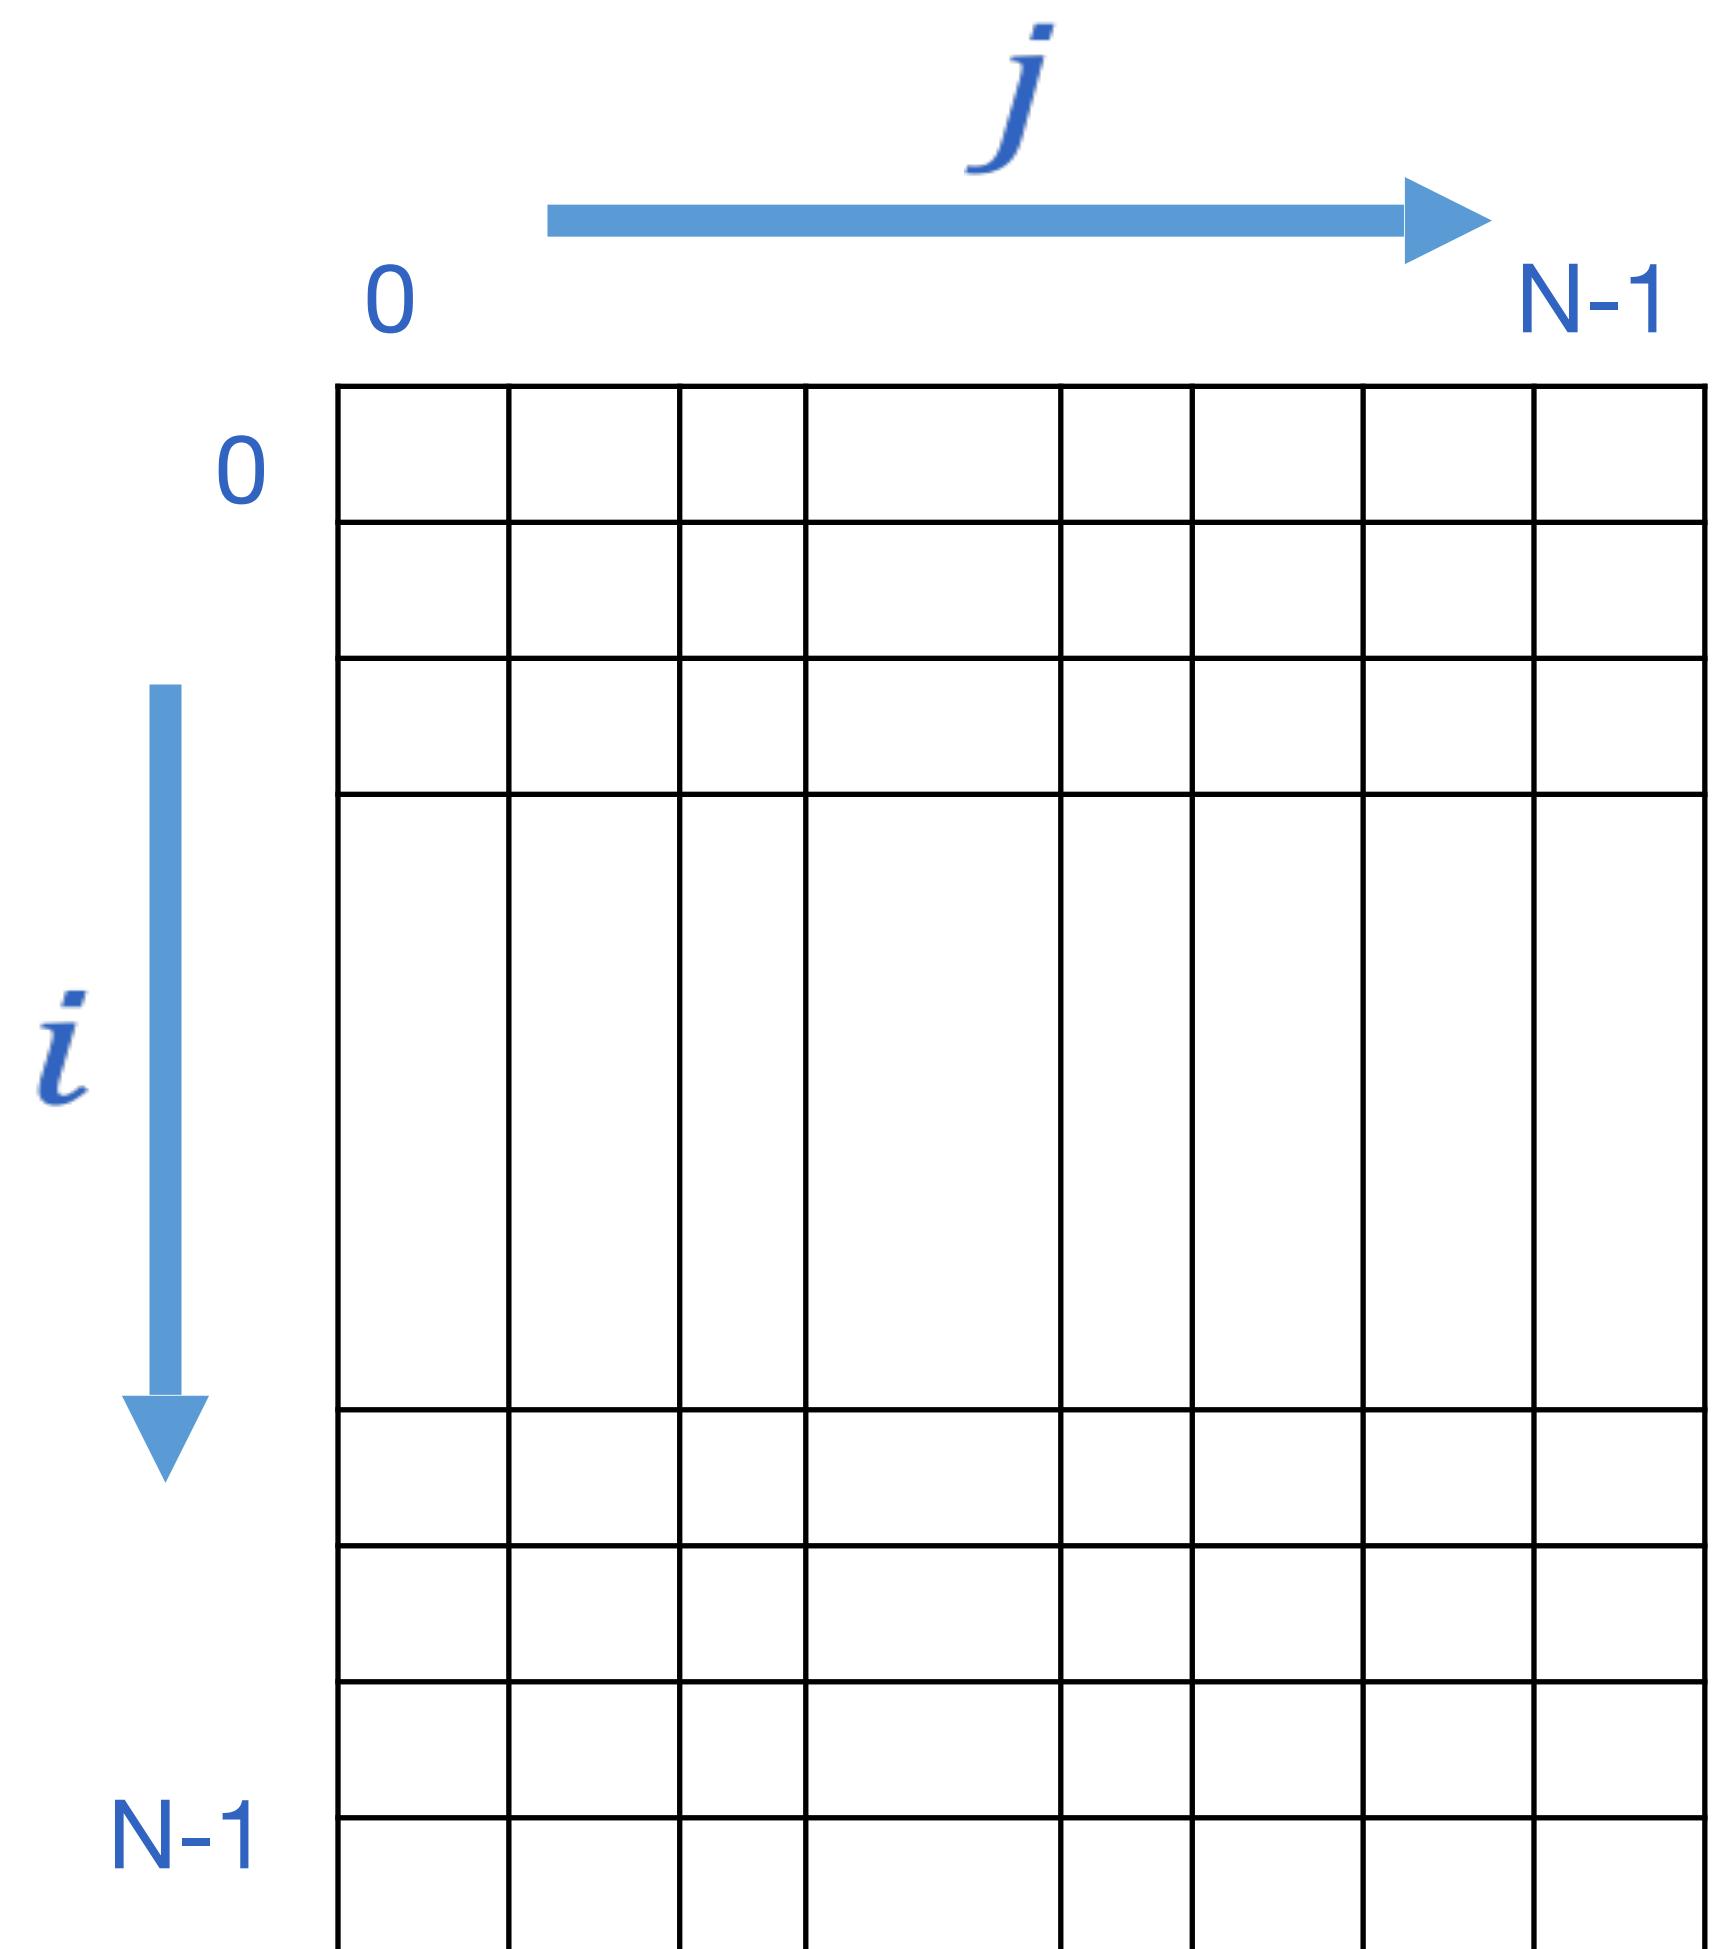
# Reference Problem

- Dense matrix multiplication
  - Basic operation in many engineering, data, and imaging processing tasks
    - Ex:, Image filtering, noise reduction, ...
  - Core operation in Neural Nets and Deep Learning
    - Image classification
    - Robot Cars
    - Machine translation
    - Fingerprint verification
    - Automatic game playing
- **dgemm**
  - double-precision floating-point general matrix-multiply
  - Standard well-studied and widely used routine
    - Part of Linpack/Lapack



# 2D-Matrices

- Square matrix of dimension  $N \times N$ 
  - $i$  indexes through rows
  - $j$  indexes through columns



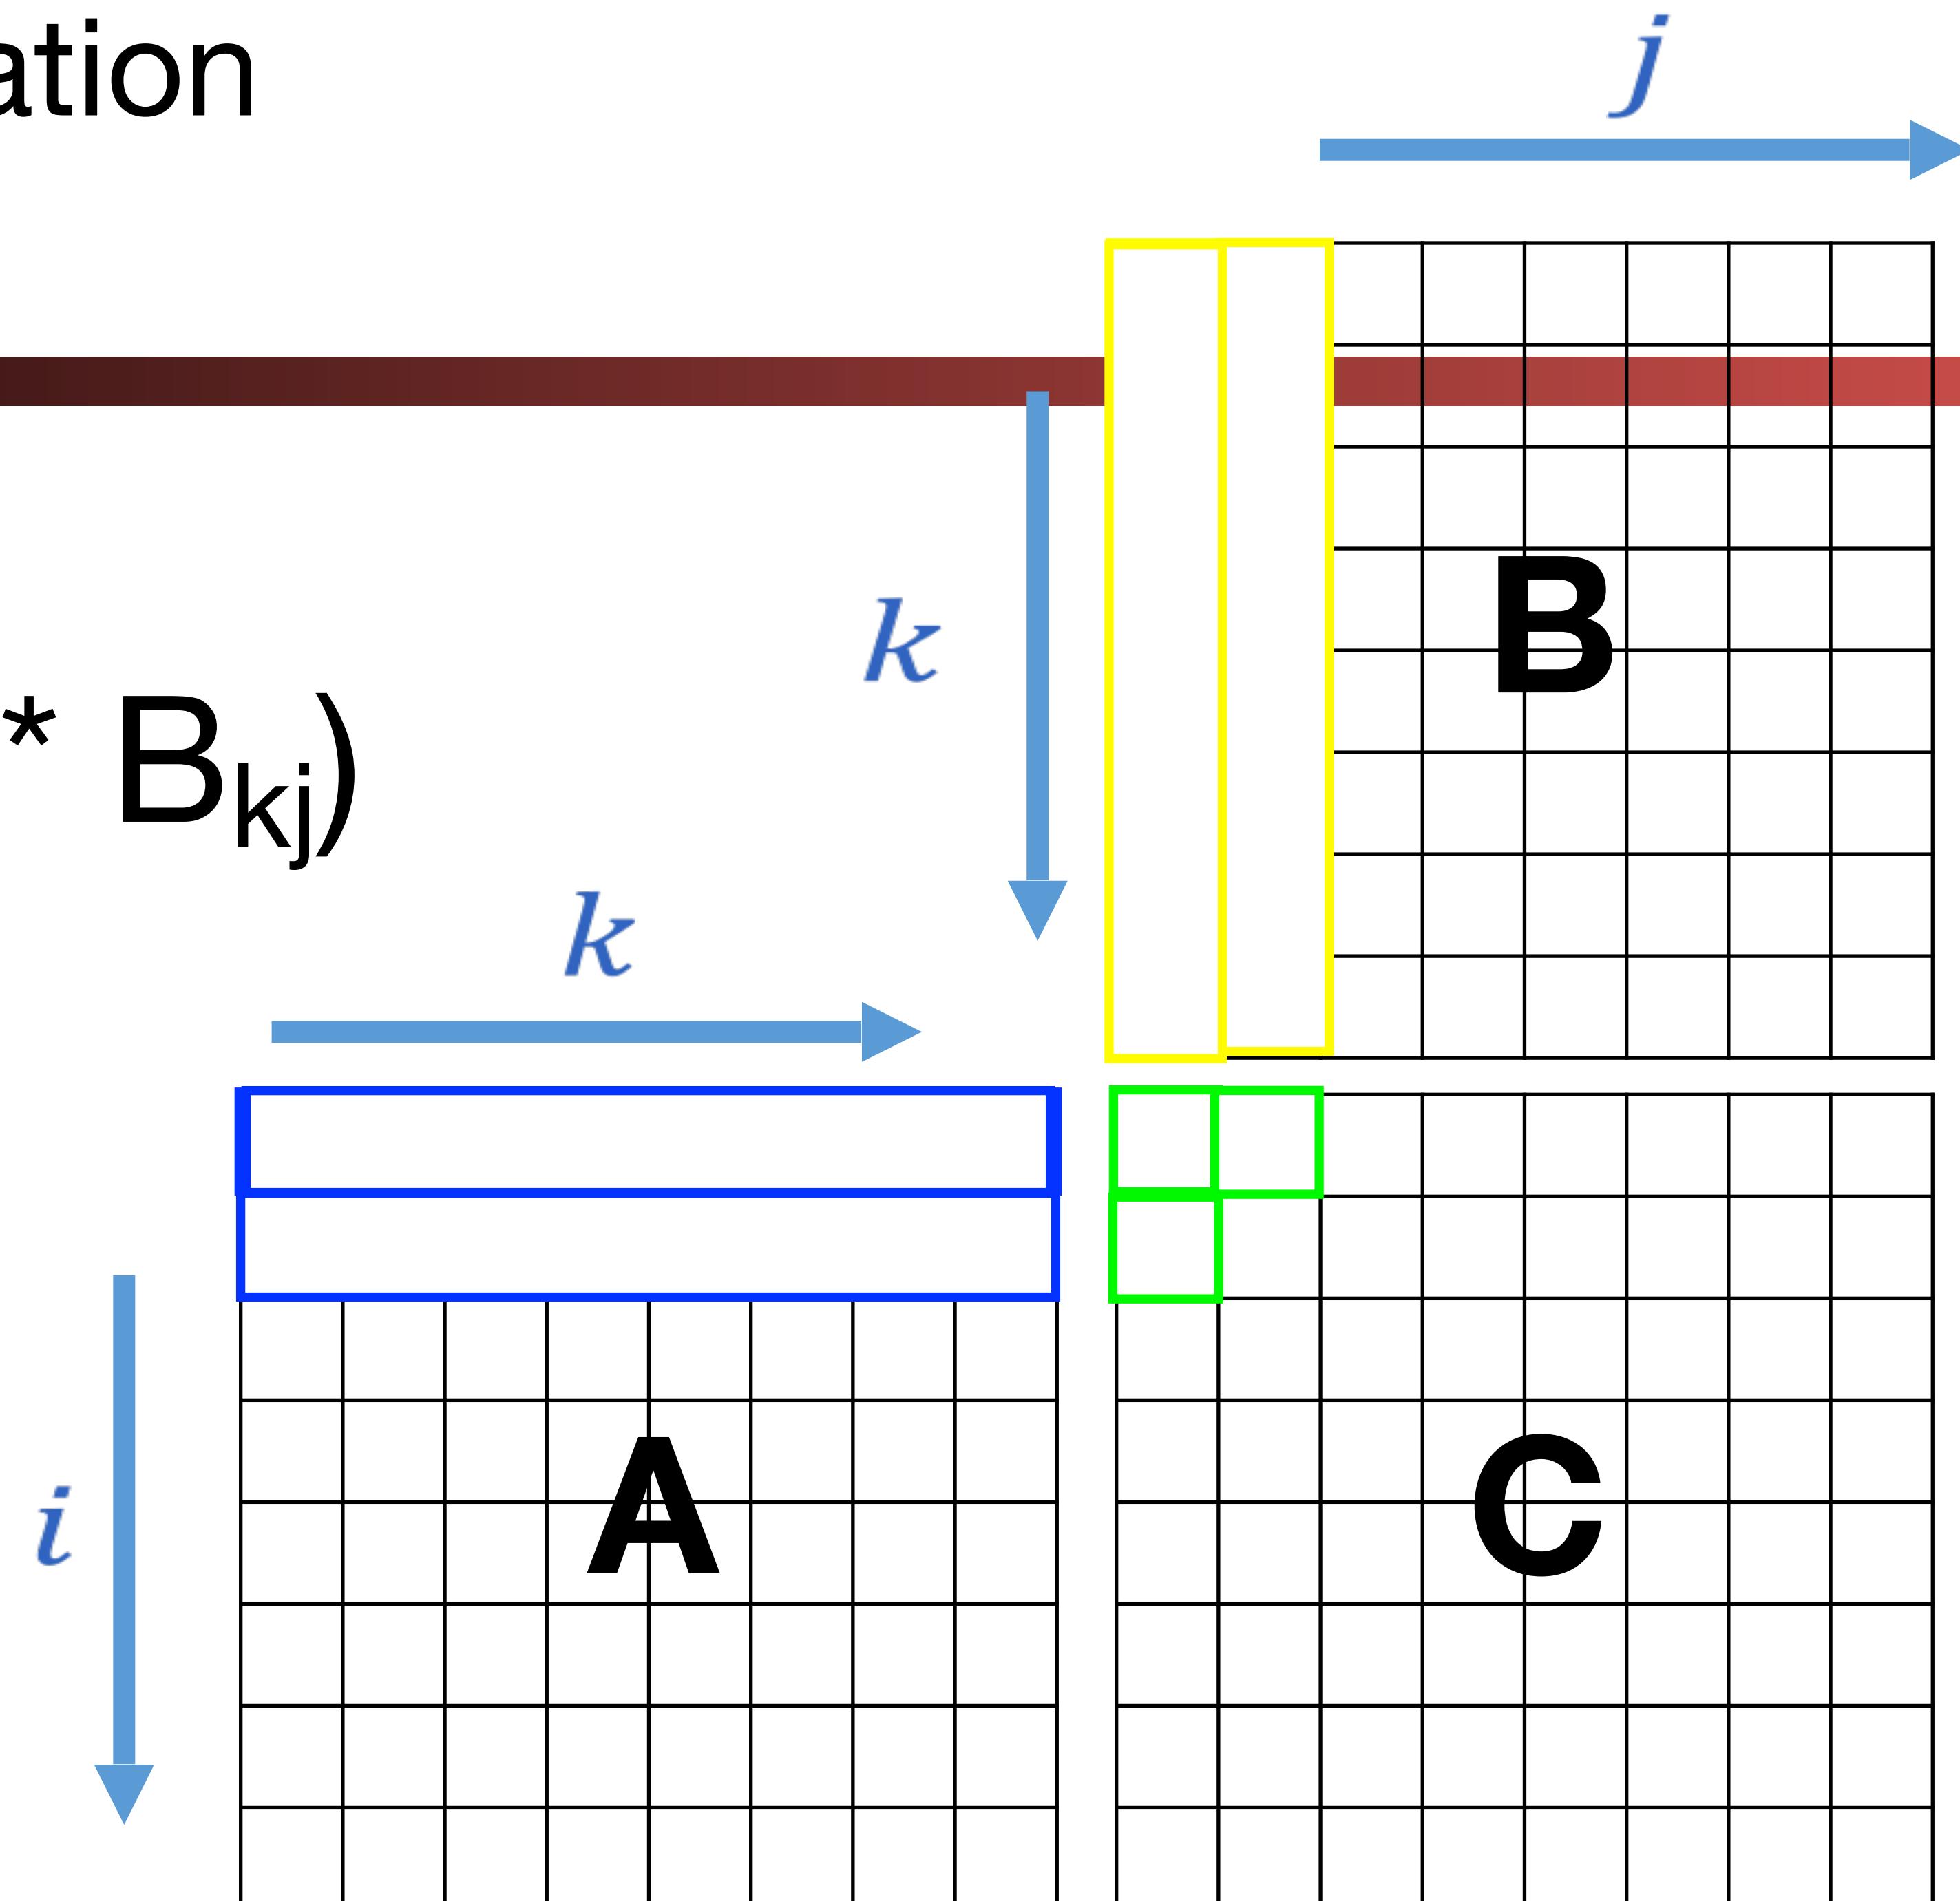
# Matrix Multiplication

Computer Science 61C

McMahon & Weaver

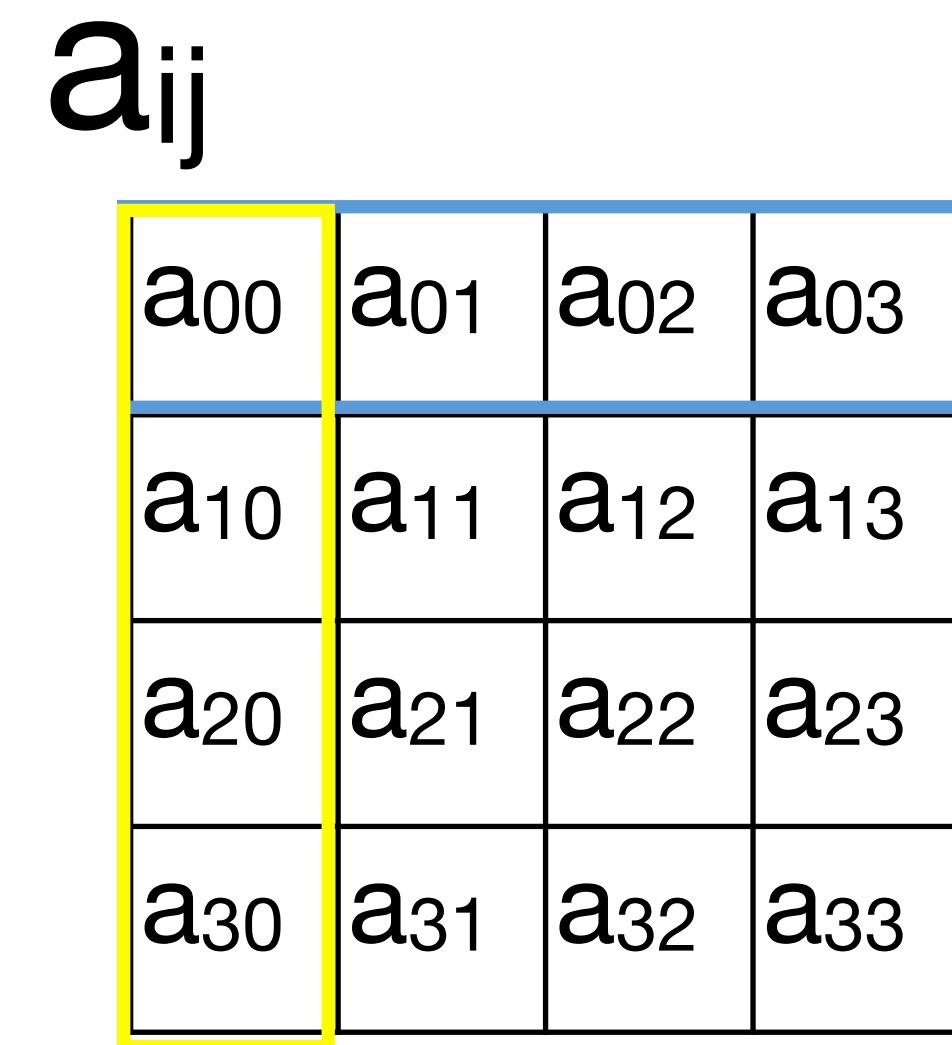
$$C = A * B$$

$$C_{ij} = \sum_k (A_{ik} * B_{kj})$$



# 2D Matrix Memory Layout

- $a[ ][ ]$  in C uses row-major
- Fortran uses column-major
- Our examples use column-major



Row-Major

Row

Row

 $a_{ij} : a[i*N + j]$ 

Column-Major

Column

$a_{31}$	
$a_{21}$	
$a_{11}$	
$a_{01}$	
$a_{30}$	
$a_{20}$	
$a_{10}$	
$a_{00}$	

 $a_{ij} : a[i + j*N]$

# Simplifying Assumptions...

- We want to keep the examples (somewhat) manageable...
- We will keep the matrixes square
  - So both matrixes are the same size with the same number of rows and columns
- We will keep the matrixes reasonably aligned
  - So size % a reasonable power of 2 == 0
- We are doing **dense** matrix multiplication
  - A related problem is "sparse" matrix multiplication, where most of the entries are 0

# **dgemm** Reference Code: Python

```
def dgemm(N, a, b, c):
    for i in range(N):
        for j in range(N):
            c[i+j*N] = 0
            for k in range(N):
                c[i+j*N] += a[i+k*N] * b[k+j*N]
```

N	Python [Mflops]
32	5.4
160	5.5
480	5.4
960	5.3

- 1 MFLOP = 1 Million floating-point operations per second (fadd, fmul)
- **dgemm(N ...)** takes  $2 \cdot N^3$  flops

- $C = A * B$
- $A, B, C$  are  $N \times N$  matrices

```
void dgemm_scalar(int N, double *a, double *b, double *c) {
    int i,j,k; double cij;
    for(i = 0; i < N; ++i) {
        for(j = 0; j < N; ++j) {
            cij = 0
            for(k = 0; k < N; ++k) {
                cij += a[i+k*N] * b[k+j*N];
            }
            c[i+j*N] = cij;
        }
    }
}
```

# Timing Program Execution

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(void) {
    // start time
    // Note: clock() measures execution time, not real time
    //       big difference in shared computer environments
    //       and with heavy system load
    clock_t start = clock();

    // task to time goes here:
    // dgemm(N, ...);

    // "stop" the timer
    clock_t end = clock();

    // compute execution time in seconds
    double delta_time = (double)(end-start)/CLOCKS_PER_SEC;
}
```

# C versus Python

N	C [GFLOPS]	Python [GFLOPS]
32	1.30	0.0054
160	1.30	0.0055
480	1.32	0.0054
960	0.91	0.0053



**Which other class gives you this kind of power?  
We could stop here ... but why? Let's do better!**

# Agenda

- 61C – the big picture
- **Parallel processing**
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

# Why Parallel Processing?

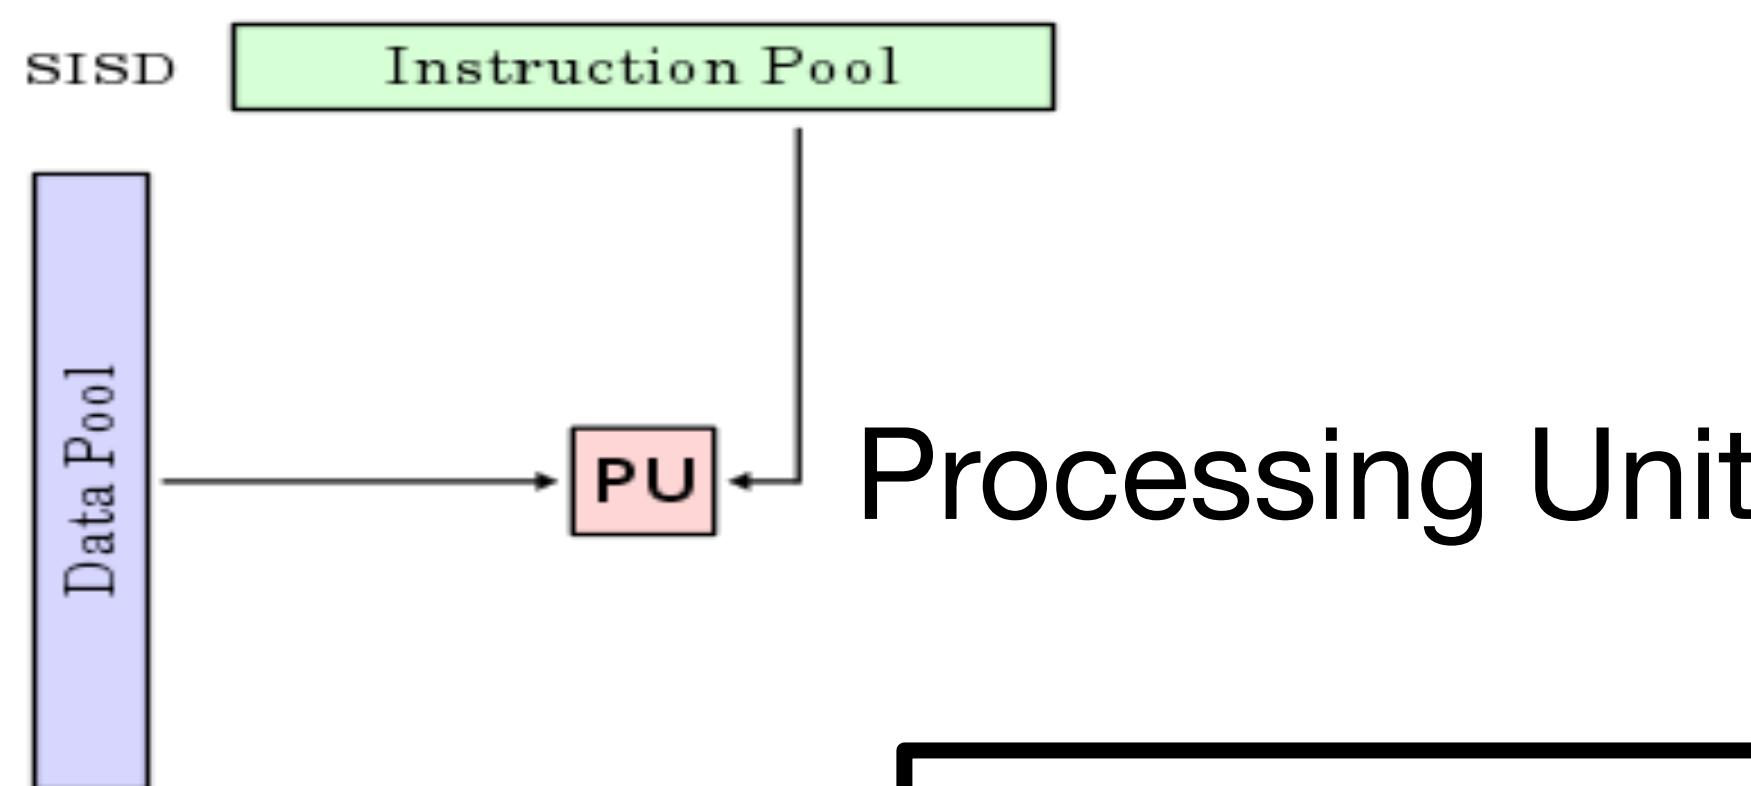
- CPU Clock Rates are no longer increasing
  - Technical & economic challenges
    - Advanced cooling technology too expensive or impractical for most applications
    - Energy costs are prohibitive
- Parallel processing is only path to higher speed
  - Compare airlines:
    - Maximum air-speed limited by economics
    - Use more and larger airplanes to increase throughput
    - (And smaller seats ...)

# Using Parallelism for Performance

- Two basic approaches to parallelism:
  - Multiprogramming
    - run multiple independent programs in parallel
    - “Easy”
  - Parallel computing
    - run one program faster
    - “Hard”
- We'll focus on parallel computing in the next few lectures

# Single-Instruction/Single-Data Stream (SISD)

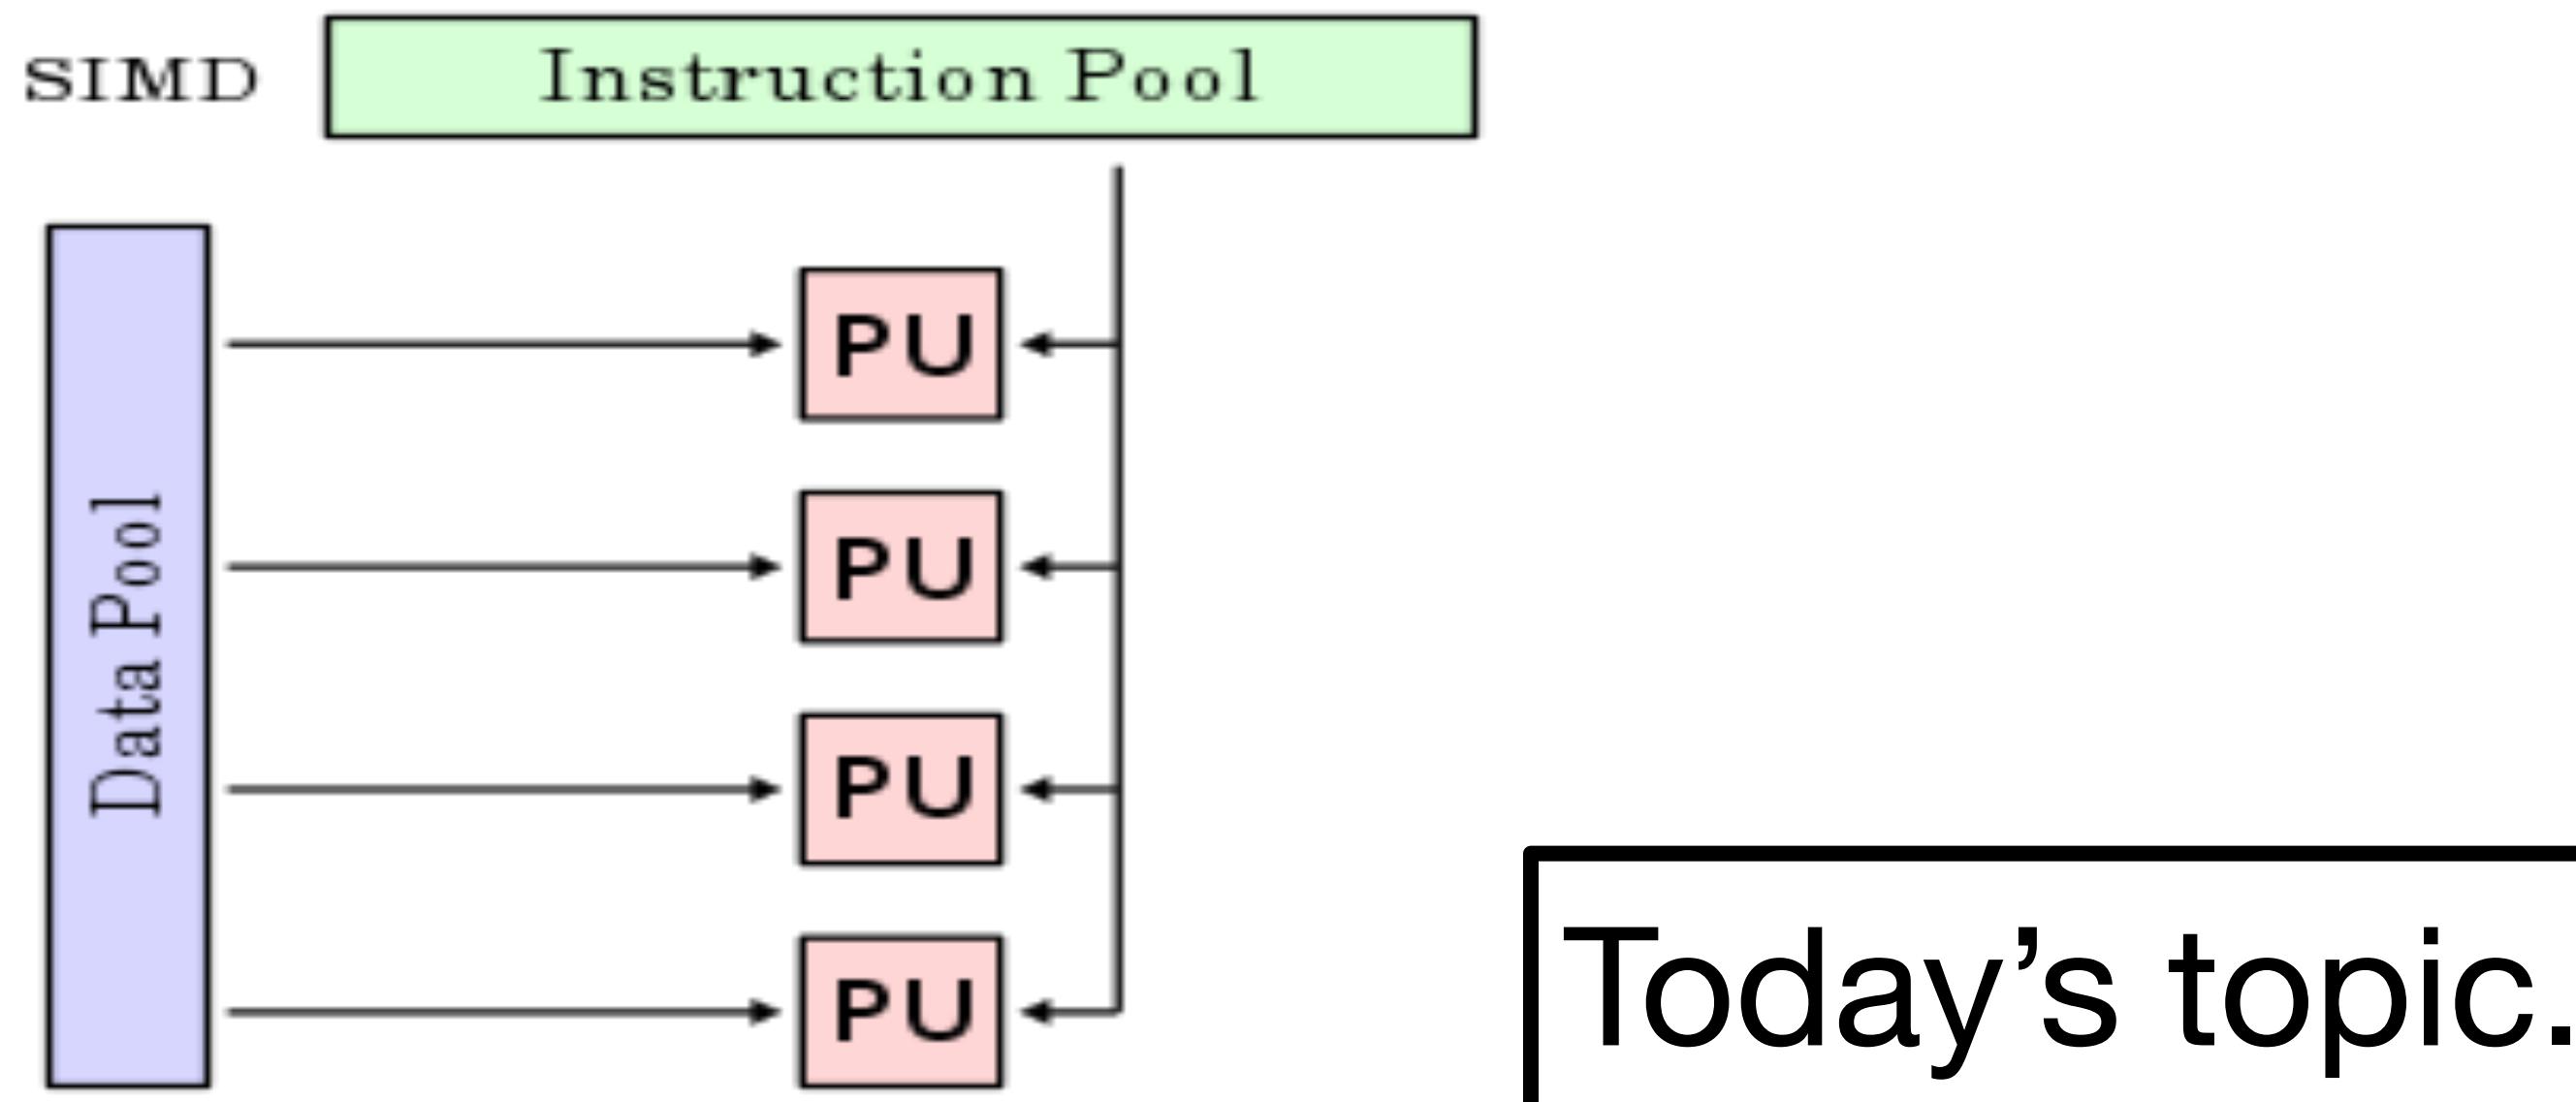
- Sequential computer that exploits no parallelism in either the instruction or data streams. Examples of SISD architecture are traditional uniprocessor machines
  - E.g. Our RISC-V processor
  - We consider superscalar as SISD because the ***programming model*** is sequential



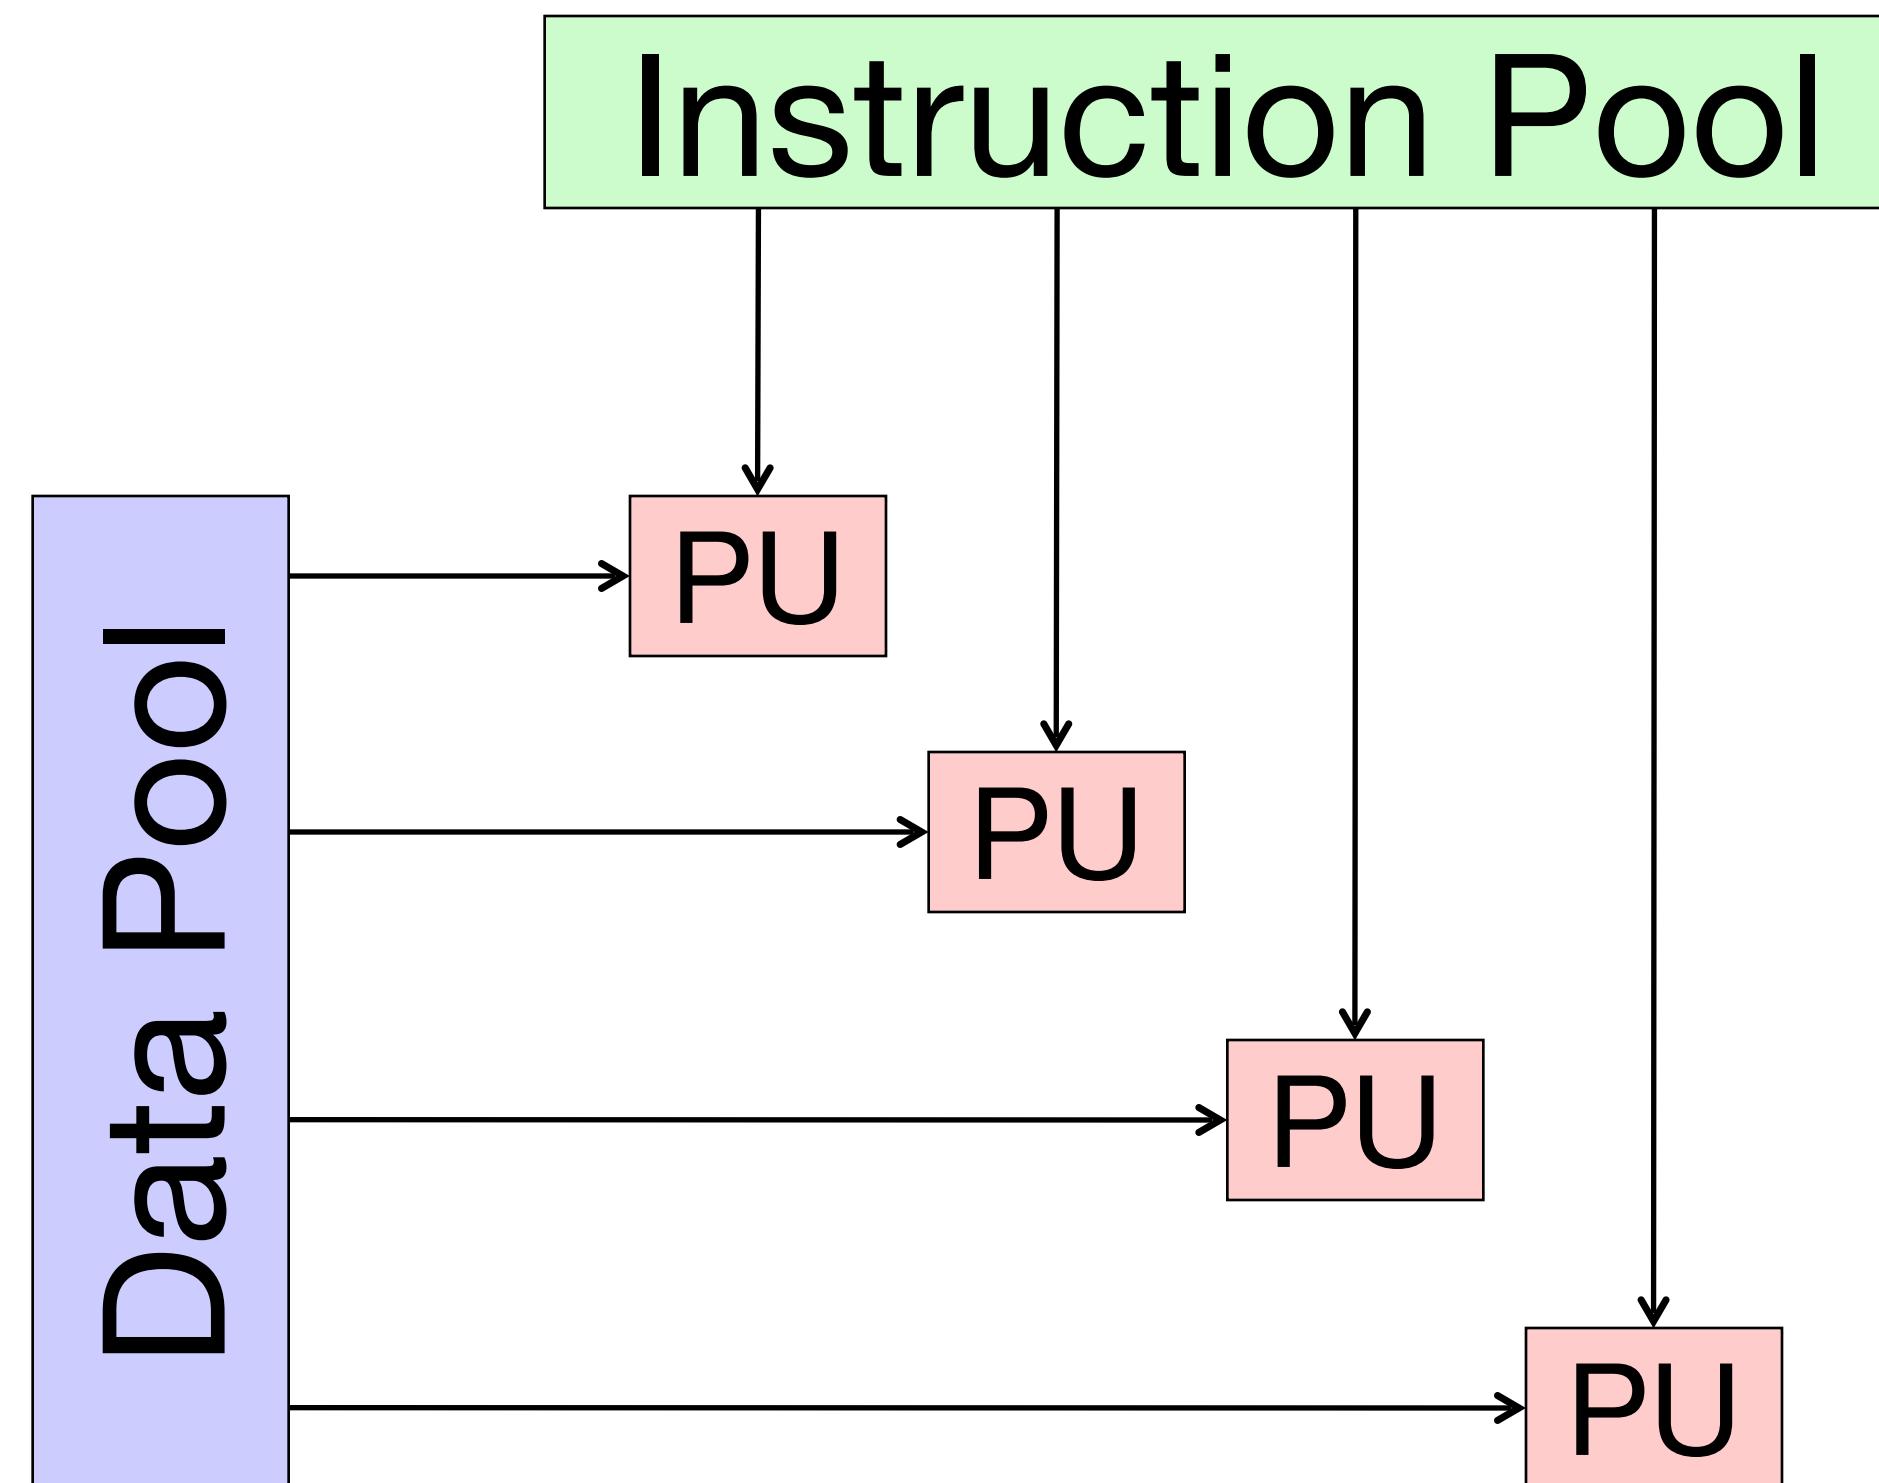
This is what we did up to now in 61C

# Single-Instruction/Multiple-Data Stream (SIMD or “sim-dee”)

- SIMD computer processes multiple data streams using a single instruction stream, e.g., Intel SIMD instruction extensions or NVIDIA Graphics Processing Unit (GPU)



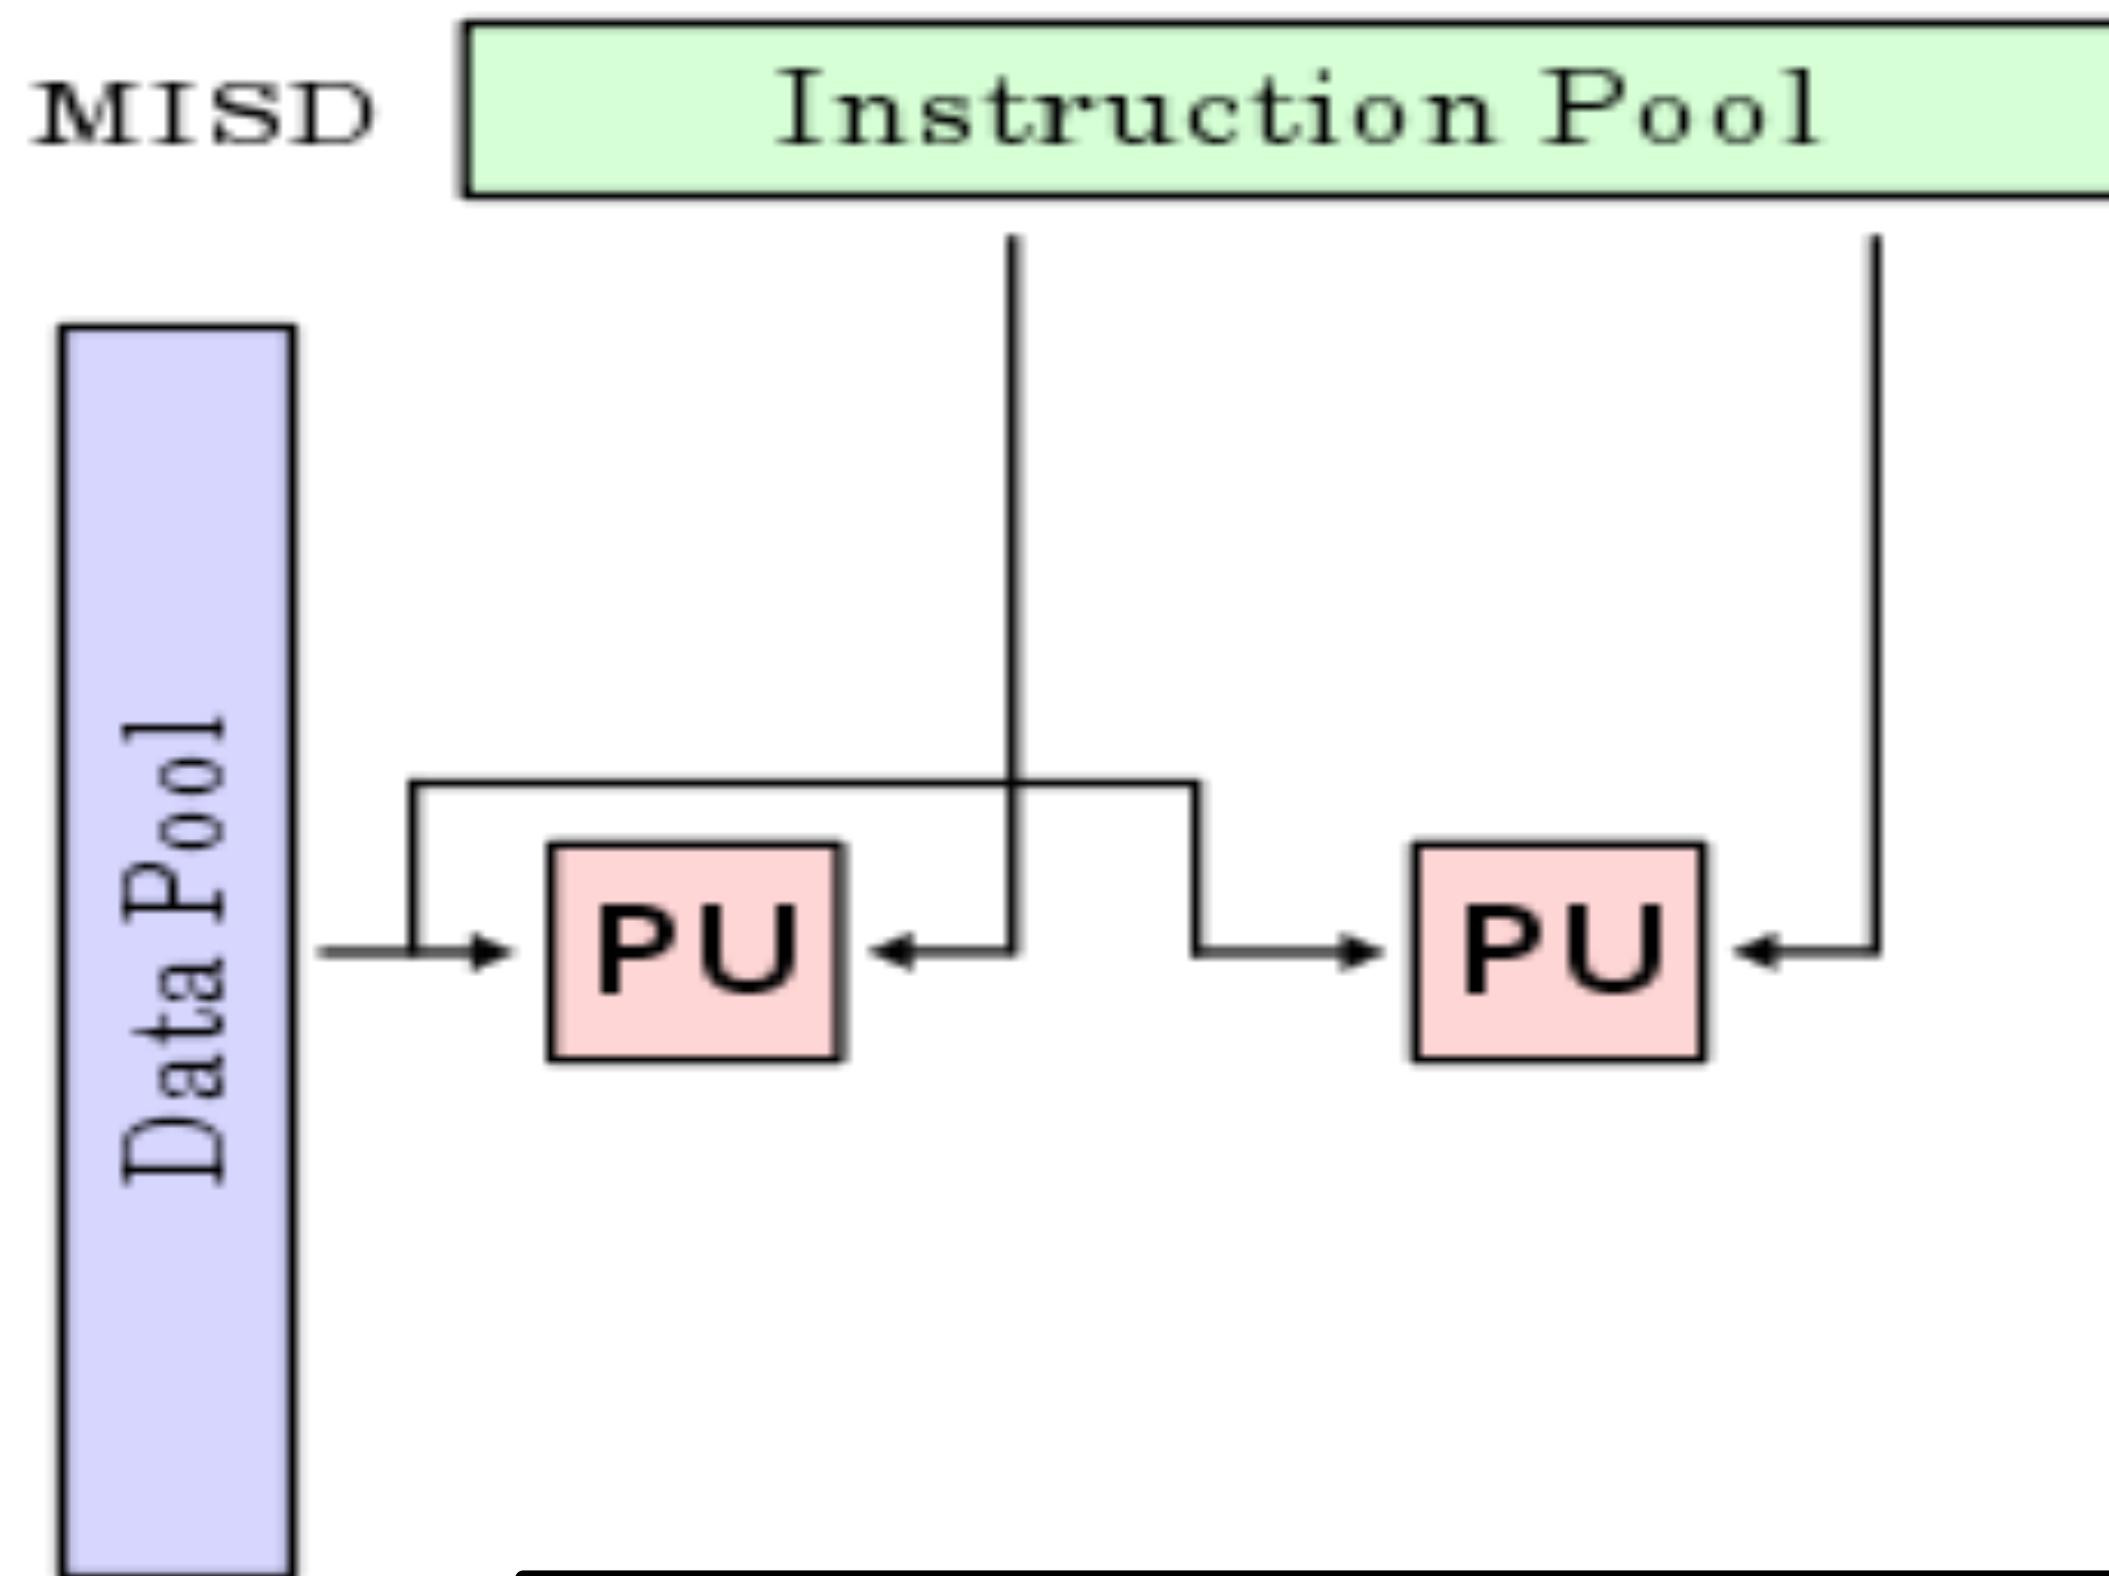
# Multiple-Instruction/Multiple-Data Streams (MIMD or “mim-dee”)



- Multiple autonomous processors simultaneously executing different instructions on different data.
  - MIMD architectures include multicore and Warehouse-Scale Computers

Topic of Lecture 22 and beyond.

# Multiple-Instruction/Single-Data Stream (MISD)



- Multiple-Instruction, Single-Data stream computer that processes multiple instruction streams with a single data stream.
- Historical significance

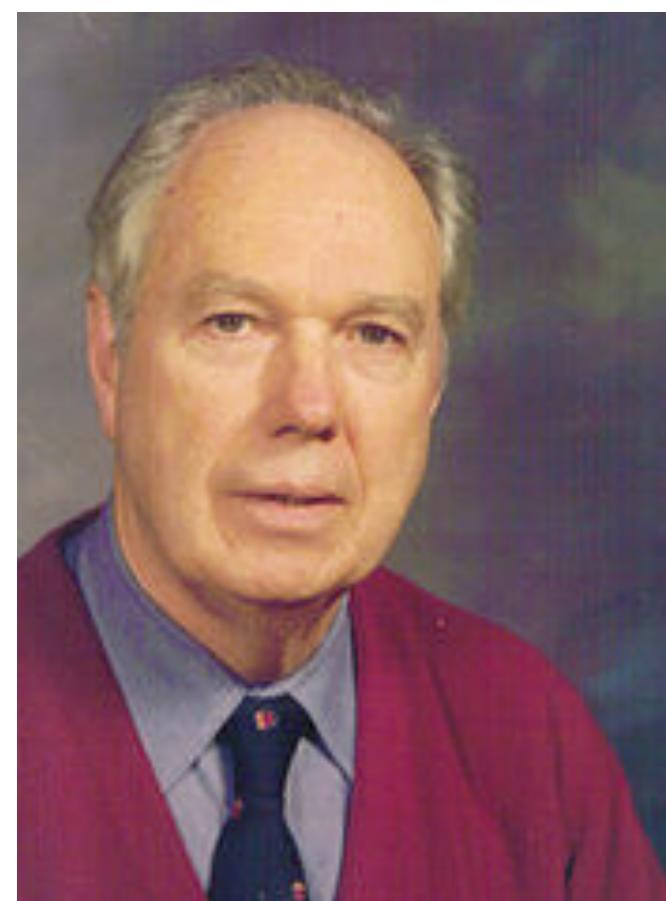
This has few applications. Not covered in 61C.

# Flynn\* Taxonomy, 1966

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345 (Clovertown)

- SIMD and MIMD are currently the most common parallelism in architectures – usually both in same system!
- Most common parallel processing programming style: Single Program Multiple Data (“SPMD”)
  - Single program that runs on all processors of a MIMD
  - Cross-processor execution coordination using synchronization primitives

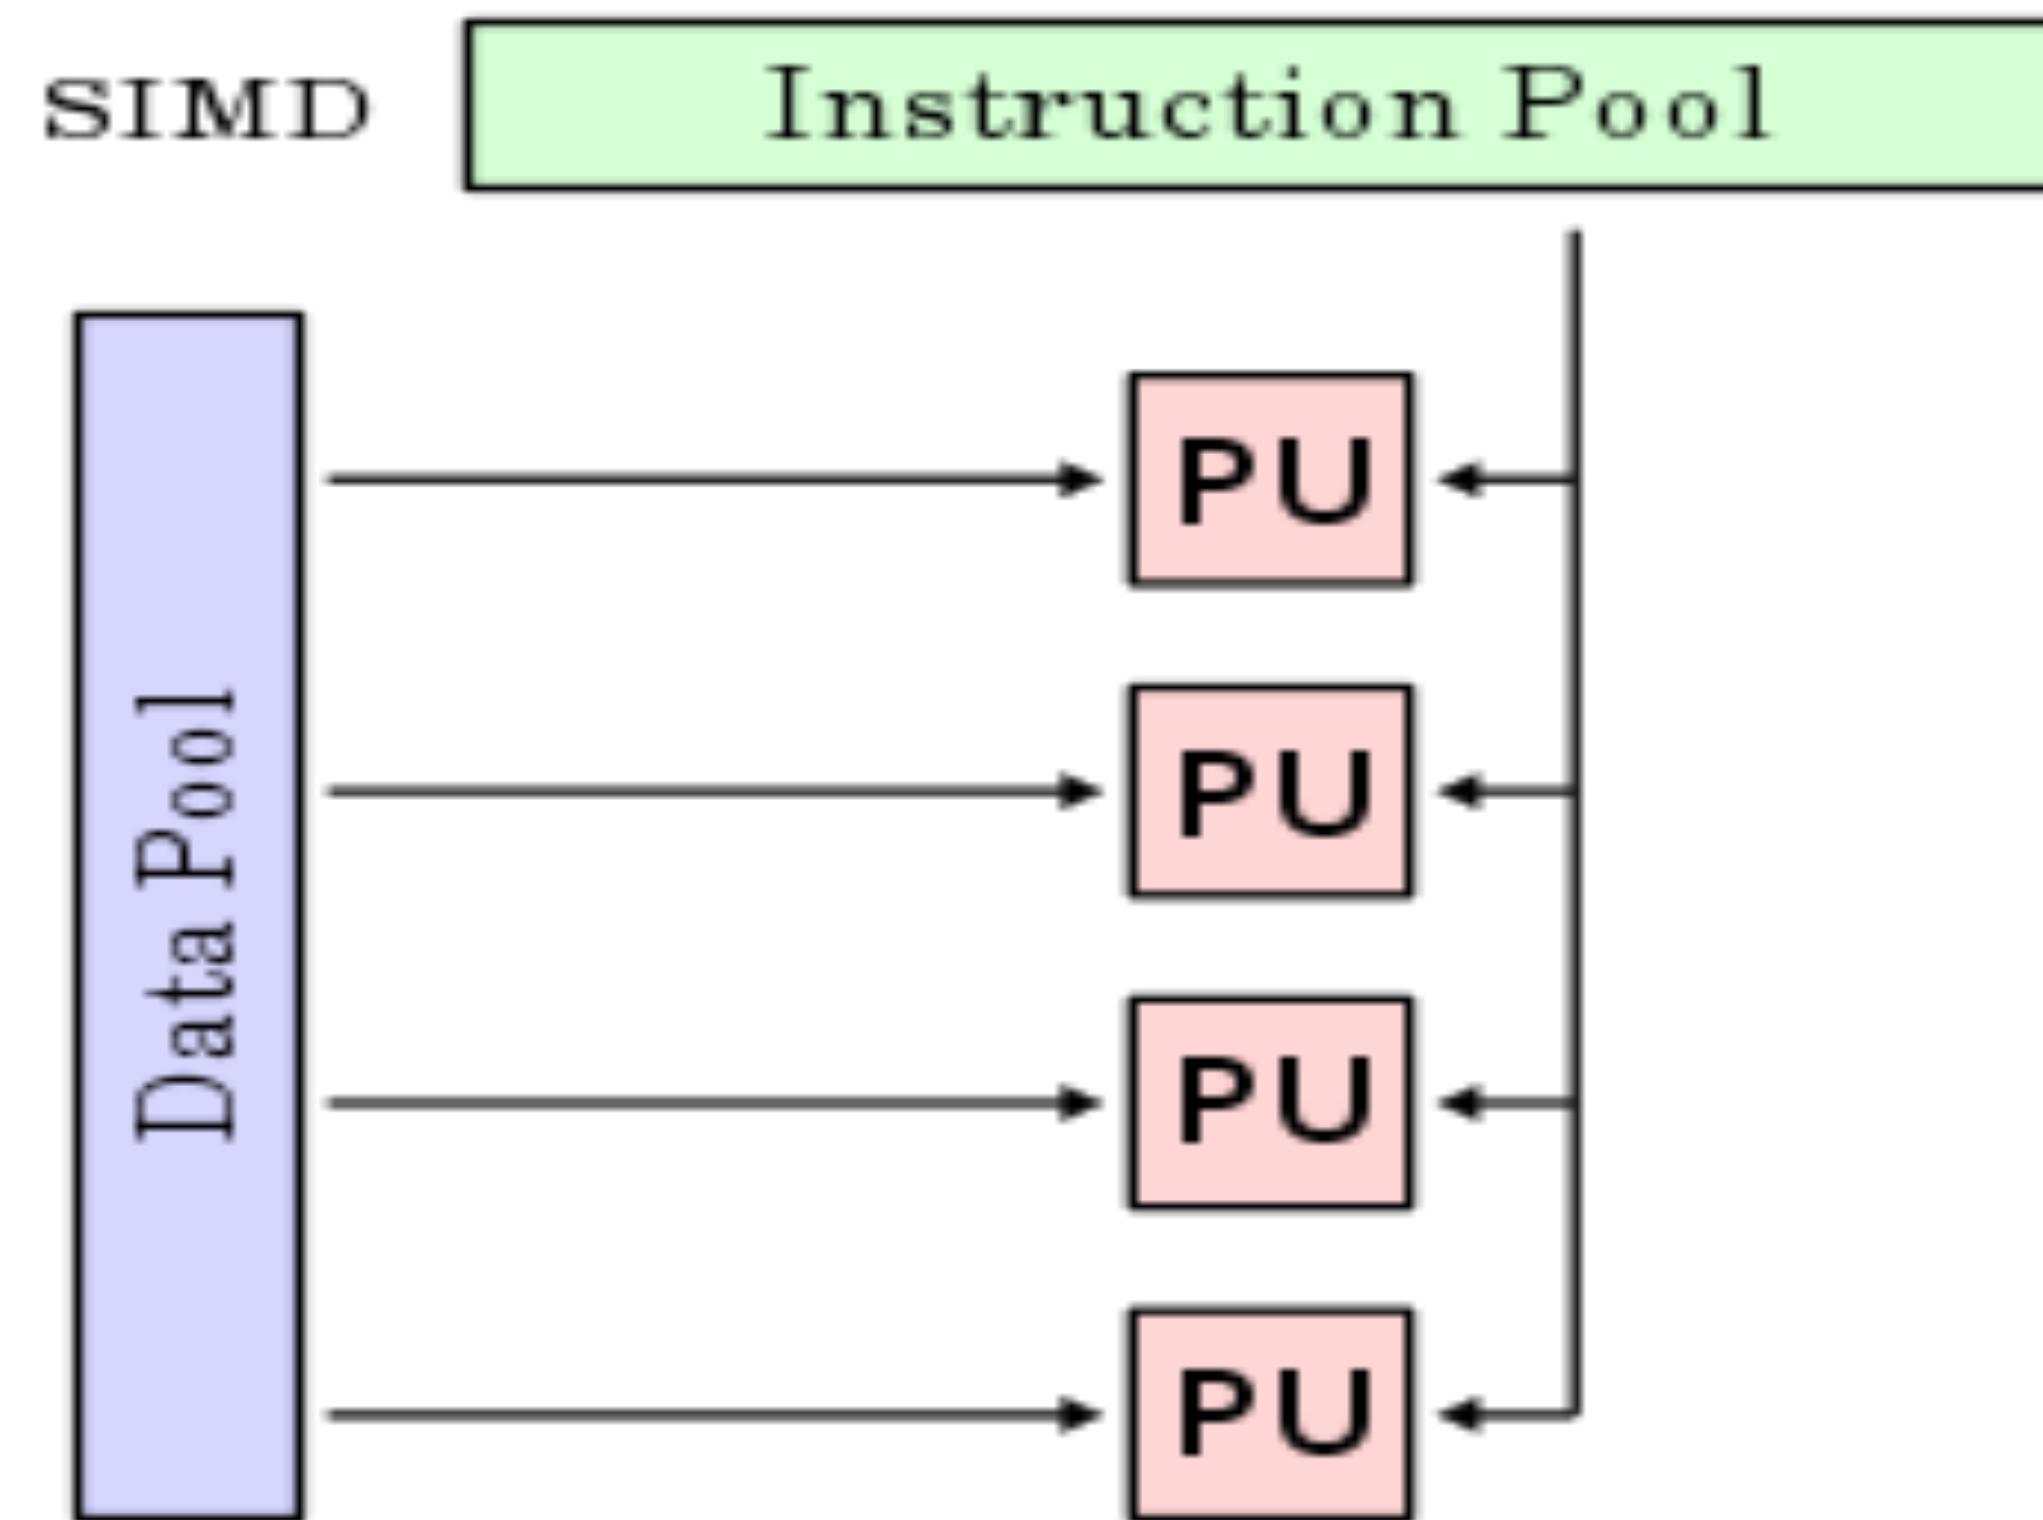
\*Prof. Michael Flynn, Stanford



# Agenda

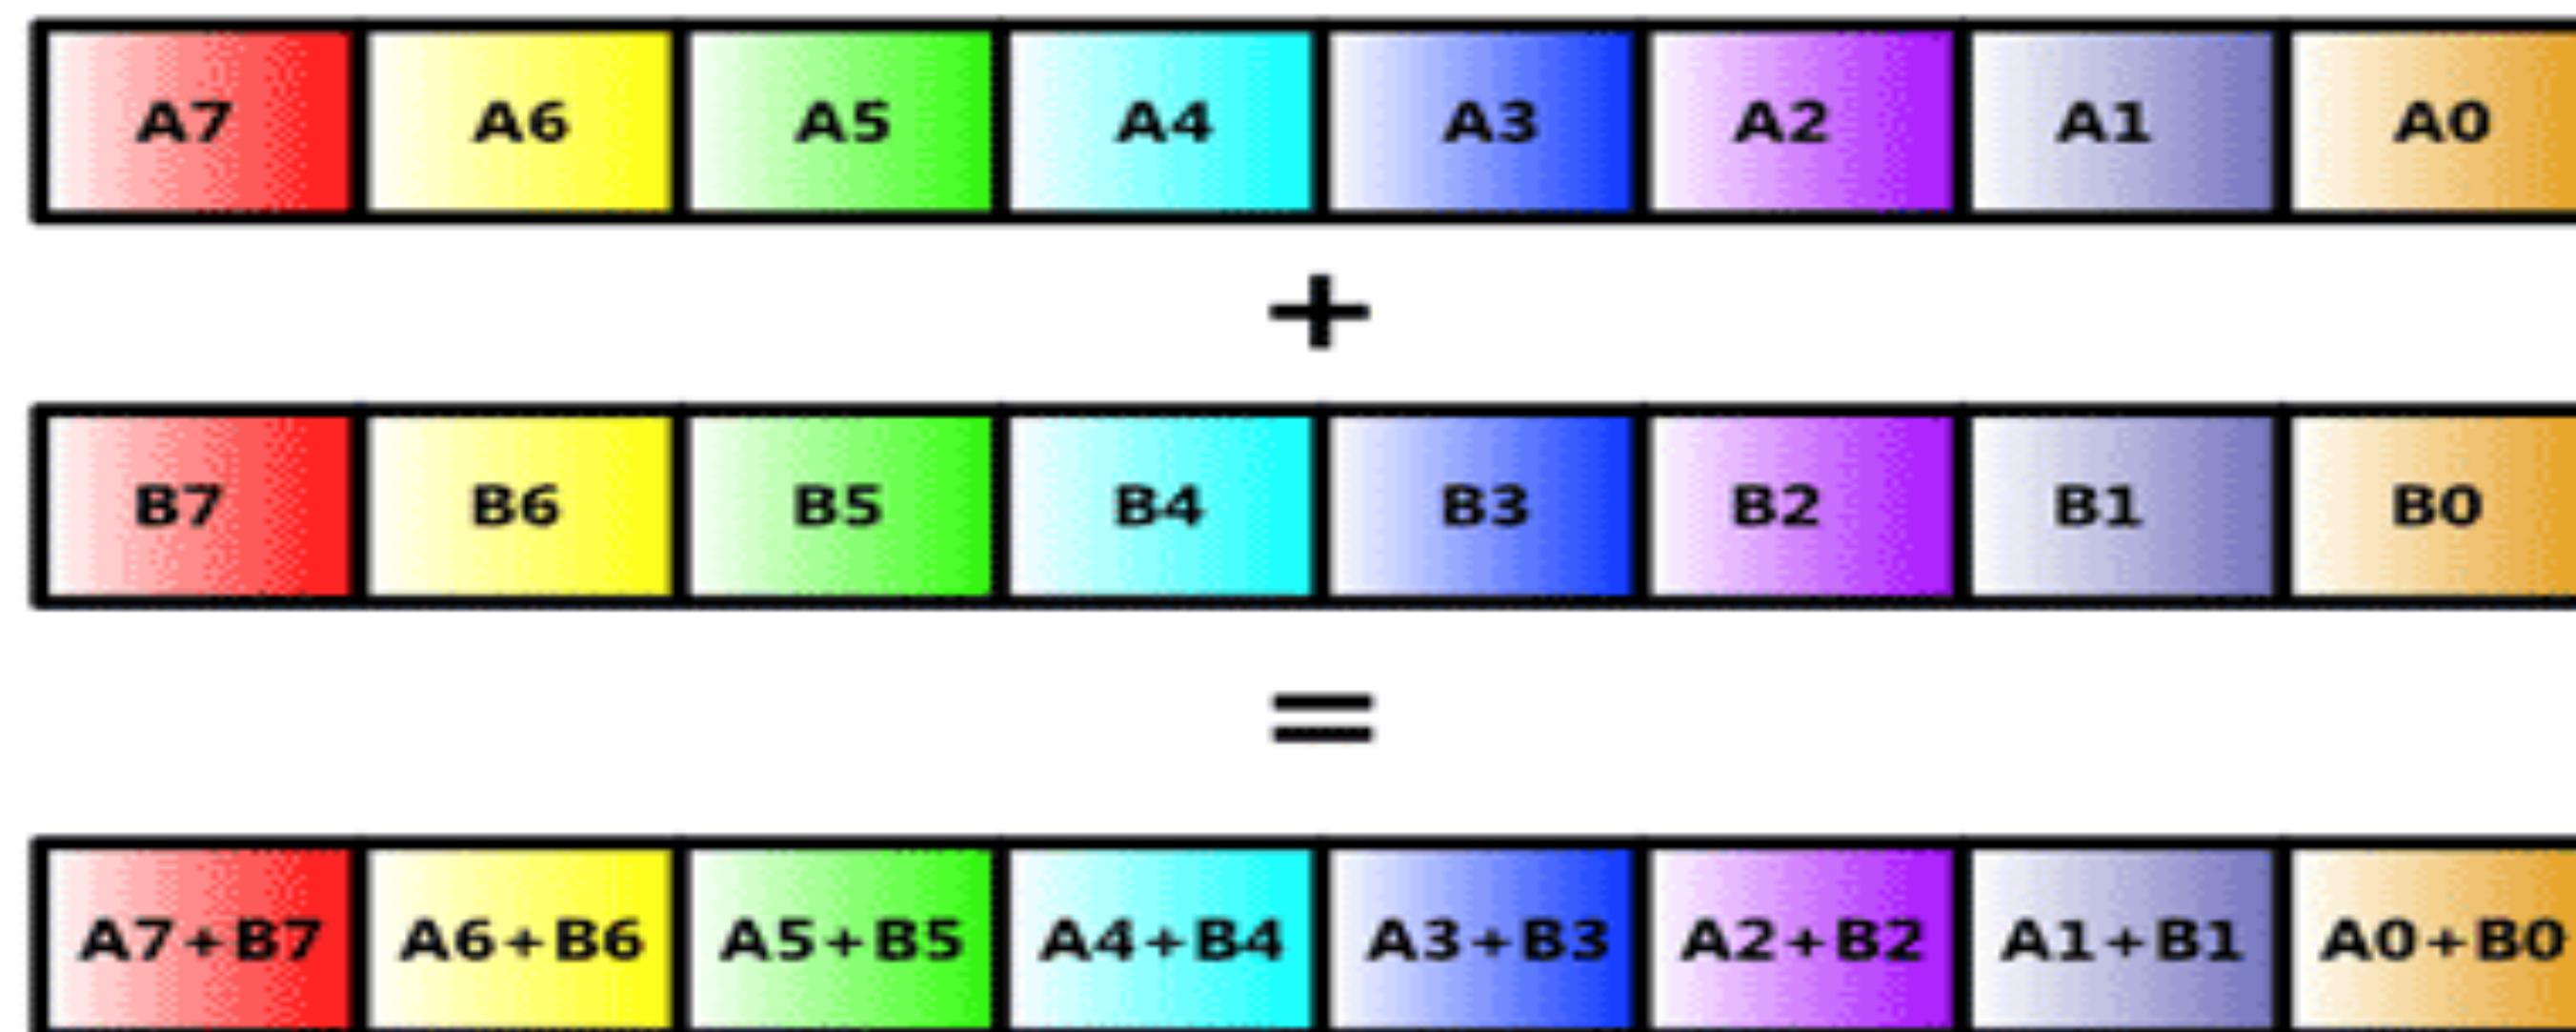
- 61C – the big picture
- Parallel processing
- **Single instruction, multiple data**
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

# SIMD – “Single Instruction Multiple Data”

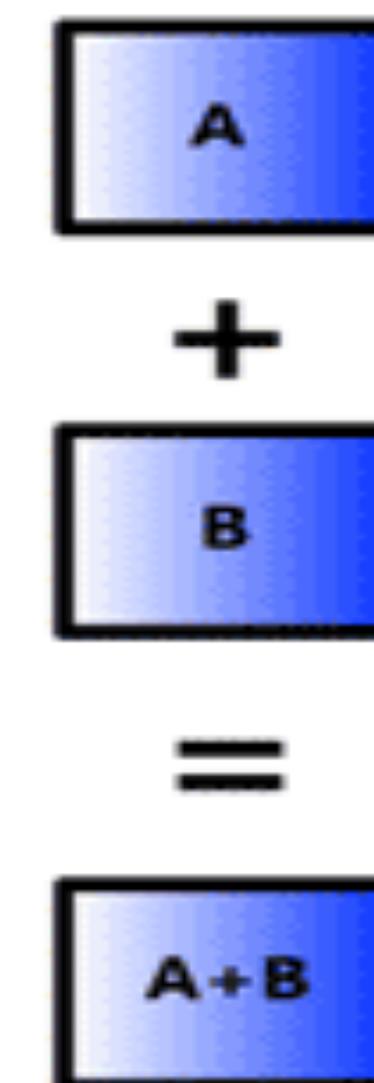


# SIMD (Vector) Mode

## SIMD Mode



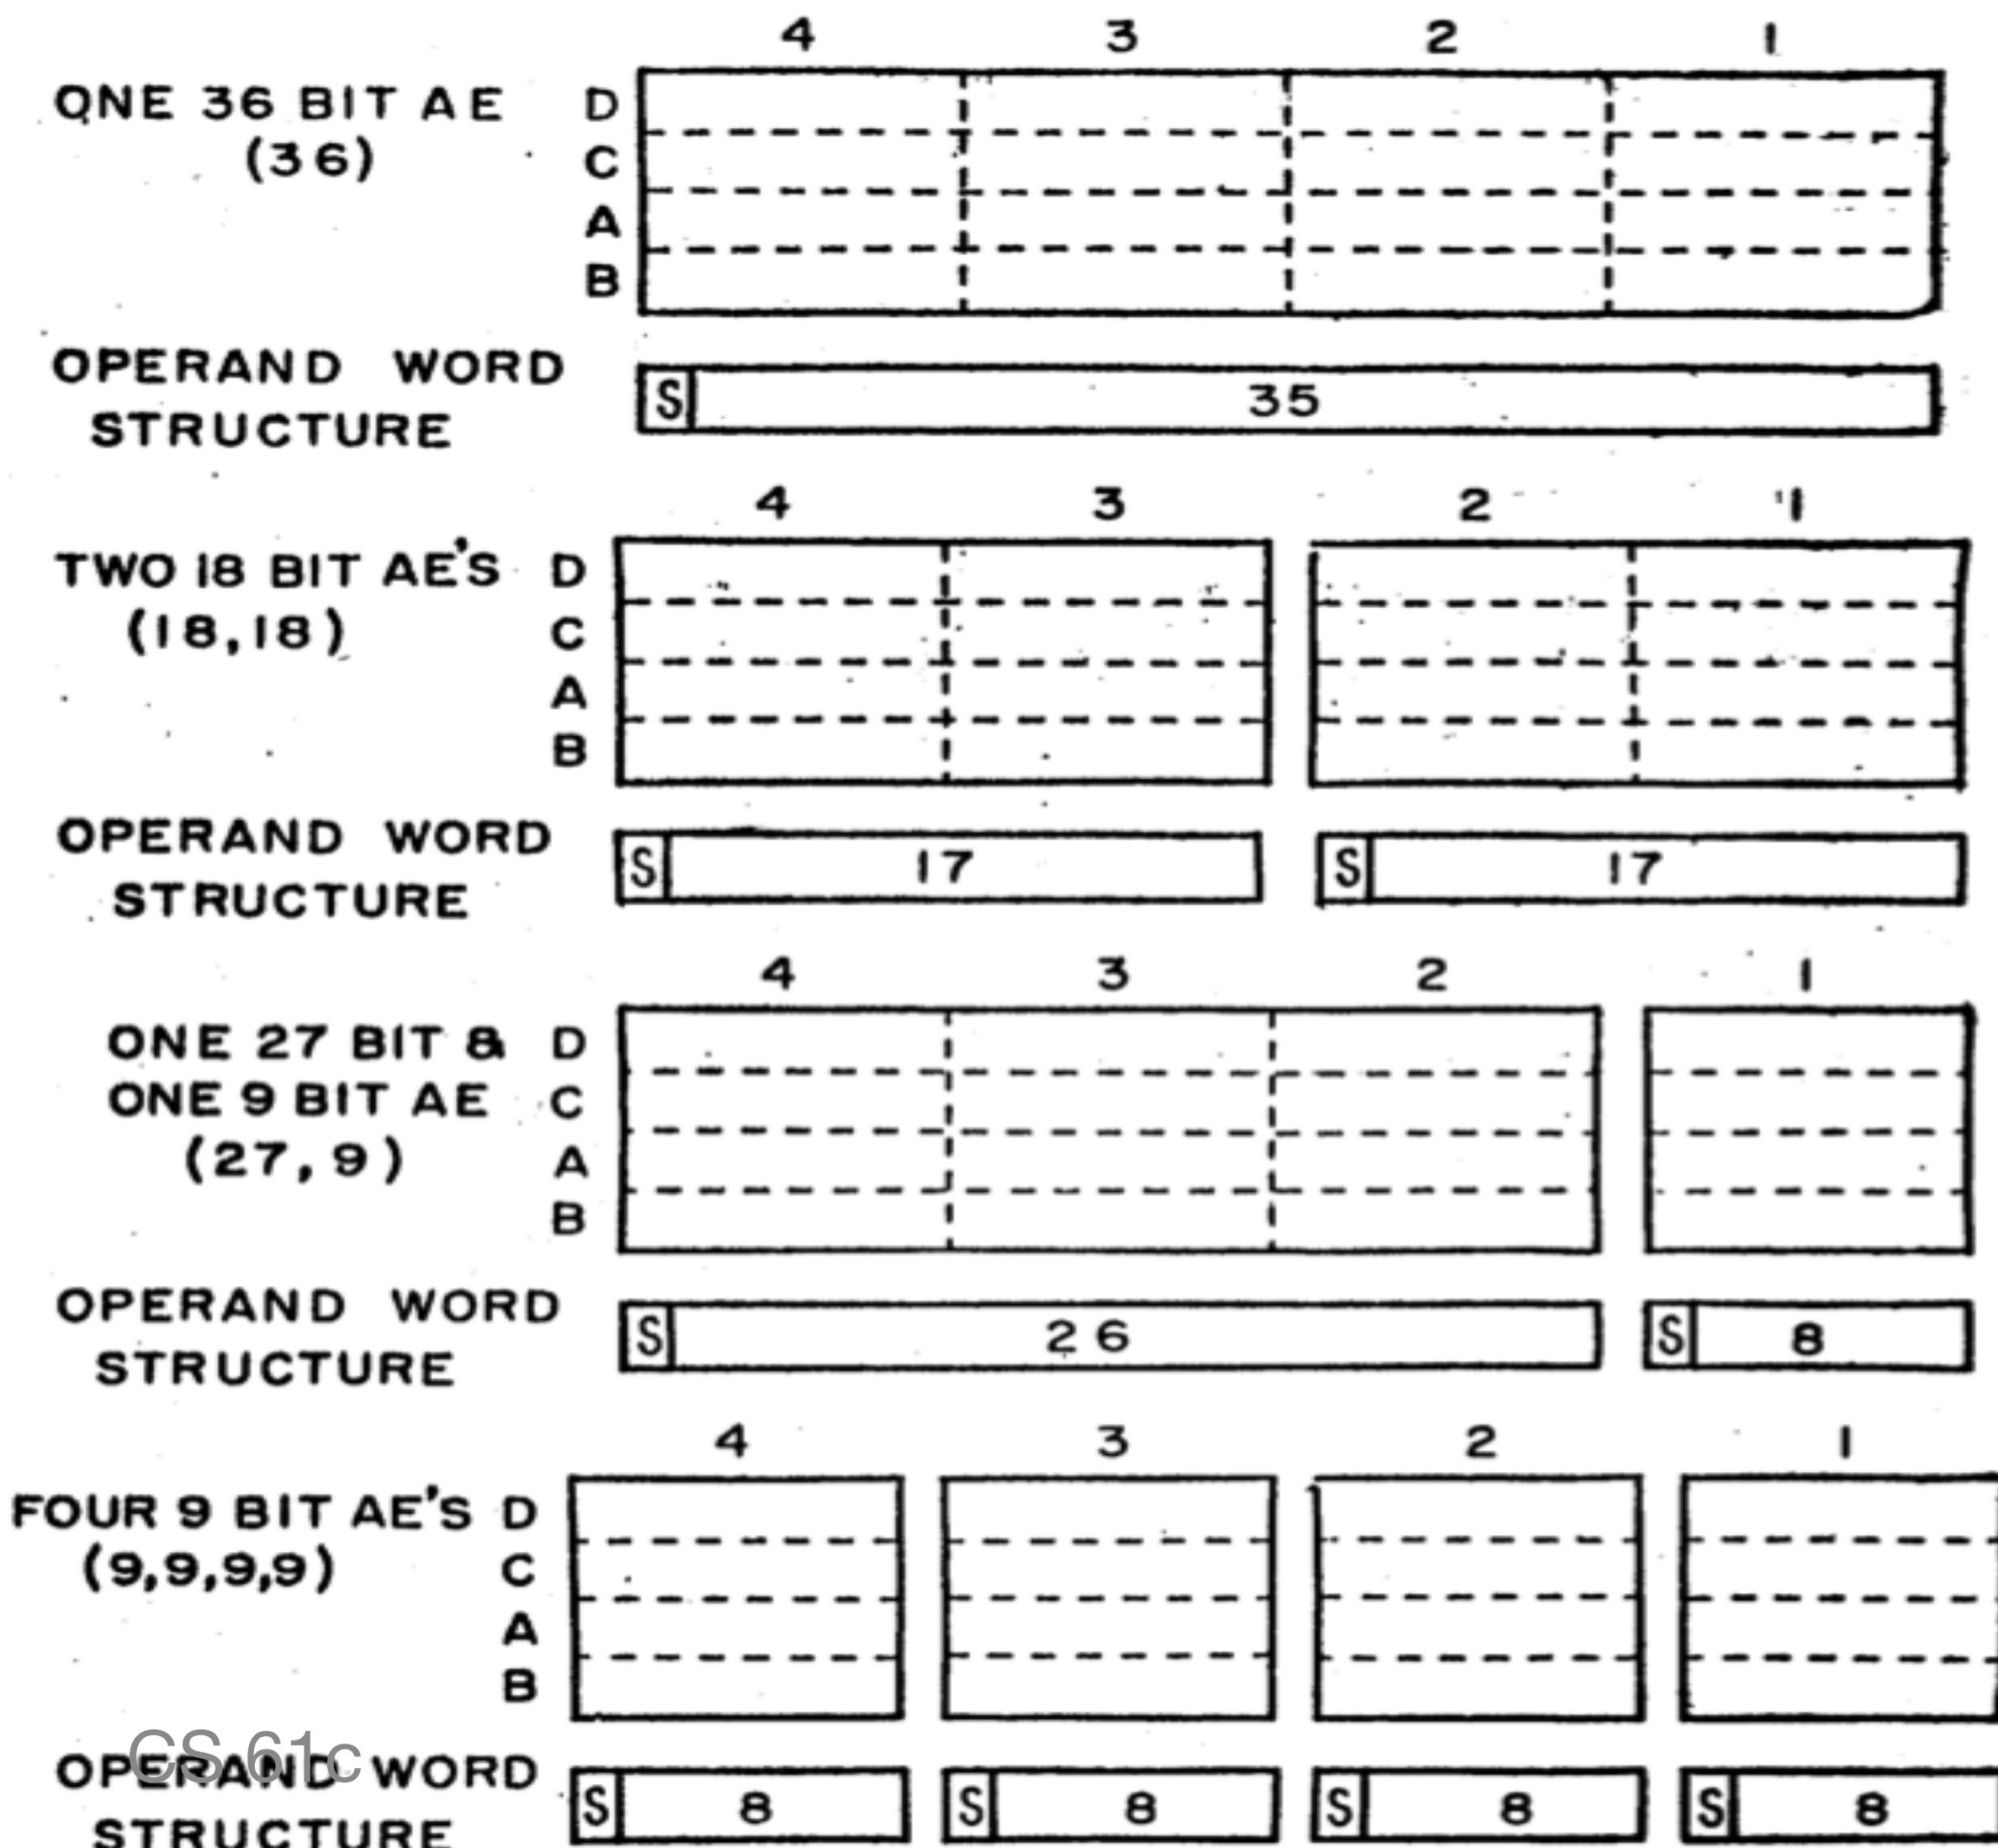
## Scalar Mode



# SIMD Applications & Implementations

- Applications
  - Scientific computing
    - Matlab, NumPy
  - Graphics and video processing
    - Photoshop, ...
  - Big Data
    - Deep learning
  - Gaming
- Implementations
  - x86
  - ARM
  - RISC-V vector extensions
  - Video cards

# First SIMD Extensions: MIT Lincoln Labs TX-2, 1957



# Intel x86 SIMD: Continuous Evolution

## MMX 1997

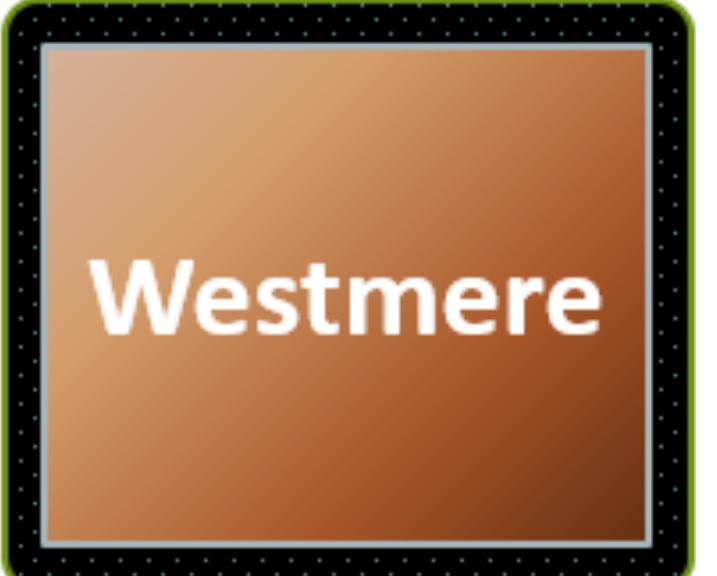
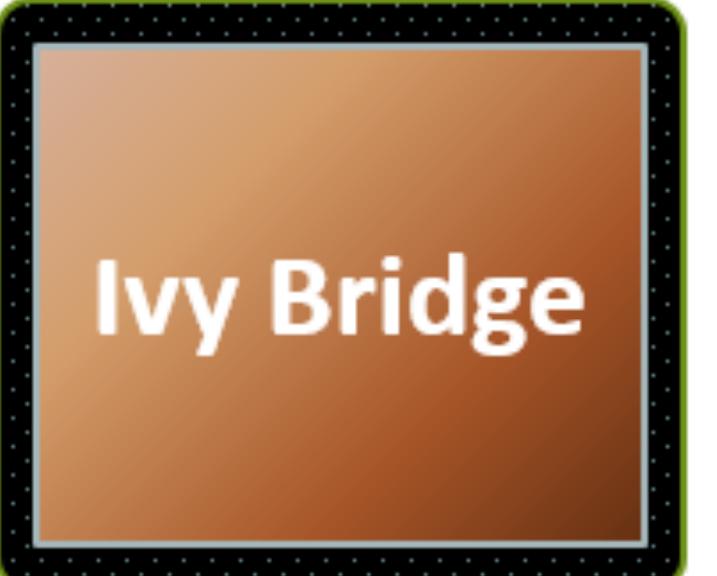
1999	2000	2004	2006	2007	2008	2009	2010\11
SSE	SSE2	SSE3	SSSE3	SSE4.1	SSE4.2	AES-NI	AVX
70 instr Single-Precision Vectors Streaming operations	144 instr Double-precision Vectors 8/16/32 64/128-bit vector integer	13 instr Complex Data	32 instr Decode	47 instr Video Graphics building blocks Advanced vector instr	8 instr String/XML processing POP-Count CRC	7 instr Encryption and Decryption Key Generation	~100 new instr. ~300 legacy sse instr updated 256-bit vector 3 and 4-operand instructions

# Intel Advanced Vector eXtensions

AVX also supported by AMD processors

Comp

on & Weaver

2011	2012	2013	2014	2015	Future
<b>AVX Registers getting wider, instruction set getting richer</b>					
<b>87 GFLOPS</b>  Westmere	<b>185 GFLOPS</b>  Sandy Bridge	<b>~225 GFLOPS</b>  Ivy Bridge	<b>~500 GFLOPS</b>  Haswell	<b>tbd GFLOPS</b>  Broadwell	<b>tbd GFLOPS</b>  Skylake
32 nm SSE 4.2 DDR3 PCIe2	32 nm <b>AVX (256 bit registers)</b> DDR3 PCIe3	22 nm	22 nm <b>AVX2 (new instructions)</b> DDR4 PCIe3	14 nm	14 nm <b>AVX 3.2 (512 bit registers)</b> DDR4 PCIe4

# Note on AVX-512...

- 512 bit vectors may be more efficient
  - Twice as much operations per instruction
  - But also requires 2x the resources over 256b
- Really heavy vector code is now done on the GPU
  - Thousands of simultaneous operations
- Intel is now disabling it on some of the newer designs
  - The big/little "Alder Lake":  
The little cores don't support AVX-512  
The big ones do... but have it disabled
- For us, the Hive is only AVX-2 anyway
  - So no big loss

# Laptop CPU Specs

```
$ sysctl -a | grep cpu
```

hw.physicalcpu: 4

hw.logicalcpu: 8

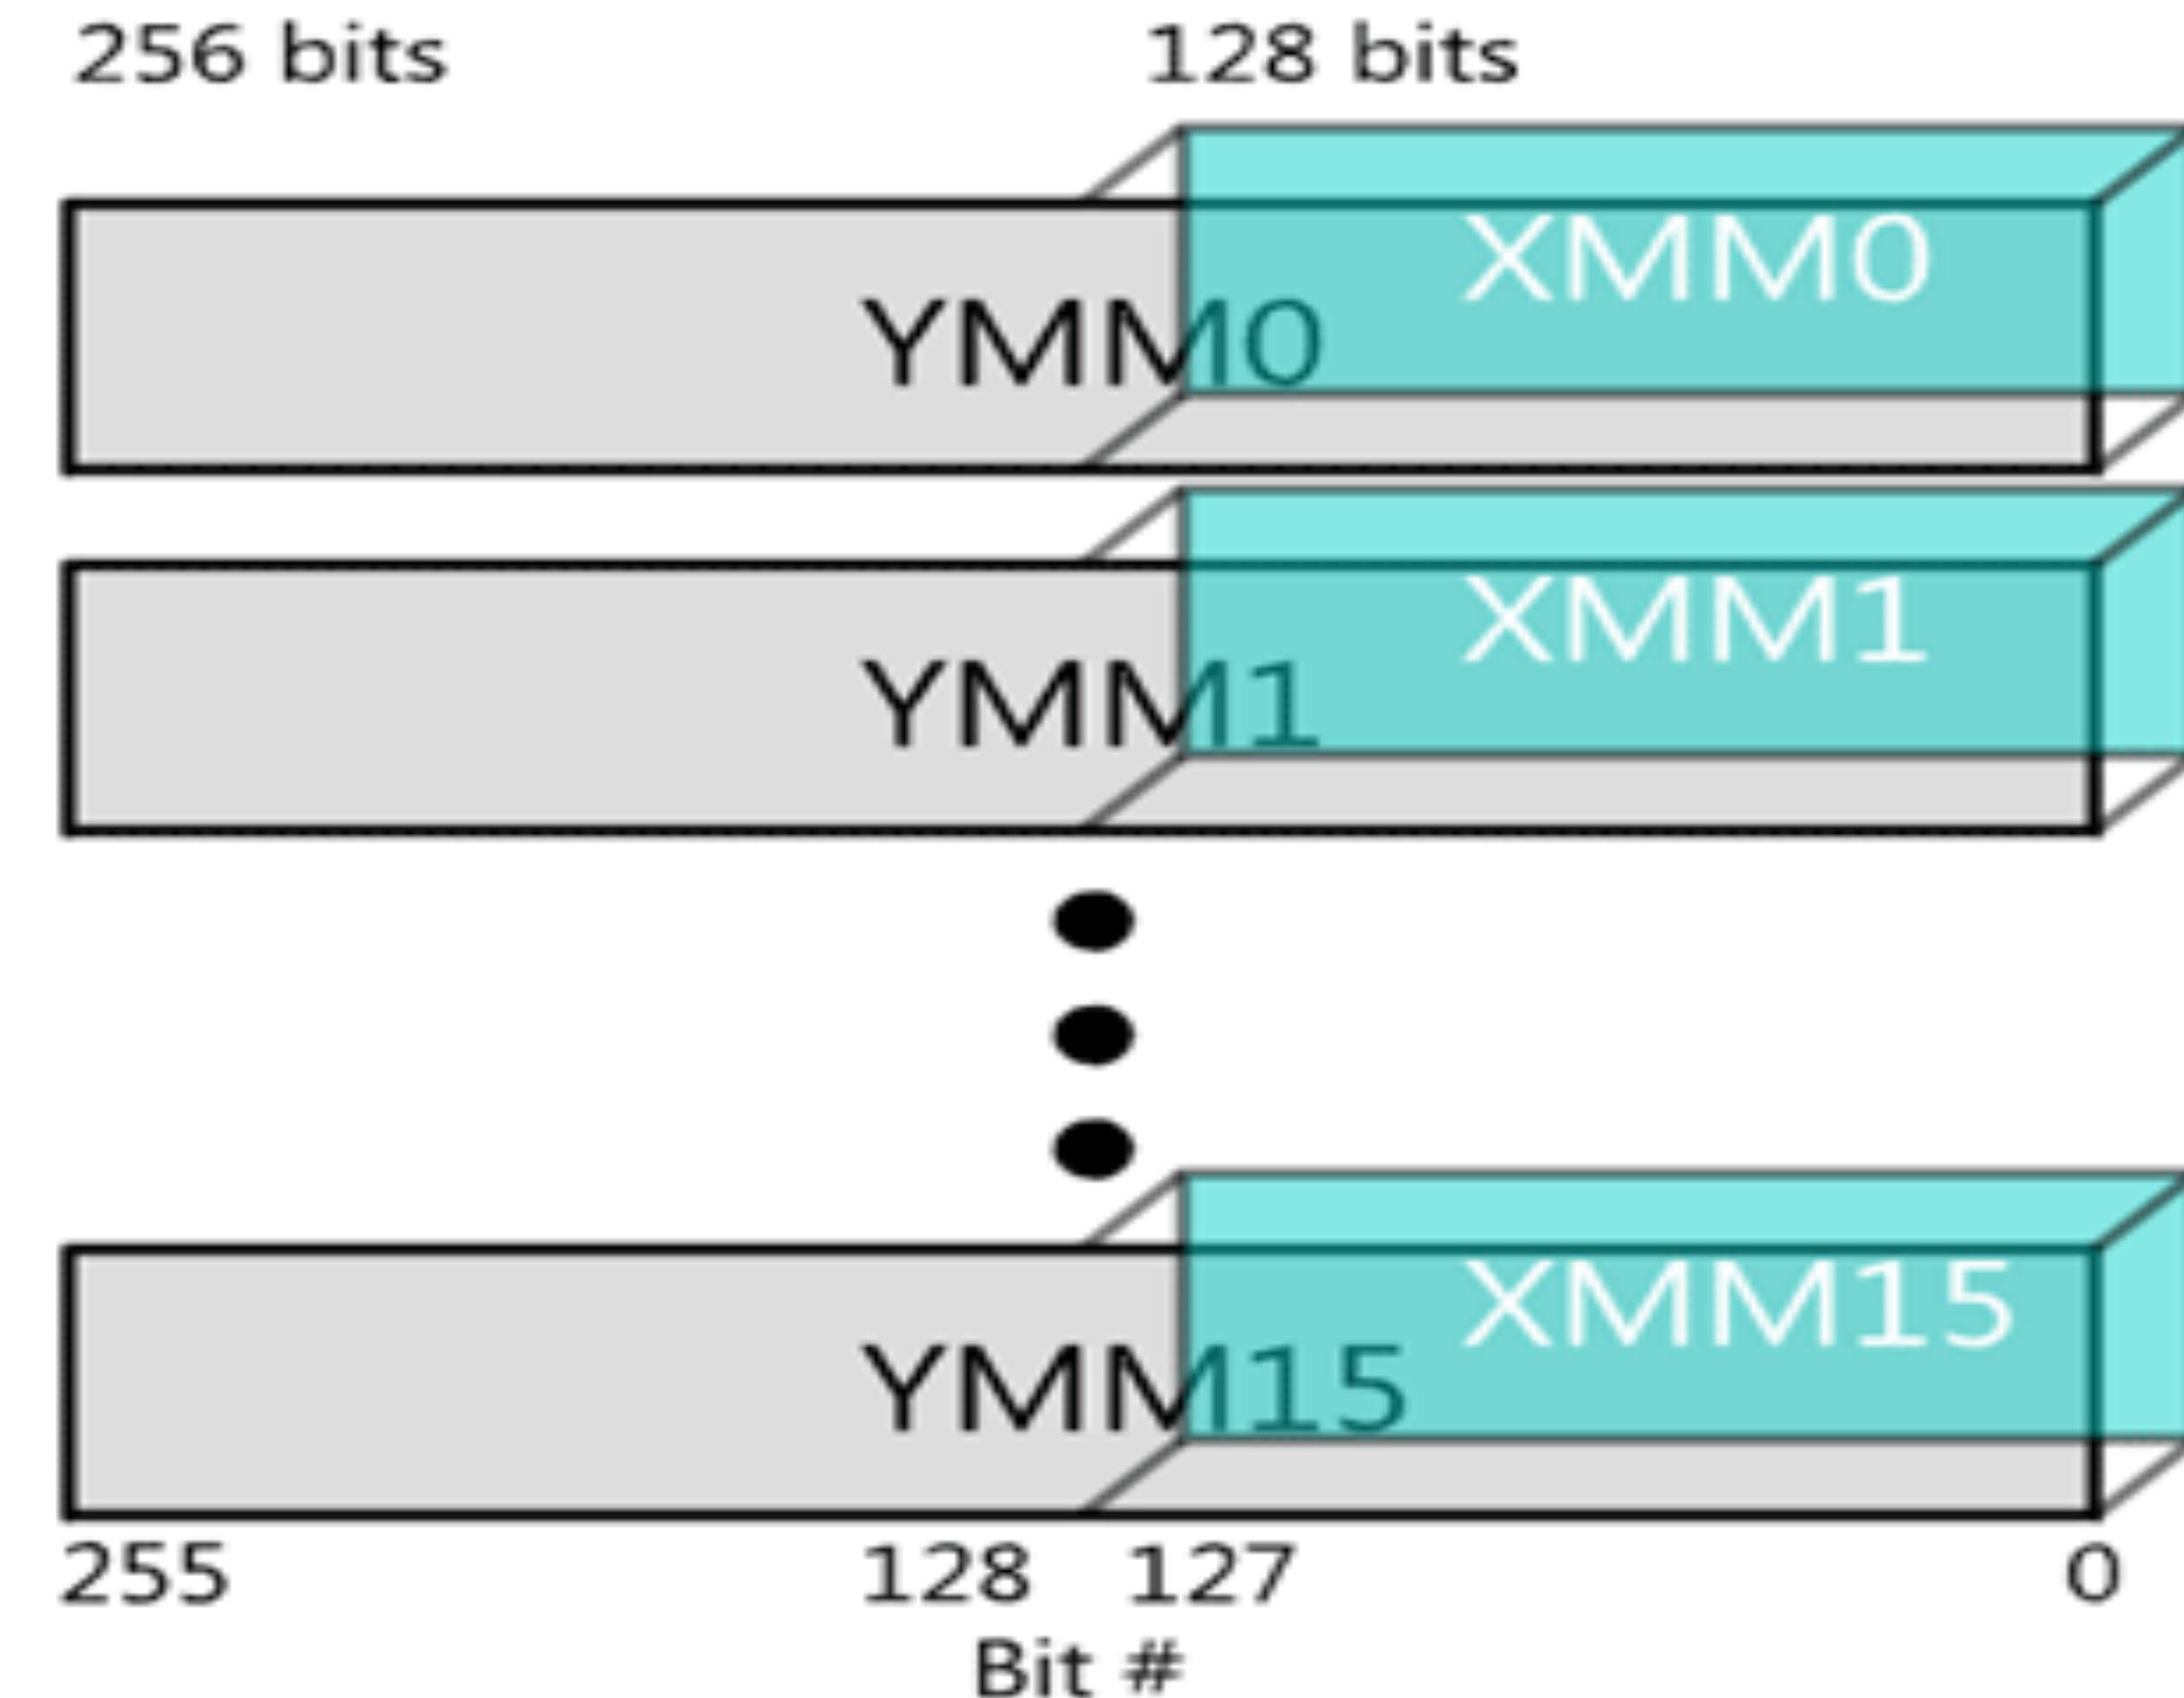
machdep.cpu.brand\_string: Intel(R) Core(TM) i5-1038NG7 CPU @ 2.00GHz

machdep.cpu.features: FPU VME DE PSE TSC MSR PAE MCE CX8 APIC SEP MTRR PGE MCA CMOV PAT PSE36 CLFSH DS ACPI MMX FXSR SSE SSE2 SS HTT TM PBE SSE3 PCLMULQDQ DTES64 MON DSCPL VMX EST TM2 SSSE3 FMA CX16 TPR PDCM SSE4.1 SSE4.2 x2APIC MOVBE POPCNT AES PCID XSAVE OSXSAVE SEGLIM64 TSCTMR AVX1.0 RDRAND F16C

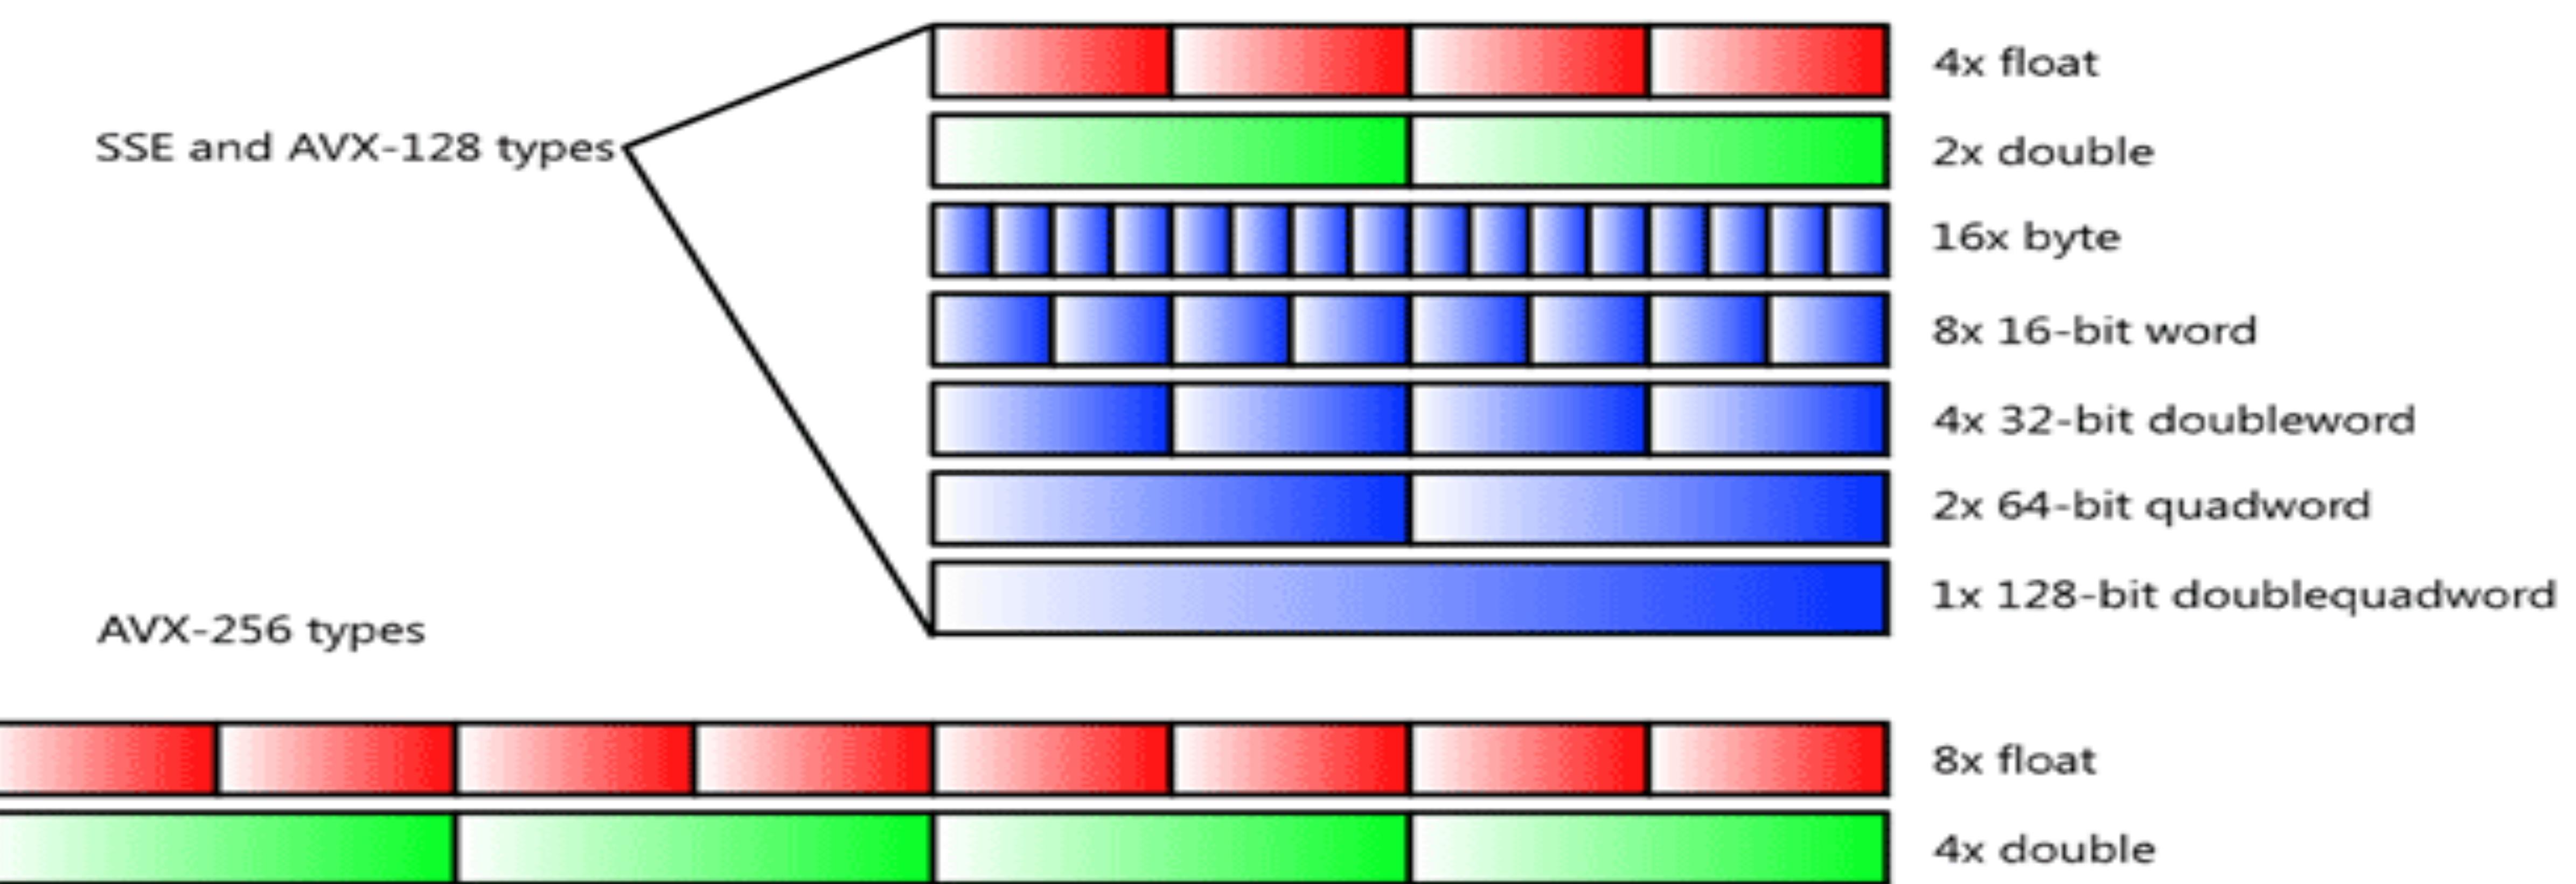
machdep.cpu.leaf7\_features: RDWRFSGS TSC\_THREAD\_OFFSET SGX BMI1 AVX2 FDPEO SMEP BMI2 ERMS INVPCID FPU\_CSDS AVX512F AVX512DQ RDSEED ADX SMAP AVX512IFMA CLFSOPT IPT AVX512CD SHA AVX512BW AVX512VL AVX512VBM1 UMIP PKU GFNI VAES VPCLMULQDQ AVX512VNNI AVX512BITALG AVX512VPOPCNTDQ RDPID SGXLC FSREPMOV MDCLEAR IBRS STIBP L1DF ACAPMSR SSBD

machdep.cpu.extfeatures: SYSCALL XD 1GBPAGE EM64T LAHF LZCNT PREFETCHW RDTSCP TSCI

# AVX SIMD Registers: 16 registers, Greater Bit Extensions Overlap Smaller Versions



# Intel SIMD Data Types



(AVX-512 available (**but not on Hive so you can't use on Proj 4**):

16x float and 8x double)...

But latest: Intel has decided to basically give up on AVX-512 going forward!

Alder Lake's "efficient" cores don't include it so it is turned off!

# Doggo Break!

Computer Science 61C

McMahon & Weaver



# Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- **SIMD matrix multiplication**
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

# Problem

- Today's compilers can generate SIMD code
- But in some cases, better results by hand (assembly)
- We will study x86 (not using RISC-V as no vector hardware widely available yet)
  - Over 1000 instructions to learn ...
  - Or to google, either one...
- Can we use the compiler to generate all non-SIMD instructions?

# x86 SIMD “Intrinsics”



The Intel Intrinsics Guide is an interactive reference tool for Intel intrinsic instructions, which are C style functions that provide access to many Intel instructions - including Intel® SSE, AVX, AVX-512, and more - without the need to write assembly code.

## Technologies

- MMX
  - SSE
  - SSE2
  - SSE3
  - SSSE3
  - SSE4.1
  - SSE4.2
  - AVX
  - AVX2
  - FMA
  - AVX-512
  - KNC
  - SVML
  - Other

## Categories

- Application-Targeted
  - Arithmetic
  - Bit Manipulation
  - Cast
  - Compare
  - Convert
  - Cryptography
  - Elementary Math

## Functions

- ## General Support

## mul\_pd

x

?

```
__m256d _mm256_mul_pd (__m256d a, __m256d b)
```

**vmulpd**

## Synopsis

```
__m256d _mm256_mul_pd (__m256d a, __m256d b)
```

```
#include <immintrin.h>
```

Instruction: `vmulpd ymm, ymm, ymm`

CPUID Flags: AVX

## Description

Multiply packed double-precision (64-bit) floating-point elements in **a** and **b**, and store the results in **dst**.

Operation

```
FOR j := 0 to 3
    i := j*64
    dst[i+63:i] := a[i+63:i] * b[i+63:i]
ENDFOR
dst[MAX:256] := 0
```

## Performance

Architecture	Latency	Throughput (CPI)
Icelake	4	0.5
Skylake	4	0.5
Broadwell	3	0.5
Haswell	5	0.5

Intrinsic

# — assembly instruction

# 4 parallel multiplies

Instructions per clock cycle (CPI = 0.5)

4 cycles latency (data hazard time)

# x86 Intrinsics AVX Data Types

## Intrinsics: Direct access to assembly from C

Type	Meaning
<code>_m256</code>	256-bit as eight single-precision floating-point values, representing a YMM register or memory location
<code>_m256d</code>	256-bit as four double-precision floating-point values, representing a YMM register or memory location
<code>_m256i</code>	256-bit as integers, (bytes, words, etc.)
<code>_m128</code>	128-bit single precision floating-point (32 bits each)
<code>_m128d</code>	128-bit double precision floating-point (64 bits each)

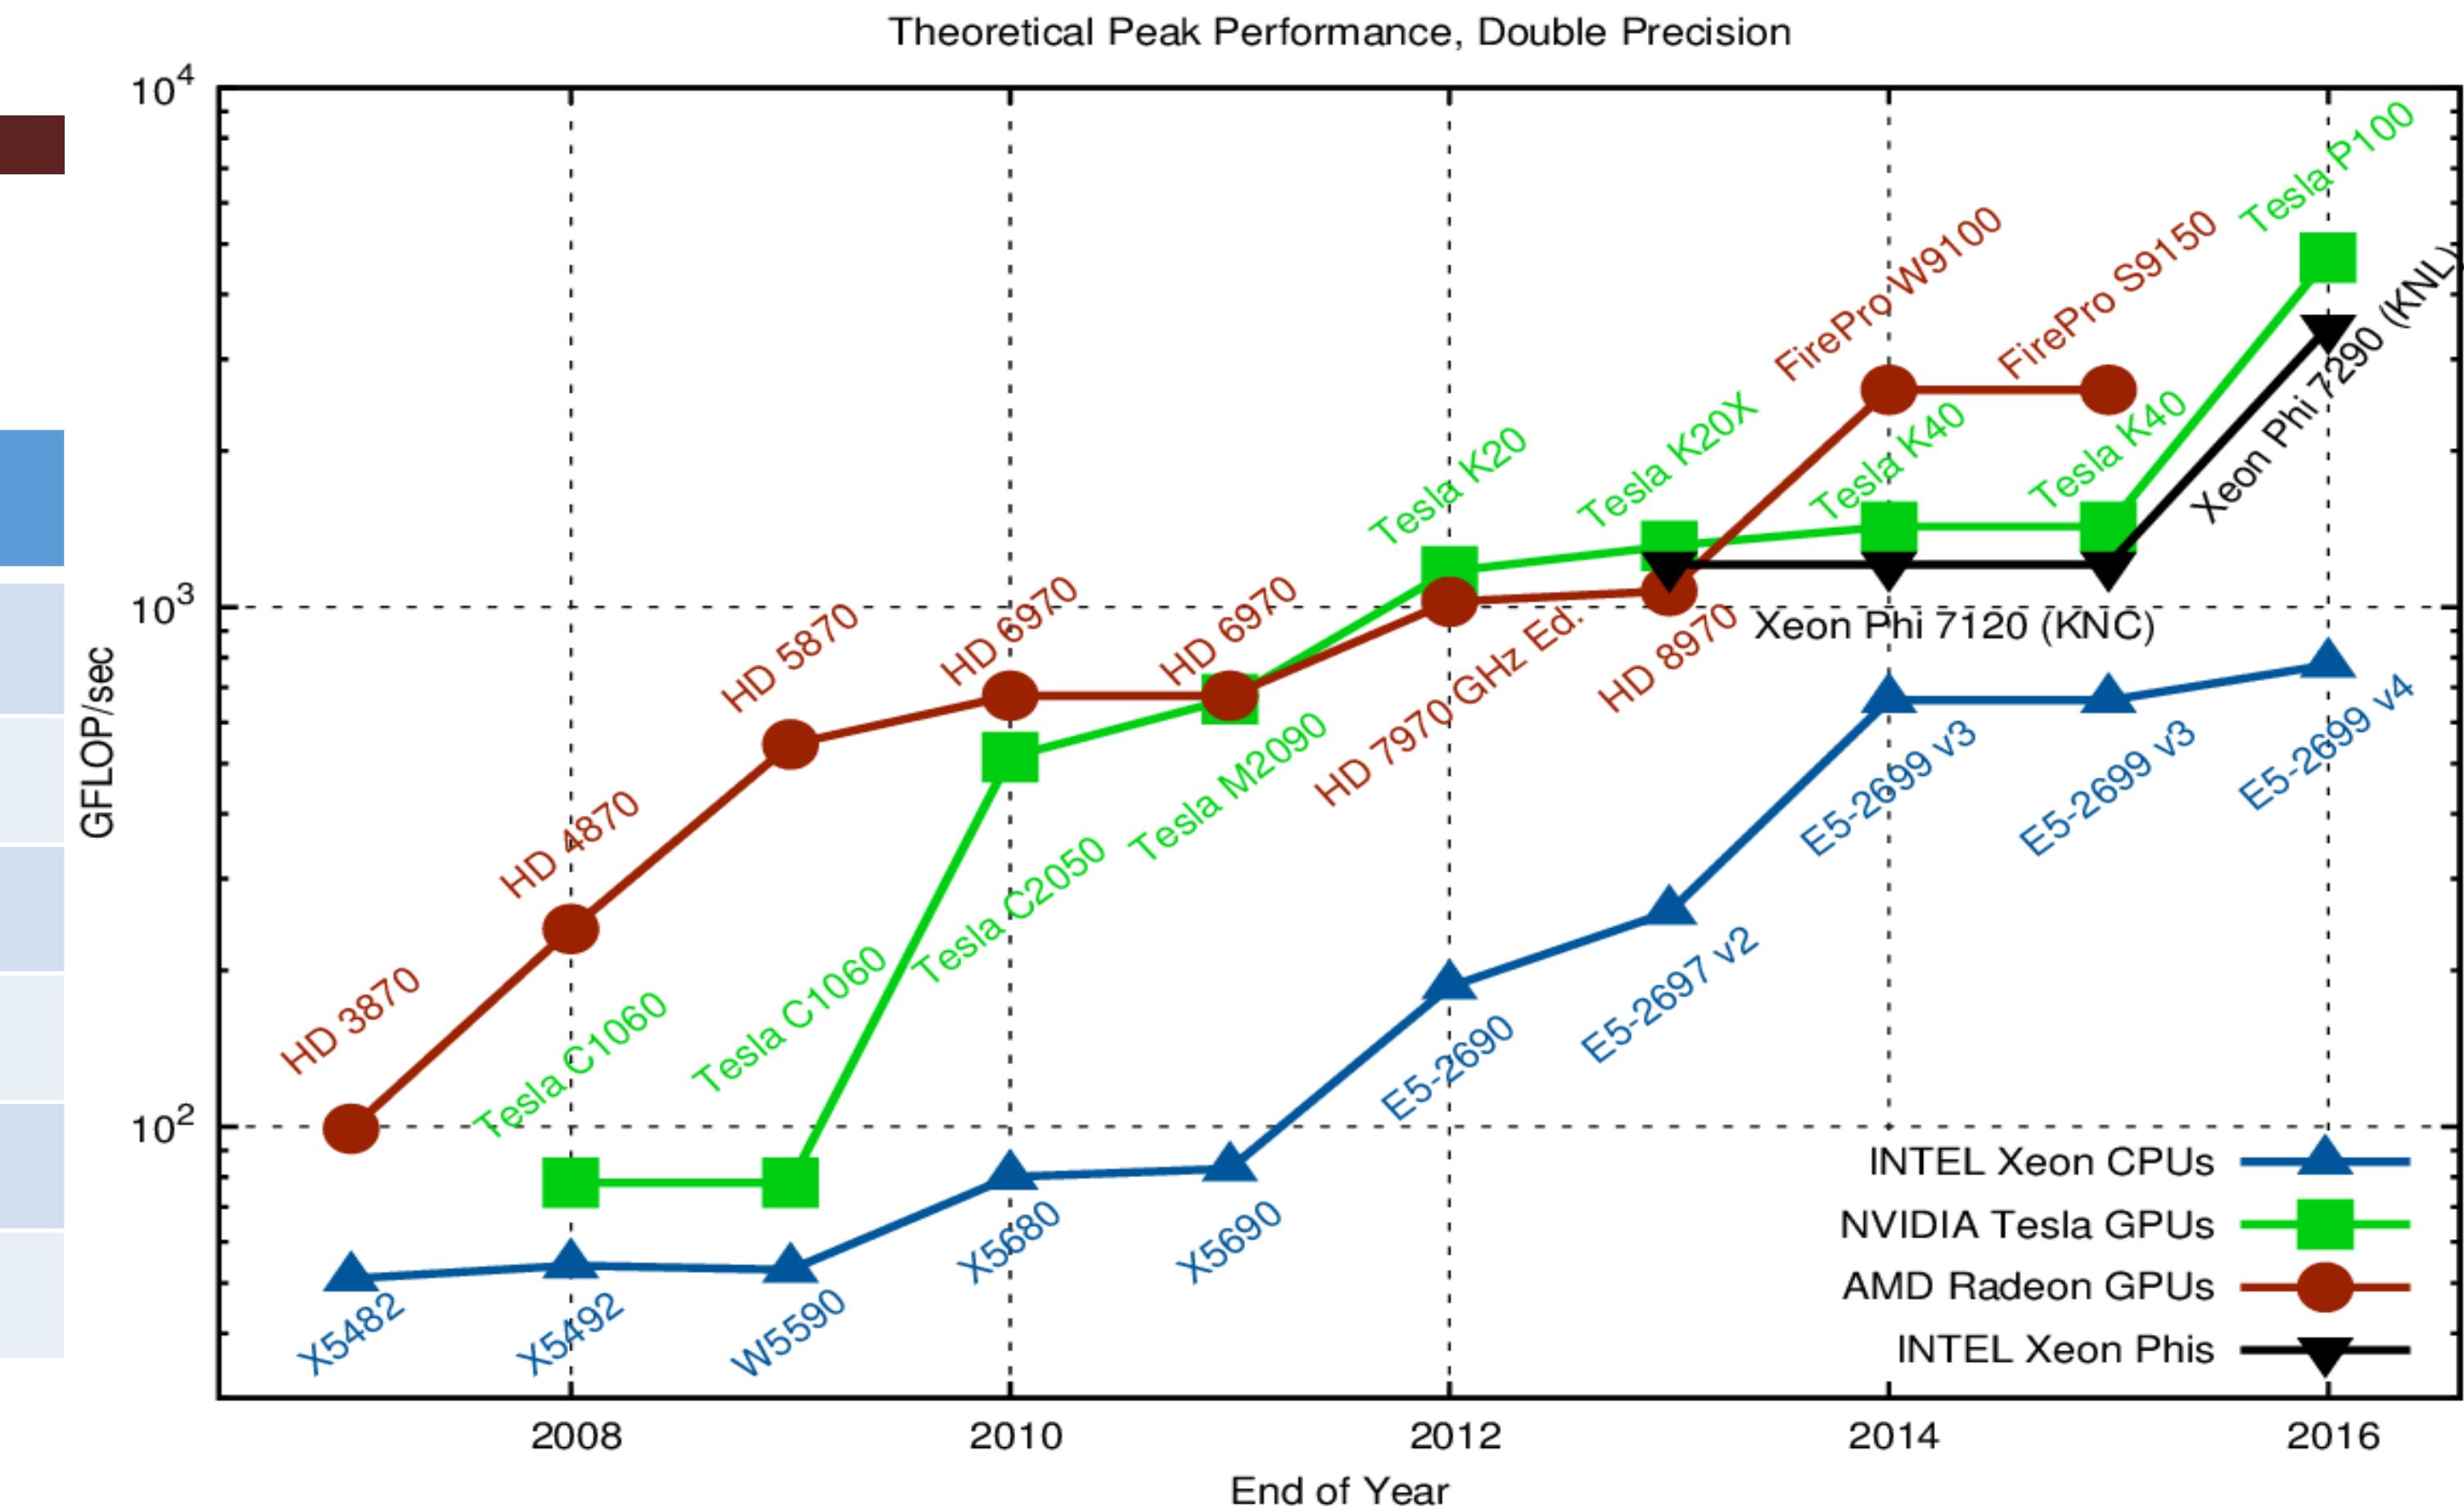
# Intrinsics AVX Code Nomenclature

Marking	Meaning
[s/d]	Single- or double-precision floating point
[i/u] nnn	Signed or unsigned integer of bit size <i>nnn</i> , where <i>nnn</i> is 128, 64, 32, 16, or 8
[ps/pd/sd]	Packed single, packed double, or scalar double
epi32	Extended packed 32-bit signed integer
si256	Scalar 256-bit integer

# Raw Double-Precision Throughput

Computer Science 61C

Characteristic	Value
CPU	i7-5557U
Clock rate (sustained)	3.1 GHz
Instructions per clock (mul_pd)	2
Parallel multiplies per instruction	4
Peak double FLOPS	<b>24.8 GFLOPS</b>



<https://www.karlrupp.net/2013/06/cpu-gpu-and-mic-hardware-characteristics-over-time/>

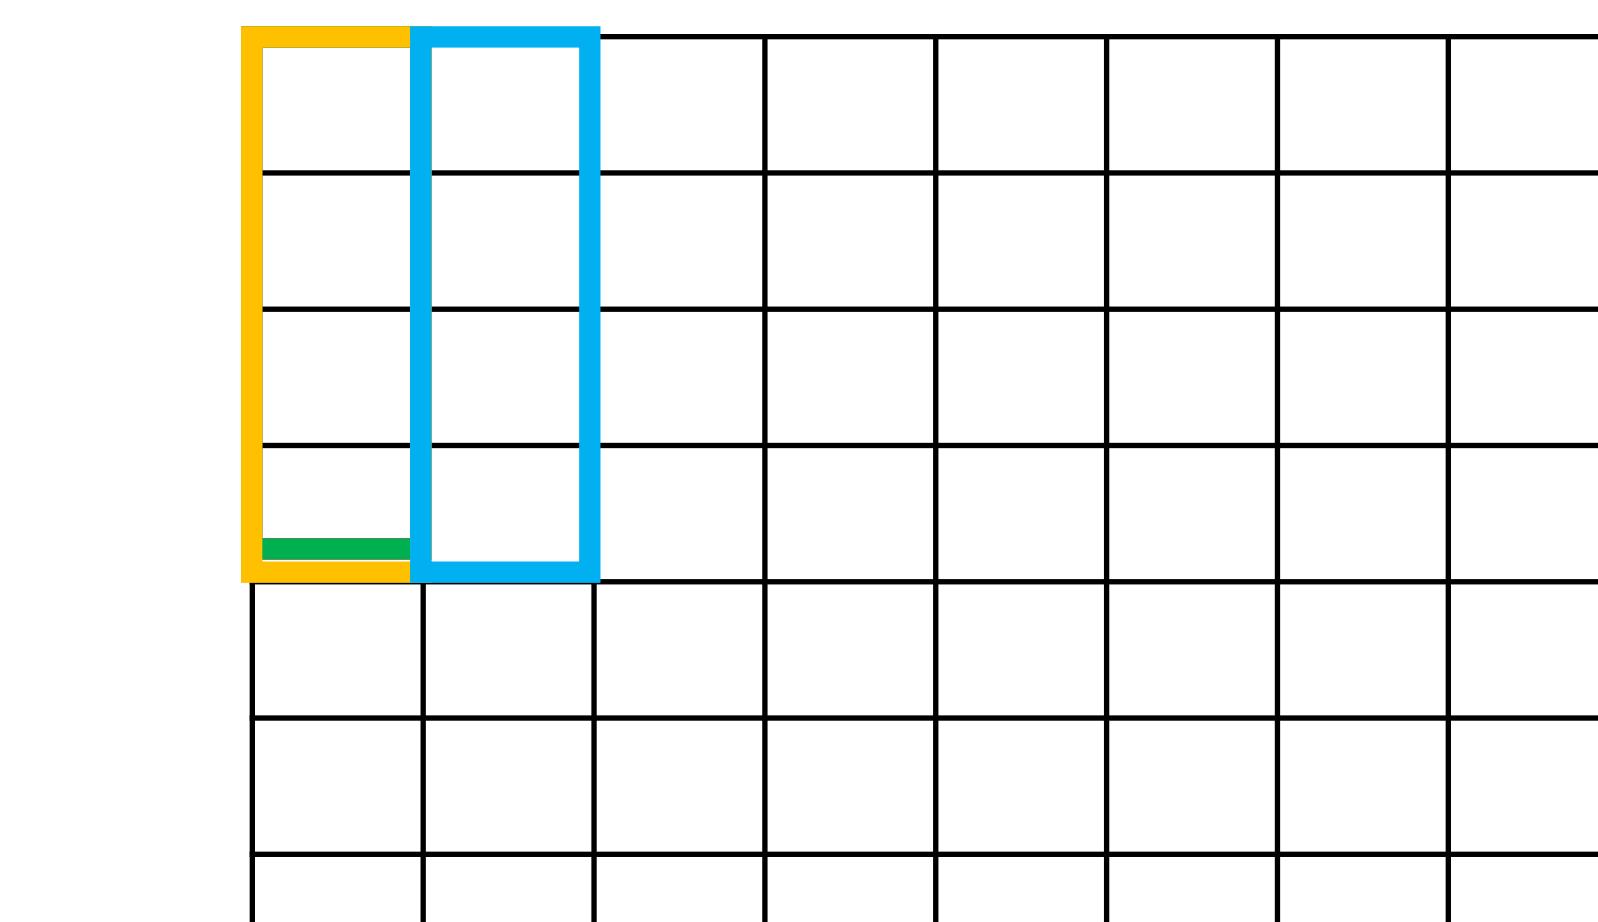
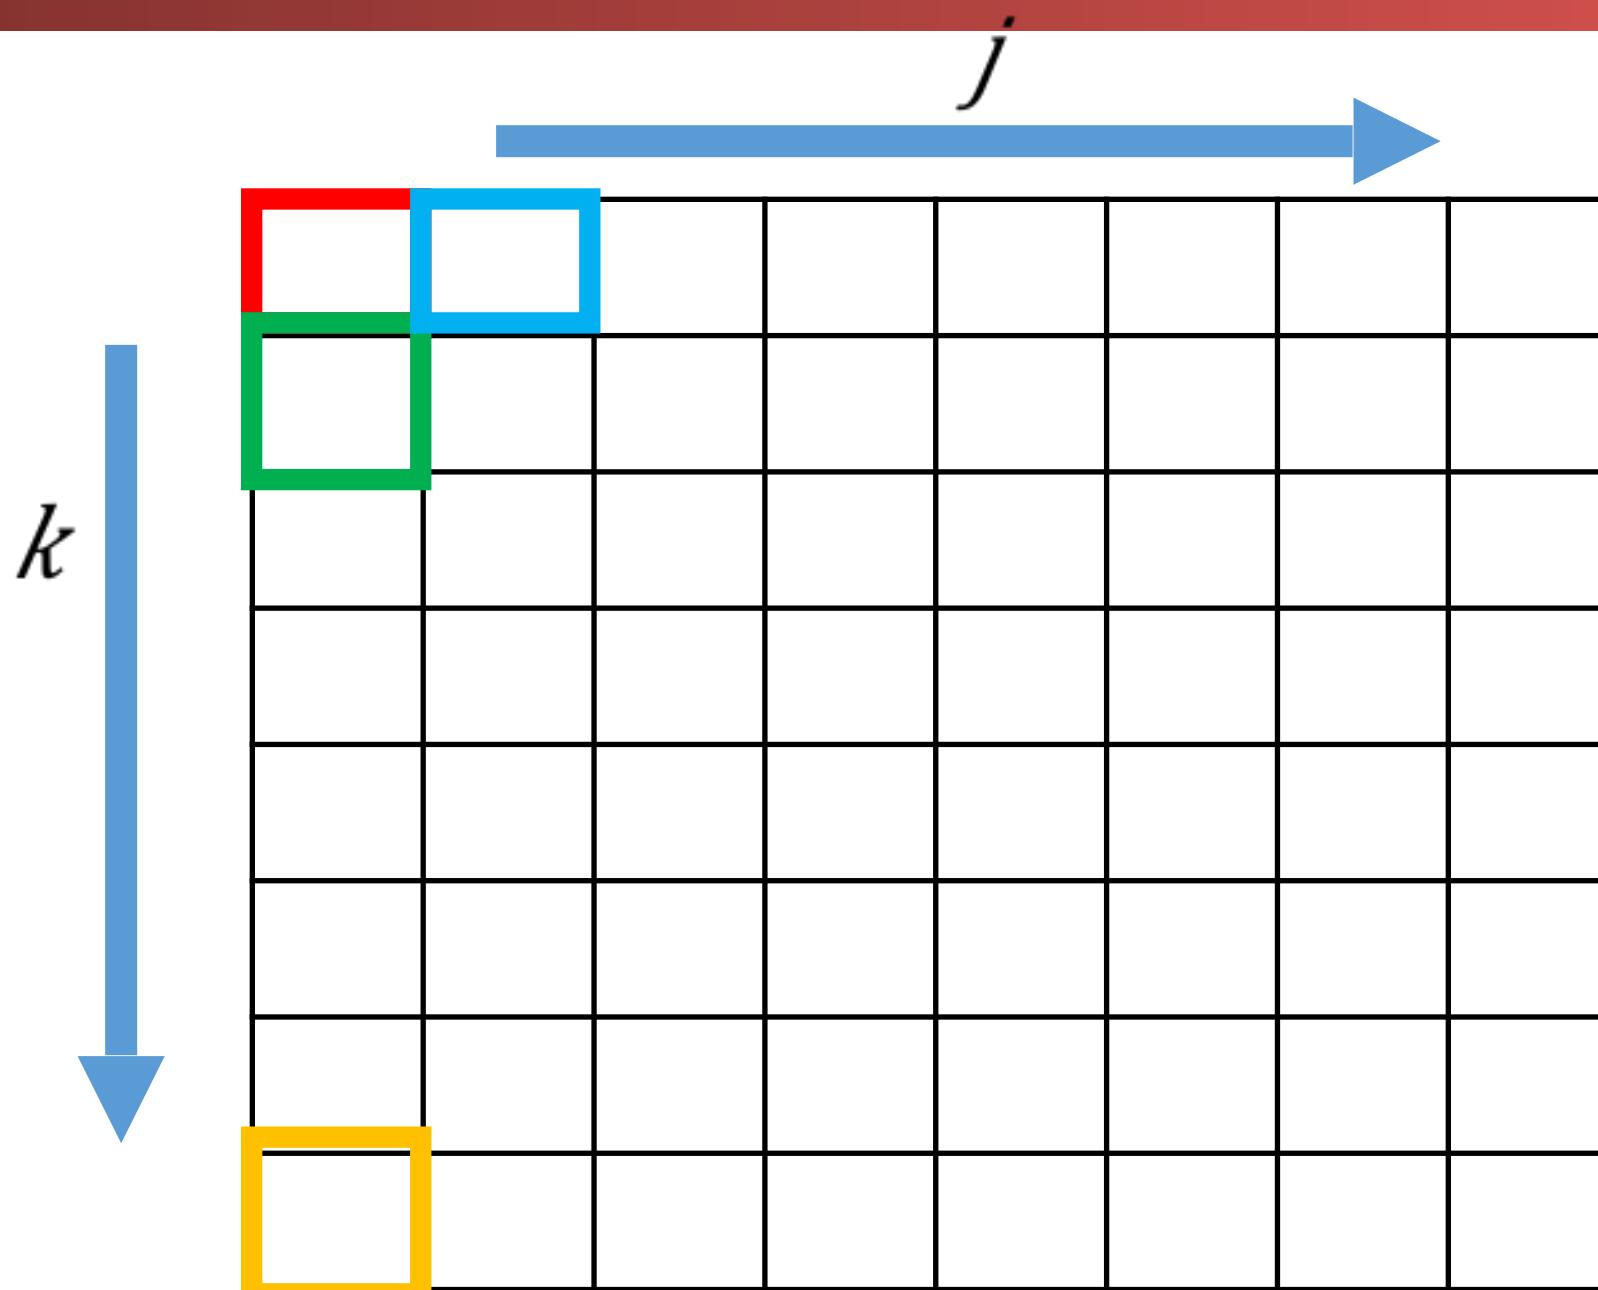
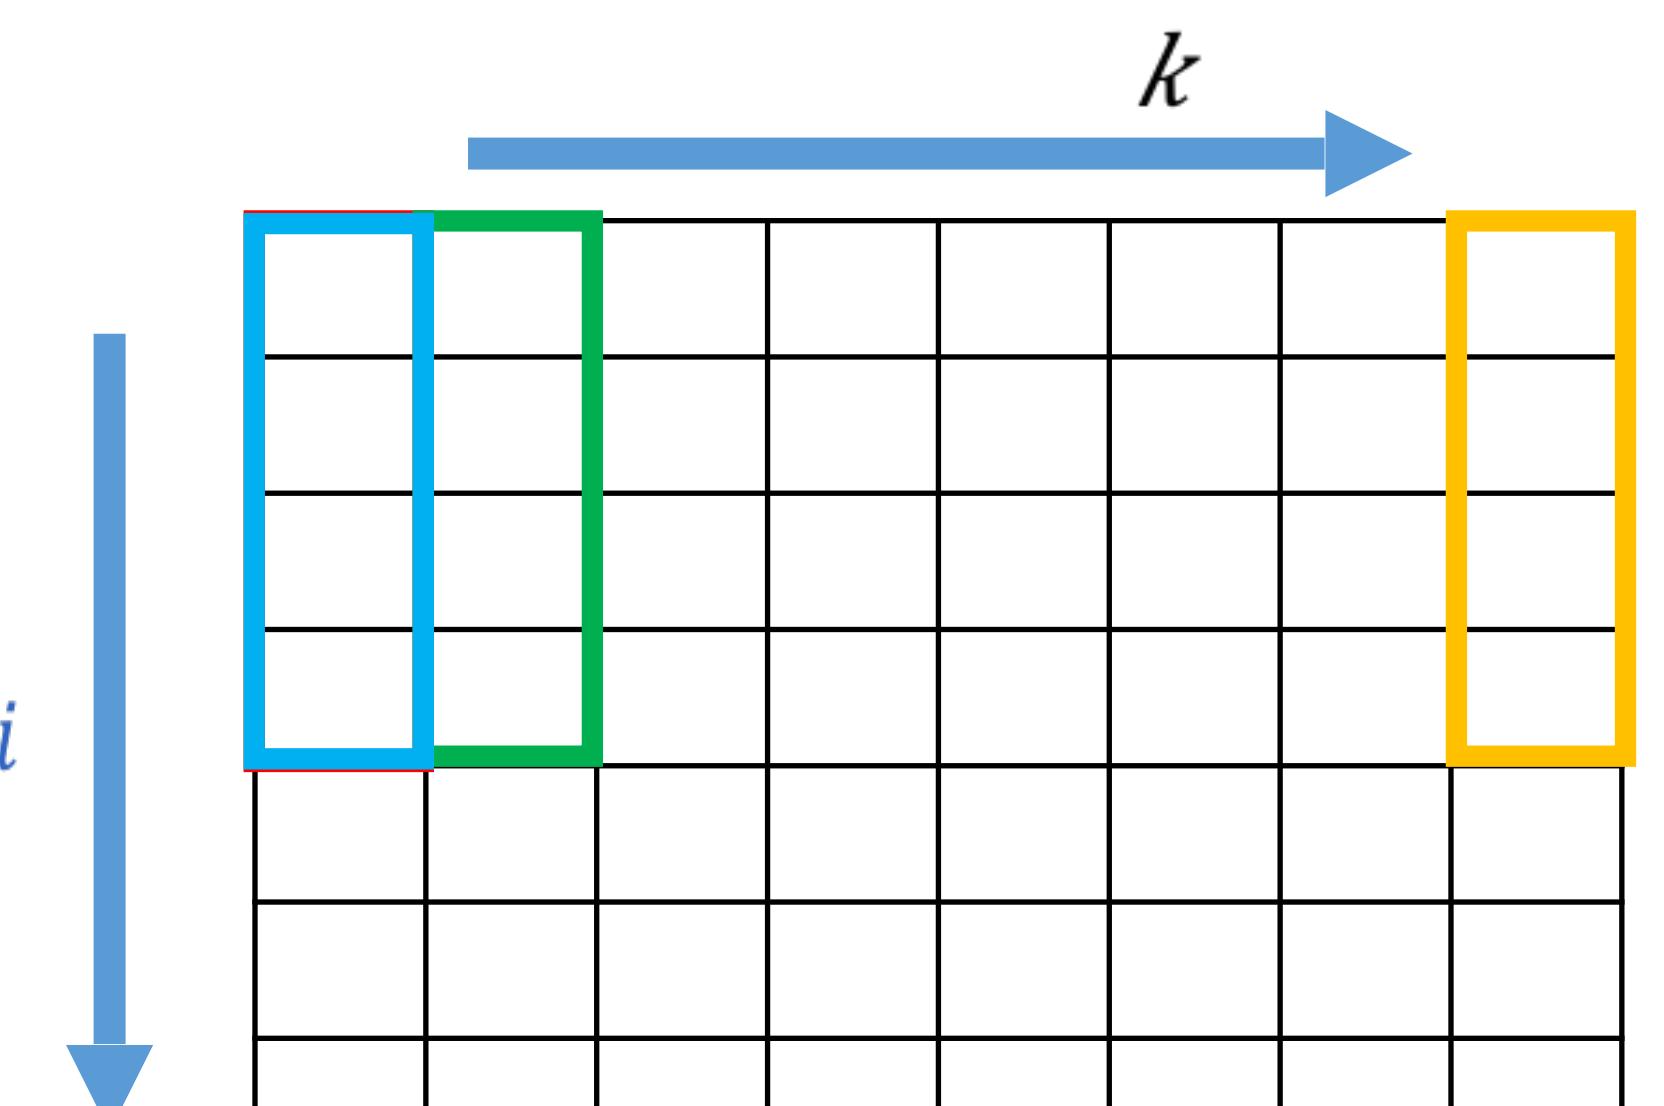
Actual performance is lower because of overhead

# Vectorized Matrix Multiplication

for i ...; **i+=4**

**Inner Loop:** for j ...

```
__m256d c0 = {0,0,0,0};  
for (int k=0; k<N; k++) {  
    c0 = _mm256_fmadd_pd(  
        _mm256_load_pd(a+i+k*N),  
        _mm256_broadcast_sd(b+k+j*N),  
        c0);  
}  
_mm256_store_pd(c+i+j*N, c0);
```



# “Vectorized” dgemm

```
void dgemm_avx(int N, double *a, double *b, double *c) {
    int i,j,k; __m256d c0;
    for(i = 0; i < N; i += 4) {
        for(j = 0; j < N; ++j) {
            c0 = {0,0,0,0}
            for(k = 0; k < N; ++k) {
                c0 = __mm256_add_pd(
                    c0,
                    __mm256_mull_pd(
                        __mm256_load_pd(a+i+k*N) ,
                        __mm256_broadcast_sd(b+k+j*N) ) );
            }
            __mm256_store_pd(c+i+j*N, c0);
        }
    }
}
```

# Performance

N	Gflops	
	scalar	avx
32	1.30	4.56
160	1.30	5.47
480	1.32	5.27
960	0.91	3.64

- 4x faster
- But still << theoretical 25 GFLOPS!

# Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- **Loop unrolling**
- Memory access strategy - blocking
- And in Conclusion, ...

# Loop Unrolling

- On high performance processors, optimizing compilers performs “loop unrolling” operation to expose more parallelism and improve performance:

```
for(i=0; i<N; i++)
    x[i] = x[i] + s;
```

- Could become:

```
for(i=0; i<N; i+=4) {
    x[i] = x[i] + s;
    x[i+1] = x[i+1] + s;
    x[i+2] = x[i+2] + s;
    x[i+3] = x[i+3] + s;
}
```

1. Expose data-level parallelism for vector (SIMD) instructions or super-scalar multiple instruction issue
2. Mix pipeline with unrelated operations to help with reduce hazards
3. Reduce loop “overhead”
4. Makes code size larger

# Amdahl's Law\* applied to `dgemm`

- Measured `dgemm` performance
  - Peak 5.5 GFLOPS
  - Large matrices 3.6 GFLOPS
  - Processor 24.8 GFLOPS
- Why are we not getting (close to) 25 GFLOPS?
  - Something else (not floating-point ALU) is limiting performance!
  - But what? Possible culprits:
    - Cache
    - Hazards
    - Let's look at both!

# “Vectorized” dgemm: Pipeline Hazards

```
void dgemm_avx(int N, double *a, double *b, double *c) {
    int i,j,k; __m256d c0;
    for(i = 0; i < N; i += 4) {
        for(j = 0; j < N; ++j) {
            c0 = {0,0,0,0}
            for(k = 0; k < N; ++k) {
                c0 = __mm256_add_pd(←
                    c0,
                    __mm256_mull_pd(
                        __mm256_load_pd(a+i+k*N),
                        __mm256_broadcast_sd(b+k+j*N)) );
            }
            __mm256_store_pd(c+i+j*N, c0);
        }
    }
}
```

“add\_pd” depends on result of “mult\_pd” which depends on “load\_pd”

# Loop Unrolling

// Loop unrolling; P&H p. 352

const int UNROLL = 4;

eaver

```
void dgemm_unroll(int n, double *A, double *B, double *C) {
    for (int i=0; i<n; i+= UNROLL*4) {
        for (int j=0; j<n; j++) {
            __m256d c[4]; ← 4 registers
            for (int x=0; x<UNROLL; x++)
                c[x] = _mm256_load_pd(C+i+x*4+j*n);
            for (int k=0; k<n; k++) {
                __m256d b = _mm256_broadcast_sd(B+k+j*n);
                for (int x=0; x<UNROLL; x++) ← Compiler does the unrolling
                    c[x] = _mm256_add_pd(c[x],
                        _mm256_mul_pd(_mm256_load_pd(A+n*k+x*4+i), b));
            }
            for (int x=0; x<UNROLL; x++)
                _mm256_store_pd(C+i+x*4+j*n, c[x]);
        }
    }
}
```

How do you verify that the generated code is actually unrolled?

# Performance

N	Gflops		
	scalar	avx	unroll
32	1.30	4.56	12.95
160	1.30	5.47	19.70
480	1.32	5.27	14.50
960	0.91	3.64	6.91

WOW!

?

# Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- **Memory access strategy - blocking**
- And in Conclusion, ...

# FPU versus Memory Access

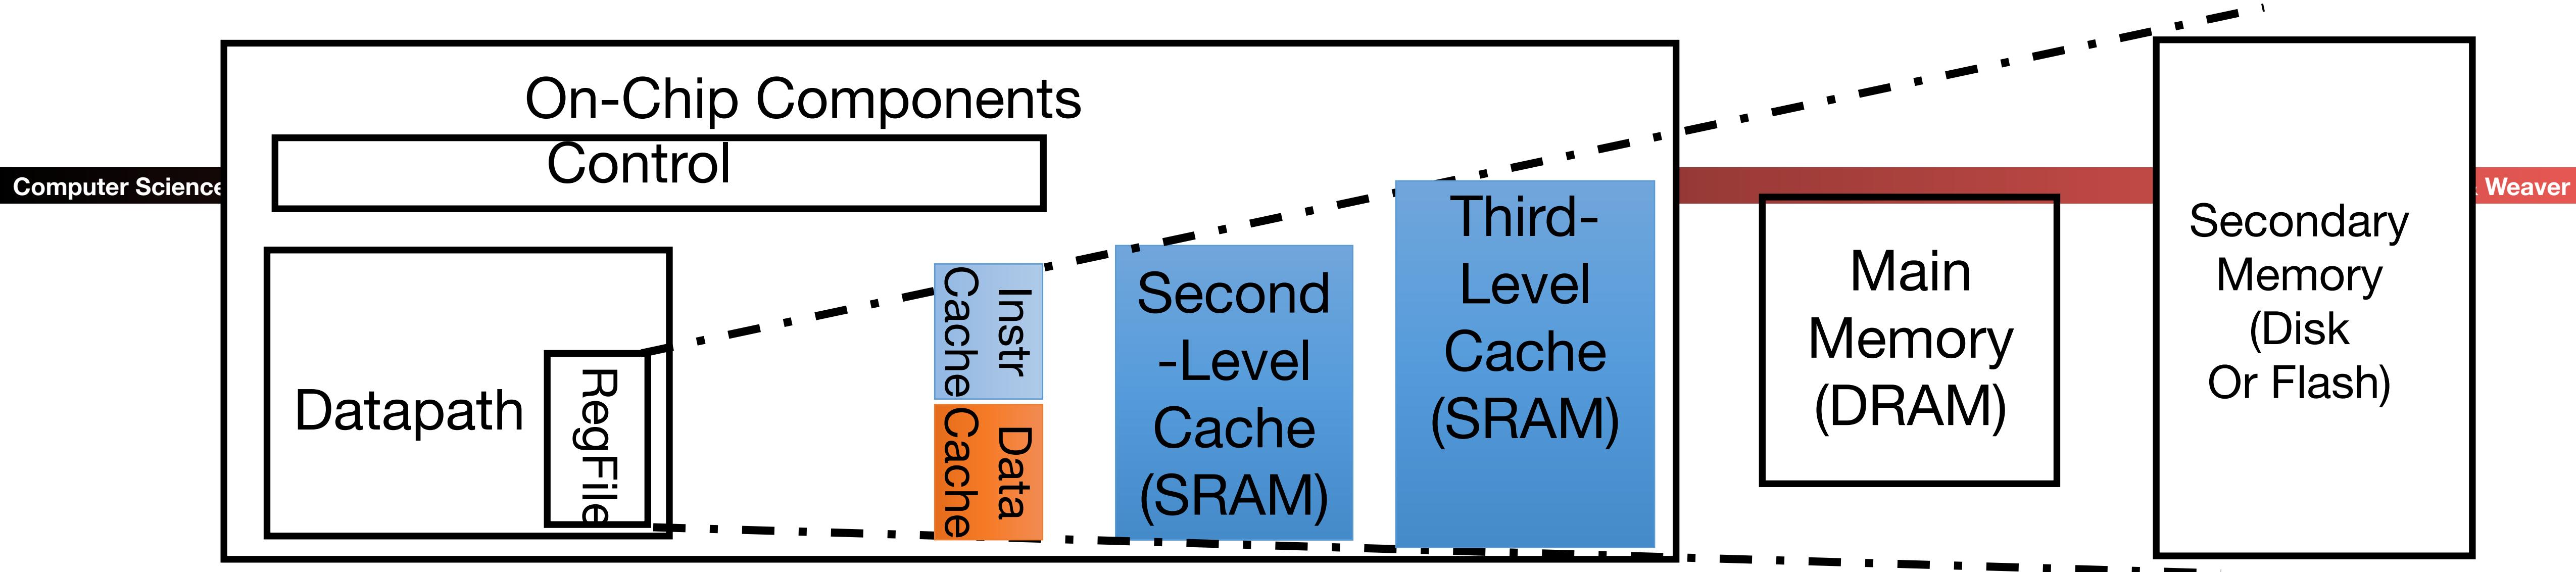
- How many floating-point operations does matrix multiply take?
  - $F = 2 \times N^3$  ( $N^3$  multiplies,  $N^3$  adds) in the straightforward case
- How many memory load/stores?
  - $M = 3 \times N^2$  (for A, B, C)
- Many more floating-point operations than memory accesses
  - $q = F/M = 2/3 * N$
  - Good, since arithmetic is faster than memory access
  - Let's check the code ...

# But memory is accessed repeatedly

- $q = F/M = 1.6!$  (1.25 loads and 2 floating-point operations)

## Inner loop:

```
for (int k=0; k<N; k++) {
    c0 = _mm256_add_pd(
        c0,           // c0 += a[i][k] * b[k][j]
        _mm256_mul_pd(
            _mm256_load_pd(a+i+k*N),
            _mm256_broadcast_sd(b+k+j*N)) );
}
```



<b>Speed (cycles):</b>	$\frac{1}{2}$ 's	1's	10's	100's-1000	1,000,000's
<b>Size (bytes):</b>	100's	10K's	M's	G's	T's
<b>Cost/bit:</b>	highest				lowest

- Where are the operands (A, B, C) stored?
- What happens as N increases?
- **Idea:** arrange that most accesses are to fast cache!

# Blocking

- Idea:
  - Rearrange code to use values loaded in cache many times
  - Only “few” accesses to slow main memory (DRAM) per floating point operation
    - -> throughput limited by FP hardware and cache, not slow DRAM
  - P&H, RISC-V edition p. 465

# Blocking Matrix Multiply

## (divide and conquer: sub-matrix multiplication)

Computer Science 61C

McMahon & Weaver

$$\begin{array}{c|c} \text{Red 3x3 block} \\ \hline \text{Large white matrix} \end{array} \times \begin{array}{c|c} \text{Blue 3x3 block} \\ \hline \text{Large white matrix} \end{array} = \begin{array}{c|c} \text{Purple 3x3 block} \\ \hline \text{Large white matrix} \end{array}$$

$$\begin{array}{c|c} \text{Red 3x3 block} \\ \hline \text{Large white matrix} \end{array} \times \begin{array}{c|c} \text{Blue 3x3 block} \\ \hline \text{Large white matrix} \end{array} = \begin{array}{c|c} \text{Purple 3x3 block} \\ \hline \text{Large white matrix} \end{array}$$

$$\begin{array}{c|c} \text{Red 3x3 block} \\ \hline \text{Large white matrix} \end{array} \times \begin{array}{c|c} \text{Blue 3x3 block} \\ \hline \text{Large white matrix} \end{array} = \begin{array}{c|c} \text{Purple 3x3 block} \\ \hline \text{Large white matrix} \end{array}$$

$$\begin{array}{c|c} \text{Red 3x3 block} \\ \hline \text{Large white matrix} \end{array} \times \begin{array}{c|c} \text{Blue 3x3 block} \\ \hline \text{Large white matrix} \end{array} = \begin{array}{c|c} \text{Purple 3x3 block} \\ \hline \text{Large white matrix} \end{array}$$

$$\begin{array}{c|c} \text{Red 3x3 block} \\ \hline \text{Large white matrix} \end{array} \times \begin{array}{c|c} \text{Blue 3x3 block} \\ \hline \text{Large white matrix} \end{array} = \begin{array}{c|c} \text{Purple 3x3 block} \\ \hline \text{Large white matrix} \end{array}$$

$$\begin{array}{c|c} \text{Red 3x3 block} \\ \hline \text{Large white matrix} \end{array} \times \begin{array}{c|c} \text{Blue 3x3 block} \\ \hline \text{Large white matrix} \end{array} = \begin{array}{c|c} \text{Purple 3x3 block} \\ \hline \text{Large white matrix} \end{array}$$

# Memory Access Blocking

```
// Cache blocking; P&H p. 556
const int BLOCKSIZE = 32;

void do_block(int n, int si, int sj, int sk, double *A, double *B, double *C) {
    for (int i=si; i<si+BLOCKSIZE; i+=UNROLL*4)
        for (int j=sj; j<sj+BLOCKSIZE; j++) {
            __m256d c[4];
            for (int x=0; x<UNROLL; x++)
                c[x] = _mm256_load_pd(C+i+x*4+j*n);
            for (int k=sk; k<sk+BLOCKSIZE; k++) {
                __m256d b = _mm256_broadcast_sd(B+k+j*n);
                for (int x=0; x<UNROLL; x++)
                    c[x] = _mm256_add_pd(c[x],
                                          _mm256_mul_pd(_mm256_load_pd(A+n*k+x*4+i), b));
            }
            for (int x=0; x<UNROLL; x++)
                _mm256_store_pd(C+i+x*4+j*n, c[x]);
        }
}

void dgemm_block(int n, double* A, double* B, double* C) {
    for(int sj=0; sj<n; sj+=BLOCKSIZE)
        for(int si=0; si<n; si+=BLOCKSIZE)
            for (int sk=0; sk<n; sk += BLOCKSIZE)
                do_block(n, si, sj, sk, A, B, C);
}
```

# Performance

N	Gflops			
	scalar	avx	unroll	blocking
32	1.30	4.56	12.95	13.80
160	1.30	5.47	19.70	21.79
480	1.32	5.27	14.50	20.17
960	0.91	3.64	6.91	15.82

# And in Conclusion, ...

- Approaches to Parallelism
  - SISD, SIMD, MIMD (next lecture)
- SIMD
  - One instruction operates on multiple operands simultaneously
- Example: matrix multiplication
  - Floating point heavy -> exploit Moore's law to make fast