




Lab 7

61C Summer 2023

SIMD

- We will be focusing on SIMD today (single instruction multiple streams of data)



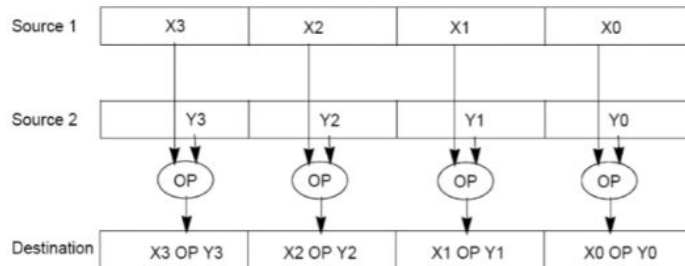


SIMD Intrinsics

- An intrinsic function is a function whose implementation is handled by the compiler
- We will be using SIMD intrinsic functions to speed up our programs

SIMD Intrinsics

- Our SISD (single instruction single data stream) instructions operate on 32 bits
 - Ex. Element wise adding two arrays of size 4 takes 4 instructions $\{1,2,3,4\} + \{1,2,3,4\} = \{2,4,6,8\}$
- SIMD functions use large (128 bits) registers to store and operate on more value at the same time
 - Now we can stuff 4 ints into 1 register and it will only take us 1 instruction to add 4 ints to 4 ints





Intel Intrinsic Functions

- Guide:
<https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>
- `__m128i _mm_setzero_si128()` - returns a 128-bit zero vector
- `__m128i _mm_loadu_si128(__m128i *p)` - returns 128-bit vector stored at pointer p
- `__m128i _mm_add_epi32(__m128i a, __m128i b)` - returns vector $(a_0 + b_0, a_1 + b_1, a_2 + b_2, a_3 + b_3)$



Intel Intrinsics

- `void _mm_storeu_si128(__m128i *p, __m128i a)` - stores 128-bit vector `a` into pointer `p`
- `__m128i _mm_cmpgt_epi32(__m128i a, __m128i b)` - returns the vector $(a_i > b_i ? 0xffffffff : 0x0 \text{ for } i \text{ from } 0 \text{ to } 3)$ (useful as mask for and)
- `__m128i _mm_and_si128(__m128i a, __m128i b)` - returns vector $(a_0 \& b_0, a_1 \& b_1, a_2 \& b_2, a_3 \& b_3)$, where `&` represents the bitwise and operator

Intel Intrinsics Examples

- We have two 4 int wide arrays, arr1, arr2, add arr1 and arr2 element wise and store it in arr1

- `__m128i vec1 = _mm_loadu_si128((__m128i*) arr1)`
- `__m128i vec2 = _mm_loadu_si128((__m128i*) arr2)`
- `__m128i result = _mm_add_epi32(vec1, vec2)`
- `_mm_storeu_si128((__m128i*) arr1, result)`

We update arr1

We have to load arrays from memory into vector registers



Another Intrinsics Example

- Given `int arr[8] = {3, 1, 4, 1, 5, 9, 2, 6};` we want to write SIMD code that adds all the elements up together

Another Intrinsics Example

- Given `int arr[8] = {3, 1, 4, 1, 5, 9, 2, 6};` we want to write SIMD code that adds all the elements up together

```
// Initialize sum vector to {0, 0, 0, 0}  
__m128i sum_vec = _mm_setzero_si128();
```

We first create `sum_vec`(4 ints wide set to all 0s) to store our sum

sum_vec:

0	0	0	0
---	---	---	---

Another Intrinsics Example

- Given `int arr[8] = {3, 1, 4, 1, 5, 9, 2, 6};` we want to write SIMD code that adds all the elements up together

```
// Initialize sum vector to {0, 0, 0, 0}
__m128i sum_vec = _mm_setzero_si128();

// Load array elements 0-3 into a temporary vector register
__m128i tmp = _mm_loadu_si128((__m128i *) arr);
```

Next we load in 4 ints (elems 0-3) to a temp vector

sum_vec:

0	0	0	0
---	---	---	---

tmp:

3	1	4	1
---	---	---	---

Another Intrinsics Example

- Given `int arr[8] = {3, 1, 4, 1, 5, 9, 2, 6};` we want to write SIMD code that adds all the elements up together

```
// Initialize sum vector to {0, 0, 0, 0}
__m128i sum_vec = _mm_setzero_si128();

// Load array elements 0-3 into a temporary vector register
__m128i tmp = _mm_loadu_si128((__m128i *) arr);
// Add to existing sum vector
sum_vec = _mm_add_epi32(sum_vec, tmp);
// sum_vec = {3, 1, 4, 1}
```

sum_vec:

3	1	4	1
---	---	---	---

tmp:

3	1	4	1
---	---	---	---

We add sum_vec and tmp together



Another Intrinsics Example

- Given `int arr[8] = {3, 1, 4, 1, 5, 9, 2, 6};` that adds all the elements up together

we want to write SIMD code

```
// Initialize sum vector to {0, 0, 0, 0}
__m128i sum_vec = _mm_setzero_si128();

// Load array elements 0-3 into a temporary vector register
__m128i tmp = _mm_loadu_si128((__m128i *) arr);
// Add to existing sum vector
sum_vec = _mm_add_epi32(sum_vec, tmp);
// sum_vec = {3, 1, 4, 1}

// Load array elements 4-7 into a temporary vector register
tmp = _mm_loadu_si128((__m128i *) (arr + 4));
```

sum_vec:

3	1	4	1
---	---	---	---

tmp:

5	9	2	6
---	---	---	---

We load in the next 4 elems of arr into tmp



Another Intrinsics Example

- Given `int arr[8] = {3, 1, 4, 1, 5, 9, 2, 6};` that adds all the elements up together

we want to write SIMD code

```
// Initialize sum vector to {0, 0, 0, 0}
__m128i sum_vec = _mm_setzero_si128();

// Load array elements 0-3 into a temporary vector register
__m128i tmp = _mm_loadu_si128((__m128i *) arr);
// Add to existing sum vector
sum_vec = _mm_add_epi32(sum_vec, tmp);
// sum_vec = {3, 1, 4, 1}

// Load array elements 4-7 into a temporary vector register
tmp = _mm_loadu_si128((__m128i *) (arr + 4));
// Add to existing sum vector
sum_vec = _mm_add_epi32(sum_vec, tmp);
// sum_vec = {3 + 5, 1 + 9, 4 + 2, 1 + 6}
```

sum_vec:

8	10	6	7
---	----	---	---

tmp:

5	9	2	6
---	---	---	---

Once again we add sum_vec and tmp



Another Intrinsics Example

- Given `int arr[8] = {3, 1, 4, 1, 5, 9, 2, 6};` we want to write SIMD code that adds all the elements up together

we want to write SIMD code

```
// Initialize sum vector to {0, 0, 0, 0}
__m128i sum_vec = _mm_setzero_si128();

// Load array elements 0-3 into a temporary vector register
__m128i tmp = _mm_loadu_si128((__m128i *) arr);
// Add to existing sum vector
sum_vec = _mm_add_epi32(sum_vec, tmp);
// sum_vec = {3, 1, 4, 1}

// Load array elements 4-7 into a temporary vector register
tmp = _mm_loadu_si128((__m128i *) (arr + 4));
// Add to existing sum vector
sum_vec = _mm_add_epi32(sum_vec, tmp);
// sum_vec = {3 + 5, 1 + 9, 4 + 2, 1 + 6}
```

sum_vec:

8	10	6	7
---	----	---	---

tmp:

5	9	2	6
---	---	---	---

```
// Create temporary array to hold values from sum_vec
// We must store the vector into an array in order to access the
int tmp_arr[4];
_mm_storeu_si128((__m128i *) tmp_arr, sum_vec);
// Collect values from sum_vec in a single integer
int sum = tmp_arr[0] + tmp_arr[1] + tmp_arr[2] + tmp_arr[3];
```

Finally, we store sum_vec into a temporary array and then add up all 4 elements of that array



Loop unrolling

- Unrolling a loop (more operations per iteration of loop) will result in slightly improved performance
- Unrolling a loop with SIMD functions will result in even better performance

Loop Unrolling Example

```
int N = 100;
int arr[N];

for (int i = 0; i < N; i += 1) {
    arr[i] = i;
}
```



```
int N = 100;
int arr[N];

for (int i = 0; i < N; i += 4) {
    arr[i] = i;
    arr[i + 1] = i + 1;
    arr[i + 2] = i + 2;
    arr[i + 3] = i + 3;
}
```


Loop Unrolling Example with tail case

```
int N = 103;  
int arr[N];  
  
for (int i = 0; i < N; i += 1) {  
    arr[i] = i;  
}
```



```
int N = 103;  
int arr[N];  
  
for (int i = 0; i < N / 4 * 4; i += 4) {  
    arr[i] = i;  
    arr[i + 1] = i + 1;  
    arr[i + 2] = i + 2;  
    arr[i + 3] = i + 3;  
}  
  
for (int i = N / 4 * 4; i < N; i += 1) {  
    arr[i] = i;  
}
```