

CS61C: Great Ideas in Computer Architecture (aka Machine Structures)

Lecture 19: Caches II: Direct-Mapped and Set Associative, Types of Misses

Instructors: Rosalie Fang, **Charles Hong**, Jero Wang

61C in the News: New Intel ISA Extensions

7/25/2023:

“Intel has announced two new x86-64 instruction sets designed to bolster and offer more performance in AVX-based workloads with their hybrid architecture of performance (P) and efficiency (E) cores [...]

[Intel has also published a technical paper detailing their new AVX10](#), enabling both Intel's performance (P) and efficiency (E) cores to support the converged AVX10/256-bit instruction set going forward. This means that Intel's future generation of hybrid desktop, server, and workstation chips will be able to support multiple AVX vectors, including 128, 256, and 512-bit vector sizes throughout the entirety of the cores holistically [...]

The idea behind APX is to **allow access to more registers** [...] with Intel claiming 10% fewer loads and 20% fewer stores when the code is compiled for APX versus the same code for x86-64 using Intel 64 [...] Another essential feature of Intel's APX is its support for **three-operand instruction formats**.”



<https://www.anandtech.com/show/18975/intel-unveils-avx10-and-apx-isas-unifying-avx512-for-hybrid-architectures->

Announcements

- P/NP deadline is Friday, July 28
 - Student support meetings available with course staff (see Ed)
 - As always, instructor OH open for support, course feedback, or anything else!
- Exam prep sections continue
 - Parallelism and some caches: Thursday 1-3pm in Soda 380, Friday 1-3pm in Cory 540AB
- HW5 due tonight!
- Labs 7 and 8 due tomorrow, Thursday 7/27


Agenda

- Direct-Mapped Cache
- Set Associate Cache
- Types of Cache Misses

Direct-Mapped Cache

Problems with Fully-Associative Caches

- While nice, an actual implementation of a fully associative cache has certain problems:
 - Hard to find a good and efficient replacement policy
 - Requires a lot of comparators/circuitry.



Problems with Fully-Associative Caches

- While nice, an actual implementation of a fully associative cache has certain problems:
 - Hard to find a good and efficient replacement policy
 - Requires a lot of comparators/circuitry
- Reason: Any slot in our cache could store our data, so we need to check all spots
- Library analogy: We don't have any organization scheme in our bookshelf, so it takes a long time to check if our data is in our bookshelf
- Cache tag checking circuit may be in CPU critical path
- Solutions?

Problems with Fully-Associative Caches

- Idea #1: Sort the data in the cache, so we can check in log time.
 - Works for a bookshelf; doesn't work in caches
 - Takes more time to move data between blocks than to check tag bits
 - Would still need a lot of comparators to do this
- Idea #2: Separate books into groups and only assign books to a cache block matching their group
 - Library analogy: Have the top shelf for books that start with "A", the second shelf for books that start with "B", etc. That way, we only need to check 1/26 of our bookshelf.

Idea 2: Picking a grouping mechanism


- Books can be grouped based on genre, name, etc. How can we group cache blocks?
- Hash the address?
 - Used in 61B to distribute data into even groups, but we don't want the extra hardware
 - Plus, we can do better than random
- Top N bits of the address?
 - Doesn't distribute blocks well
 - In practice, we'll be using consecutive blocks of memory, and chances are that consecutive blocks will have the same top N bits of memory address
- **Bottom N bits of the tag:** Works well with consecutive blocks of memory
 - Also scales to arbitrarily many groups without much additional circuitry
 - Note: Number of groups must be a power of 2.

Index Bits

- Bottom N bits of the tag: Works well with consecutive blocks of memory
 - Also scales to arbitrarily many cache slots without much additional circuitry
- A memory address can be divided into a tag, **an index**, and an offset
- Ex. Let's say we have a block size of 4096 and 16 distinct groups
 - Ex. **0xDEADBEEF**
 - This address is in Block 0xDEADB
 - This block has tag 0xDEAD and index 0xB
 - We want the 0xEEFth byte in that block
- Determining the number of bits in each component of a cache address is called the "**TIO** breakdown", for "Tag, Index, Offset".

Direct-Mapped Cache

- Each slot in a direct-mapped cache is assigned a unique index; if we have 32 slots, one will be for index 0, one for index 1, ... , and one for index 31
- When a memory access occurs:
 - Step 1: Check **the unique slot that could contain our data**. If the tag matches and the valid bit is on, cache hit.
 - Step 2: Otherwise, cache miss
- A direct-mapped cache is parameterized on two aspects: The block size, and the number of blocks that can be stored in the cache
 - Same as fully-associative
- Library analogy: You buy a bookshelf with N slots, each holding 1 book. Each book can only go in 1 specific slot
 - For a given book, only need to check one slot of your bookshelf
 - However, you might have to return books while slots are still empty.



Direct-Mapped Cache: Example Memory Access (1/2)

- Say we have a direct-mapped cache with 4 byte blocks and 4 blocks storage
- Addresses are 10 bits
- To access address 0x3CD:
 - Split into TIO:
 - 4 bytes blocks -> O = 2 bits
 - 4 indexes -> I = 2 bits
 - 0x3CD == 0b11 1100 11 01
 - For index 3, tag matches, and valid bit is on, so hit
 - We get the 1st byte of the block, which is 0xDE.

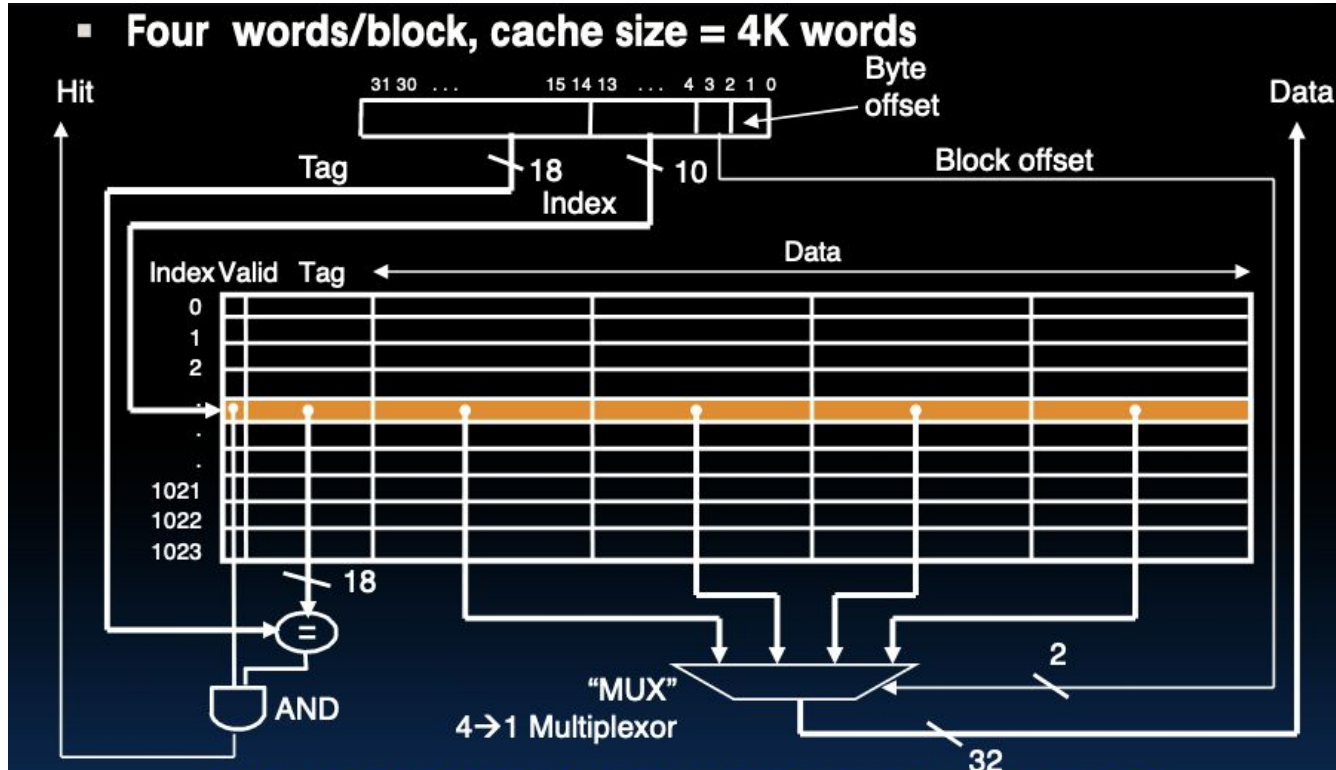
Index	Tag	V	Data
0	0b10 0110	0	0xDE 0xAD 0xBE 0xEF
1	0b10 1100	0	0x01 0x23 0x45 0x67
2	0b01 1010	1	0xAB 0xAD 0xCA 0xFE
3	0b11 1100	1	0xAC 0xDE 0xFF 0x61

Direct-Mapped Cache: Example Memory Access (2/2)

- To access address 0x3C1:
 - Split into TIO:
 - 4 bytes blocks -> O = 2 bits
 - 4 indexes -> I = 2 bits
 - 0x3CD == 0b11 1100 00 01
 - For index 0, tag doesn't match, so miss, even though the correct tag is in another index
 - Note: same tag + different index is a different memory address.

Index	Tag	V	Data
0	0b10 0110	1	0xDE 0xAD 0xBE 0xEF
1	0b10 1100	0	0x01 0x23 0x45 0x67
2	0b01 1010	1	0xAB 0xAD 0xCA 0xFE
3	0b11 1100	1	0xAC 0xDE 0xFF 0x61

Direct-Mapped Cache Circuit



Direct-Mapped Cache Pros and Cons

- Pros:

- Simpler circuit
 - Compare one tag rather than all tags
 - Faster accesses
 - Scales better than fully associative caches
- No replacement policy
 - Only one block can be evicted, so no need to store LRU

- Cons:

- If an index isn't used, we end up with free space in our cache that we can't use
- Let's see how this behaves in Matrix Multiply with 4 blocks.

Caching Example: Write-Back Direct-Mapped

- 16-bit address, 16-byte blocks
- TIO?
 - 10, 2, 4

Index	Tag	V	Dirty	Data
0	0x100	0	0	0xBF 0x16 0x88 0x2B
1	0x133	0	0	0x3B 0x18 0xF1 0xB3
2	0x156	0	0	0xE6 0x57 0x49 0xEE
3	0x0E9	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

Caching Example: Write-Back Direct-Mapped

- Start: Cache starts cold
 - Note: Index field, and no LRU field
 - Index isn't actually stored, but we'll show it anyway
 - Note: Matrix B has been transposed into Bt to optimize cache efficiency
- 0 misses, 0 hits

Index	Tag	V	Dirty	Data
0	0x100	0	0	0xBF 0x16 0x88 0x2B
1	0x133	0	0	0x3B 0x18 0xF1 0xB3
2	0x156	0	0	0xE6 0x57 0x49 0xEE
3	0x0E9	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

Caching Example: Write-Back Direct-Mapped

- Access A[0]: 0x1000
 - Index 0b00 = 0, Tag 0x040 (top 10 bits)
 - Miss, so replace the block in index 0

Index	Tag	V	Dirty	Data
0	0x100	0	0	0xBF 0x16 0x88 0x2B
1	0x133	0	0	0x3B 0x18 0xF1 0xB3
2	0x156	0	0	0xE6 0x57 0x49 0xEE
3	0x0E9	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

Caching Example: Write-Back Direct-Mapped

- Access B[0]: 0x2000
 - Index 0, Tag 0x060
 - Miss, so replace the block in index 0
 - Uh oh

Index	Tag	V	Dirty	Data
0	0x040	1	0	0x01 0x02 0x03 0x04
1	0x133	0	0	0x3B 0x18 0xF1 0xB3
2	0x156	0	0	0xE6 0x57 0x49 0xEE
3	0x0E9	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

Caching Example: Write-Back Direct-Mapped

- Access A[1]: 0x1001
 - Index 0, Tag 0x040
 - Miss, so replace the block in index 0
 - Oh no

Index	Tag	V	Dirty	Data
0	0x060	1	0	0x11 0x15 0x19 0x1D
1	0x133	0	0	0x3B 0x18 0xF1 0xB3
2	0x156	0	0	0xE6 0x57 0x49 0xEE
3	0x0E9	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

Caching Example: Write-Back Direct-Mapped

- B[1]: Miss, A[2]: Miss, B[2]: Miss, ... , C[0]: Miss
- 9 misses, 0 hits...
- A[0]: Miss, B[4]: Miss, but now address is 0x2010, which is index 1

Index	Tag	V	Dirty	Data
0	0x040	1	0	0x01 0x02 0x03 0x04
1	0x133	0	0	0x3B 0x18 0xF1 0xB3
2	0x156	0	0	0xE6 0x57 0x49 0xEE
3	0x0E9	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

Caching Example: Write-Back Direct-Mapped

- A[1]: Hit, B[5]: Hit, ... , C[1]: Miss, evicting Block 0x1000
- While calculating C[1]: 3 misses, 6 hits
 - Better!

Index	Tag	V	Dirty	Data
0	0x040	1	0	0x01 0x02 0x03 0x04
1	0x060	1	0	0x12 0x16 0x1A 0x1E
2	0x156	0	0	0xE6 0x57 0x49 0xEE
3	0x0E9	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

250			

Caching Example: Write-Back Direct-Mapped

- In general: if the element of C is on the diagonal, 9 misses, 0 hits
- Otherwise, 3 misses, 6 hits
- Total: 72 misses, 72 hits, or 50% hit rate
 - Back to pre-transpose hit rate!

Index	Tag	V	Dirty	Data
0	0x080	1	1	0xFA 0x04 0x?? 0x??
1	0x060	1	0	0x12 0x16 0x1A 0x1E
2	0x156	0	0	0xE6 0x57 0x49 0xEE
3	0x0E9	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

250			

Problem with Direct-Mapped Caches

- If we access two distinct segments of memory that access the same index, we end up continuously evicting blocks
 - This is known as **thrashing**, and causes a significant increase in misses
- Library analogy: may have unnecessary evictions if we don't have an even distribution of books
 - Not many books will start with "X", for example, so that space will likely go unused.
- Problem: How do we get the benefit of direct-mapped caches without the downsides?

Set Associative Cache

Meeting in the Middle: The Set-Associative Cache


- Direct-mapped caches thrash a lot, but fully associative caches use a lot of comparators
- Solution: Instead of just one slot per index, assign a small number of slots per index
- The **associativity** of a cache is the number of slots assigned to each index.

N-way Set Associative Cache (1/3)

- An N-way set associative cache is parameterized on **three** aspects: The block size, the number of blocks that can be stored in the cache, and the associativity N.
 - The other two types were parameterized on block size and number of blocks only.
- Blocks are split into groups of N, and each group of blocks is assigned a unique index
- When a memory access occurs:
 - Step 1: Check **the slots that could contain our data**. If the tag matches and the valid bit is on for any of these slots, cache hit
 - Step 2: Otherwise, cache miss.

N-way Set Associative Cache (2/3)

- Library analogy: You buy a bookshelf. You set it up so that each shelf can only hold a certain type of book, and promise never to put books in the wrong shelf, even if it wastes space. That way, you only need to check a few blocks on a given shelf.



N-way Set Associative Cache: Example

- Say we have a 2-way set-associative cache with 4 byte blocks and 4 blocks storage
- Addresses are 10 bits
- To access address 0x3CD:
 - Split into TIO:
 - 4 bytes blocks -> O = 2 bits
 - $4/2 = 2$ indexes -> I = 1 bits
 - $0x3CD == 0b1111001101$
 - For index 1, the tag matches one of our blocks, and the valid bit is on, so we get the data
 - We get the 1st byte of the block, which is 0xDE.


I	Tag	V	Data
0	0b100 0110	1	0xDE 0xAD 0xBE 0xEF
0	0b101 1100	0	0x01 0x23 0x45 0x67
1	0b011 1010	1	0xAB 0xAD 0xCA 0xFE
1	0b111 1001	1	0xAC 0xDE 0xFF 0x61

Calculating TIO Breakdown

- # of Offset bits = $\log_2(\text{block size in bytes})$
- # of Index bits = $\log_2(\text{\# of sets})$
 - Direct-mapped: # of sets = # of blocks
 - Fully associative: # of sets = 1
- # of Tag bits = Address length - # of Index bits - # of Offset bits

4-way Set Associative Cache Circuit

- 1024 blocks
- $1024 / 4 = 256$ sets
- 4B blocks
- 32-bit address
- TIO?
 - 22, 8, 2.



N-way Set Associative Cache (3/3)

- Direct-mapped and Fully Associative caches can be considered special cases of the N-way set associative cache:
 - Direct-mapped == 1-way set associative
 - Fully associative == N = block count-way set associative
- Generally ends up with most of the performance of a fully associative cache and most of the circuit simplicity of a direct-mapped cache
- Let's see it in action with a 2-way set associative cache with 4 blocks (2 sets).

Caching Example: Write-Back LRU 2-way Set Associative

- Start: Cache starts cold
- A[0]: Miss, B[0]: Miss (but this time no eviction)
- 2 misses, 0 hits

I	Tag	LRU	V	D	Data
0	0x100	0	0	0	0xBF 0x16 0x88 0x2B
0	0x333	0	0	0	0x3B 0x18 0xF1 0xB3
1	0x156	0	0	0	0xE6 0x57 0x49 0xEE
1	0x0E9	0	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

Caching Example: Write-Back LRU 2-way Set Associative

- A[1] Hit, B[1] Hit, ... , B[3] Hit, C[0] Miss with eviction
 - Still have a bit of thrashing, but it's much less now
- Total: 6 Hits, 3 Misses
 - Same as Fully Associative cache!

I	Tag	LRU	V	D	Data
0	0x080	2	1	0	0x01 0x02 0x03 0x04
0	0x100	1	1	0	0x11 0x15 0x19 0x1D
1	0x156	0	0	0	0xE6 0x57 0x49 0xEE
1	0x0E9	0	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

Caching Example: Write-Back LRU 2-way Set Associative

- A[0] miss (evict block 0x200 with tag 0x100 and index 0), B[4] miss (index 1)
- A[1] - B[7] hit, C[1] hit
- Total for this element: 2 misses, 7 hits

I	Tag	LRU	V	D	Data
0	0x180	1	1	1	0xFA 0x?? 0x?? 0x??
0	0x100	2	1	0	0x11 0x15 0x19 0x1D
1	0x156	0	0	0	0xE6 0x57 0x49 0xEE
1	0x0E9	0	0	0	0xB5 0x81 0x67 0x3F

A			
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt	17	21	25	29
	18	22	26	30
	19	23	27	31
	20	24	28	32

C			
250			

Caching Example: Write-Back LRU 2-way Set Associative

- Computing C[2]: A[0] hits, B[8] misses, evicting block 0x300 ... , C[2] misses
- Total for this element: 2 misses, 7 hits
- C[3]: Same as C[1], so 2 misses, 7 hits

I	Tag	LRU	V	D	Data
0	0x180	1	1	1	0xFA 0x04 0x?? 0x??
0	0x080	2	1	0	0x01 0x02 0x03 0x04
1	0x100	1	1	0	0x12 0x16 0x1A 0x1E
1	0x0E9	0	0	0	0xB5 0x81 0x67 0x3F

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

250	260		

Caching Example: Write-Back LRU 2-way Set Associative

- C[4]: Similar to C[0], except this time, the A and C blocks are in index 1
 - 3 misses, 6 hits

I	Tag	LRU	V	D	Data
0	0x080	2	1	0	0x01 0x02 0x03 0x04
0	0x180	1	1	1	0xFA 0x04 0x0E 0x18
1	0x100	2	1	0	0x12 0x16 0x1A 0x1E
1	0x101	1	1	0	0x14 0x18 0x1C 0x20

A			
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt	17	21	25	29
	18	22	26	30
	19	23	27	31
	20	24	28	32

C			
250	260	270	280

Caching Example: Write-Back LRU 2-way Set Associative

- C[5]: Acts like C[2] (B block on same index as A and C block)
 - 2 misses, 7 hits
- C[6],C[7]: 2 misses, 7 hits each

I	Tag	LRU	V	D	Data
0	0x100	1	1	0	0x11 0x15 0x19 0x1D
0	0x180	2	1	1	0xFA 0x04 0x0E 0x18
1	0x100	2	1	0	0x12 0x16 0x1A 0x1E
1	0x181	1	1	1	0x6A 0x?? 0x?? 0x??

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

250	260	270	280
618			

Caching Example: Write-Back LRU 2-way Set Associative

- C[8]-C[11]: Same as first row
- C[12]-C[15]: Same as second row
- Total: 36 misses, 108 hits
- 75% hit rate; not as good as FA, but much better than DM!

I	Tag	LRU	V	D	Data
0	0x100	2	1	0	0x11 0x15 0x19 0x1D
0	0x101	1	1	0	0x13 0x17 0x1B 0x1F
1	0x081	2	1	0	0x05 0x06 0x07 0x08
1	0x181	1	1	1	0x6A 0x84 0x9E 0xB8

A

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bt

17	21	25	29
18	22	26	30
19	23	27	31
20	24	28	32

C

250	260	270	280
618	658	670	696
986	1028	1070	1112
1354	1412	1470	1528

Types of Cache Misses

Miss Classifications: Definitions in this class

- **Compulsory Misses:** Misses that would occur even if the cache was a fully associative cache with infinite space
- **Capacity Misses:** Misses that are neither conflict nor compulsory
- **Conflict Misses:** Misses that would not have happened if the cache was fully associative under at least one consistent replacement policy
- **Consistent replacement policy:** A replacement policy that keeps all the data the lower-associativity cache would
 - Any replacement policy such that the data in the Fully Associative cache is a superset of the data in our actual cache.
 - This definition guarantees that we don't have "conflict hits", but also tends to classify misses more as capacity misses than as conflict misses.

Miss Classifications: Simplified Equivalent Definitions

- If a miss occurs, look at the most recent time that that block was evicted
- If that block was never evicted before, it's a **compulsory miss**
- If the block was evicted when the cache was full (no empty spaces), it's a **capacity miss**
- If the block was evicted when the cache was not full (at least one empty space), it's a **conflict miss**
- Let's look again at the 2-way set associative cache example.

Caching Example: Hit Classifications

- Start: Cache starts cold
 - We'll restrict the cache to only the parts that are important (no data/dirty/etc, block instead of tag).
 - Also will write the status of each block on the right
 - Blocks will be ordered by LRU within a set (unlike an actual cache)
- First miss: A[0]
 - Compulsory miss

Index	Block	Index	Block
0	–	1	–
0	–	1	–

Block	Status	Block	Status
A[0]	NA	B[8]	NA
A[4]	NA	B[12]	NA
A[8]	NA	C[0]	NA
A[12]	NA	C[4]	NA
B[0]	NA	C[8]	NA
B[4]	NA	C[12]	NA

Caching Example: Hit Classifications

- Next miss is B[0]
 - Compulsory Miss
- Next miss is C[0]
 - Compulsory Miss, evicts A[0]
 - A[0] was evicted even though there were empty slots, so the next access of A[0] will be a conflict miss

Index	Block	Index	Block
0	A[0]	1	-
0	-	1	-

Block	Status	Block	Status
A[0]	In Cache	B[8]	NA
A[4]	NA	B[12]	NA
A[8]	NA	C[0]	NA
A[12]	NA	C[4]	NA
B[0]	NA	C[8]	NA
B[4]	NA	C[12]	NA

Caching Example: Hit Classifications

- A[0]
 - Conflict miss
 - Evicts B[0] (will cause conflict miss)
- B[4]
 - Compulsory miss
- B[8]
 - Compulsory miss
 - Evicts C[0] (conflict)

Index	Block	Index	Block
0	B[0]	1	-
0	C[0]	1	-

Block	Status	Block	Status
A[0]	Conflict	B[8]	NA
A[4]	NA	B[12]	NA
A[8]	NA	C[0]	In Cache
A[12]	NA	C[4]	NA
B[0]	In Cache	C[8]	NA
B[4]	NA	C[12]	NA

Caching Example: Hit Classifications

- C[2]
 - Conflict miss
 - Evicts A[0] (conflict)
- A[0]
 - Conflict miss
 - Evicts B[8] (conflict)
- B[12]
 - Compulsory miss
 - At this point, all blocks are used, so any further evictions will cause capacity misses

Index	Block	Index	Block
0	A[0]	1	B[4]
0	B[8]	1	-

Block	Status	Block	Status
A[0]	In Cache	B[8]	In Cache
A[4]	NA	B[12]	NA
A[8]	NA	C[0]	Conflict
A[12]	NA	C[4]	NA
B[0]	Conflict	C[8]	NA
B[4]	In Cache	C[12]	NA

Caching Example: Hit Classifications

- A[4]
 - Compulsory miss
 - Evicts B[4] (will cause a capacity miss next)
- B[0]
 - Conflict miss
 - Evicts C[0] (capacity)
- C[4]
 - Compulsory miss
 - Evicts B[12] (capacity)

Index	Block	Index	Block
0	C[0]	1	B[4]
0	A[0]	1	B[12]

Block	Status	Block	Status
A[0]	In Cache	B[8]	Conflict
A[4]	NA	B[12]	In Cache
A[8]	NA	C[0]	In Cache
A[12]	NA	C[4]	NA
B[0]	Conflict	C[8]	NA
B[4]	In Cache	C[12]	NA

Caching Example: Hit Classifications

- B[4]
 - Capacity miss
 - Evicts A[4] (capacity)
- A[5]
 - Capacity miss
 - Evicts C[4] (capacity)
- C[5]
 - Capacity miss
 - Evicts B[4] (capacity)
 - This is the downside of this definition (this feels like it should be a conflict miss)

Index	Block	Index	Block
0	A[0]	1	A[4]
0	B[0]	1	C[4]

Block	Status	Block	Status
A[0]	In Cache	B[8]	Conflict
A[4]	In Cache	B[12]	Capacity
A[8]	NA	C[0]	Capacity
A[12]	NA	C[4]	In Cache
B[0]	In Cache	C[8]	NA
B[4]	Capacity	C[12]	NA

Summary

- **Direct-mapped caches** allow for a simpler implementation of caches, but may decrease cache hit rate
- **Set associative caches** are a middle ground between fully associative and direct-mapped caches
- We classify cache misses into **compulsory, capacity, and conflict misses** to analyze performance more easily (tomorrow)
- Tomorrow: Analyzing cache performance, and advanced cache techniques!