

CS61C: Great Ideas in Computer Architecture (aka Machine Structures)

Lecture 20: Caches II: Performance, Multilevel, Coherence

Instructors: Rosalie Fang, **Charles Hong**, Jero Wang

Announcements

- P/NP deadline is Friday, July 28
 - Student support meetings available with course staff (see Ed)
 - As always, instructor OH open for support and course feedback!
- Discussion cancellations
 - Imran's Monday 12pm - 1pm discussion is cancelled
 - Imran's Monday 1pm - 2pm discussion will now be taught by Ben
 - Ben's Wednesday 12pm - 1pm discussion is cancelled starting August 2nd
- Labs 7 and 8 due today, Thursday 7/27
- Project 3B due next Tuesday 8/1

Agenda

- Analyzing Cache Performance
- Multi-level Caches
- Caches in Real CPUs
- Cache Coherence

Analyzing Cache Performance

Analyzing Cache Performance: Terminology

- The **hit rate** of a cache is the percentage of accesses that result in a hit
- The **miss rate** of a cache is the percentage of accesses that result in a miss
- **Hit time** is how long it takes to check the cache. The **miss penalty** is how long it takes to access main memory after a miss.
 - On cache hit: access time = hit time
 - On cache miss: access time = hit time + miss penalty
- The **Average Memory Access Time** (or **AMAT**) is the **average** amount of time it takes for **one memory access**, given a program.

Computing AMAT (1/2)

- If our hit rate was $\frac{3}{4}$ (in yesterday's example), a hit takes 10 cycles, and a miss takes 100 cycles, what speedup did our cache provide?
- First step: find new AMAT
 - $\frac{3}{4}$ of our memory accesses were **hits**, and therefore took **10 cycles each**
 - $\frac{1}{4}$ of our memory accesses were **misses**, and therefore took $100 + 10 = \mathbf{110 \text{ cycles each}}$
 - The average time for **4 memory accesses** is $3 * (\text{hit time}) + 1 * (\text{miss time}) = \mathbf{140 \text{ total cycles}}$
 - AMAT (for **1 memory access**) is $140/4 = \mathbf{35 \text{ cycles}}$
- Our original AMAT would have been 100 cycles per access
 - Speedup is $100/35 = \mathbf{2.86x}$

Computing AMAT (2/2)

- Alternate formula:

$$\text{AMAT} = \text{hit time} + (\text{miss rate} * \text{miss penalty})$$

- $\text{AMAT} = 10 + (\frac{1}{4} * 100) = \mathbf{35 \text{ cycles}}$
- Speedup is $100/35 = \mathbf{2.86x}$


Computing Hit Rate

- Computing AMAT can get tricky for certain programs
 - Need to determine the hit rate of a program
- General Tips:
 - Typically hit rate starts low, because the cache is cold
 - After warming up, the cache normally reaches a "steady state", where a pattern of hits and misses appears every iteration of some outer loop, until we completely exhaust the memory in a block
 - Once we start a new block, the cache acts cold again, and the cycle repeats.
 - Once you find this cycle, it's possible to use that to get the total hit rate.
- More practice: Homework

Optimizing Cache Hardware: Block Size

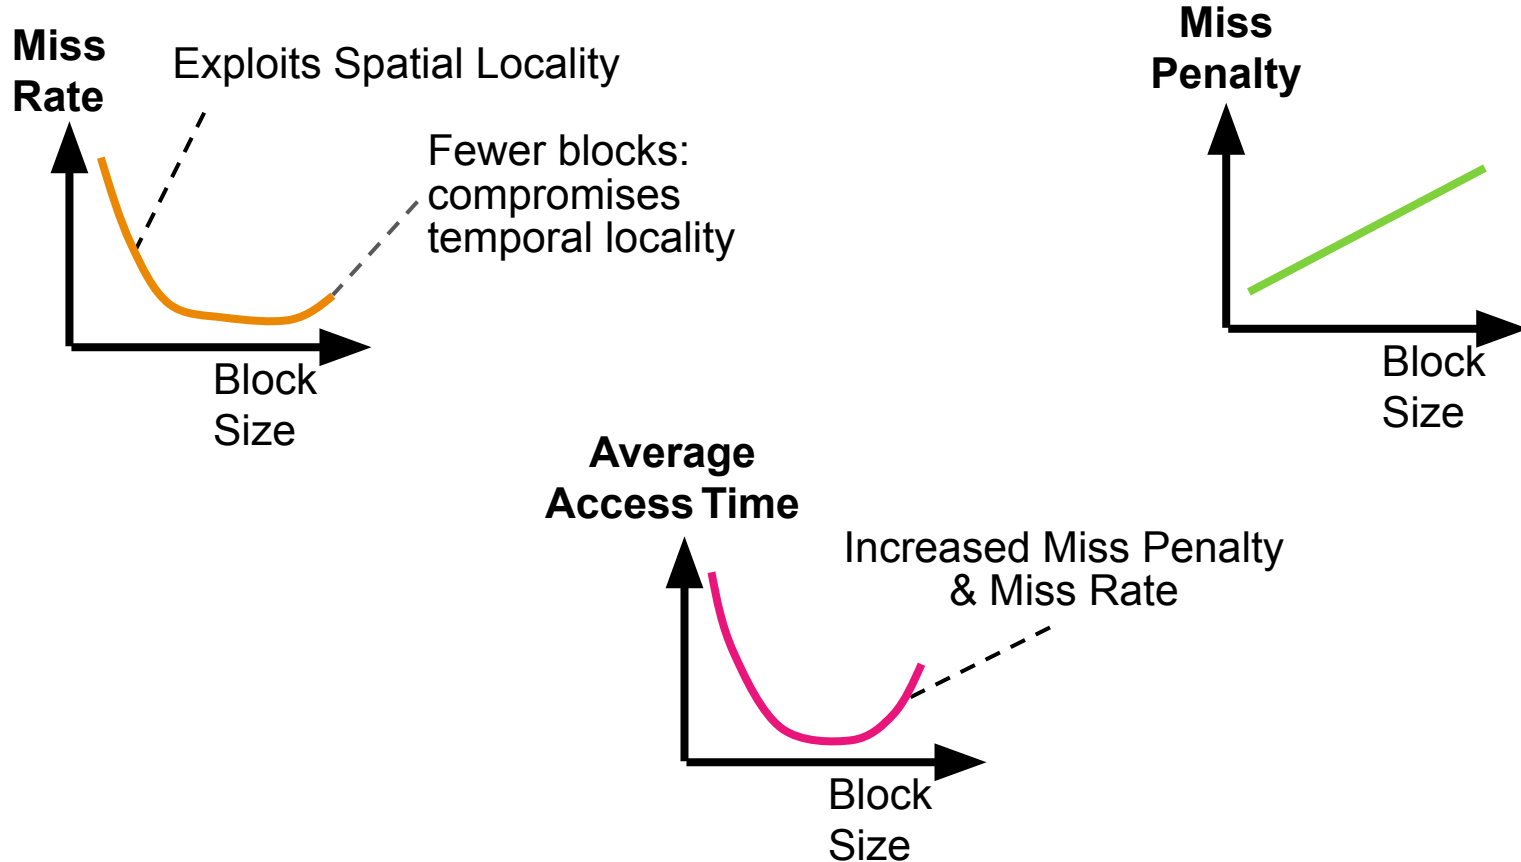
- Benefits of Larger Block Size
 - **Spatial locality**: if we access a given word, we're likely to access other nearby words soon
 - Reduces compulsory misses
 - Works well for programs
 - Works well for sequential array accesses
- Drawbacks of Larger Block Size
 - Larger block size means larger miss penalty
 - On a miss, takes longer time to load a new block from next level
 - Effect on miss rate?
 - Next slide →

Extreme Example: One Big Block



- Cache Size = 1024 bytes, Block Size = 1024 bytes
 - Only **ONE** entry (row) in the cache!
- If item accessed, likely accessed again soon
 - But unlikely will be accessed again immediately!
- The next access will likely to be a miss again
 - Continually loading data into the cache but discard data (force out) before it gets reused
 - Nightmare for cache designer: **Ping Pong Effect**
 - Increases conflict misses

Block Size Tradeoff Conclusions



Cache Optimizations: Hardware


- Larger block size
 - + reduces compulsory misses
 - increases conflict misses and miss penalty
- Larger cache size
 - + reduces capacity misses
 - increases hit time
- Higher associativity
 - + reduces conflict misses
 - increases hit time

Cache Optimizations: Software

- If **compulsory misses** are most common
 - Prefetching
- If **capacity misses** are most common
 - Using too much memory: reduce the "working set"
 - Split code into parts that access only as much memory as your cache can store
- If **conflict misses** are most common
 - Reduce the number of distinct chunks of memory you're accessing
 - E.g. By transposing your B matrix in matrix multiply


Multilevel Caches

Multilevel Caches (1/2)



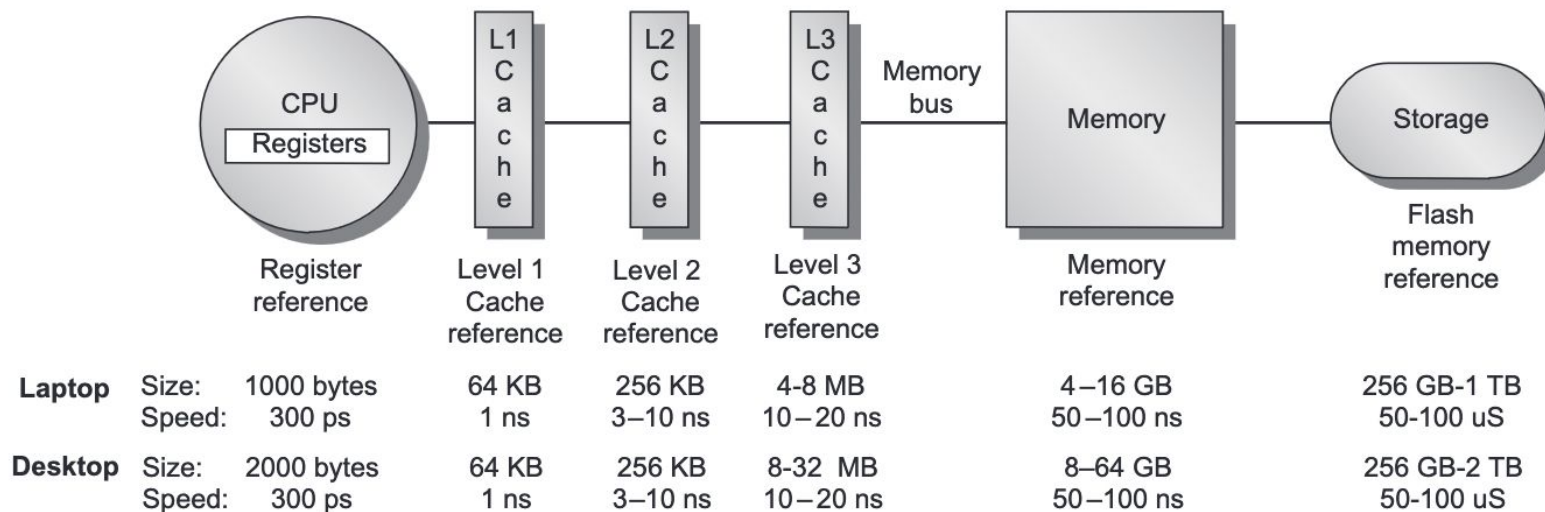
- Current setup: If cache hit, then fast access. If cache miss, then slow access
- Problem:
 - Making the cache bigger/more complex increases hit rate, but also increases hit time
 - With just one cache, we need to make some tough trade-offs.

Multilevel Caches (2/2)



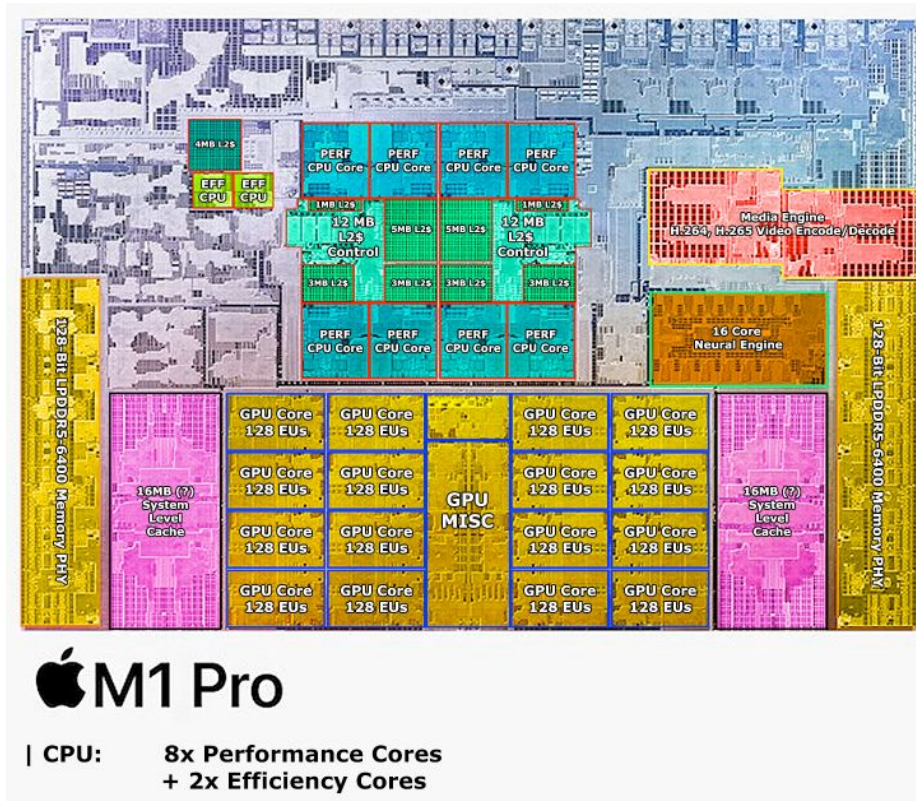
- Solution: Add **multiple layers of caches**.
 - Check L1 cache. If hit, we're done
 - If miss, check L2 cache. If hit, we're done, and bring data up to L1 cache
 - If miss, check L3 cache. If hit, we're done, and bring data up to L1 and L2 cache
 - ...
- Each level of cache is larger than the previous and holds all of the data in the previous cache
- Library analogy: We buy a bookshelf for right next to us, then we buy a larger bookshelf to put in our garage, then make a shed to store more books
 - Slight difference: to act exactly like a cache, each shelf/shed would need a copy of each book stored in the inner level.

Multilevel Cache Sizes/Latencies



Multilevel Cache Organization

- Generally, the L1 cache is small, but attached to the core
 - Each core gets its own L1 cache
 - Fast access, but also fairly small
 - Split into instruction and data caches
- L2 cache is larger and slightly slower, but usually a bit farther away
 - Sometimes shared, sometimes on the core.



Multilevel Cache Hit/Miss Times

- Using the numbers from the previous slide as examples:
 - L1 hit time = 1 ns
 - L2 hit time = 1 ns + 10 ns = 11 ns
 - L3 hit time = 1 ns + 10 ns + 20 ns = 31 ns
 - L3 miss time = 1 ns + 10 ns + 20 ns + 100 ns = 131 ns
- Hit/miss rate is calculated only for accesses that "reach" that cache
 - L2 hit rate is % of L2 accesses that hit, so L1 hits don't count as L2 hits OR L2 misses.

AMAT in Multilevel Caches

- Let's say we have a memory system with the following properties. What would be the AMAT of this system?
- 50% of accesses are L1 hits
 - 50% of accesses take 1ns
- 50% of accesses are L1 misses
 - 75% of this 50% are L2 hits (37.5% of total)
 - 87.5% of accesses are now accounted for
- 12.5% of accesses are L2 misses
 - 80% of 12.5% = 10% are L3 hits
 - 97.5% of accesses are now accounted for
- 2.5% of accesses are L3 misses
 - And therefore access DRAM
- Total: $1 \text{ ns} + 0.5 * 10 \text{ ns} + 0.125 * 20 \text{ ns} + 0.025 * 100 \text{ ns}$
- 11 ns access time \approx 9x speedup.

	Hit time	Hit rate
L1 Cache	1 ns	50%
L2 Cache	10 ns	75%
L3 Cache	20 ns	80%
DRAM	100 ns	100%

AMAT in Multilevel Caches

- Another way to think about this:

$$\text{AMAT} = \text{hit time} + \text{miss rate} * \text{miss penalty}$$

$$\text{miss penalty} = \text{AMAT of next level}$$

$$1 \text{ ns} + 0.5 * (10 \text{ ns} + 0.25 * (20 \text{ ns} + 0.2 * 100 \text{ ns}))$$

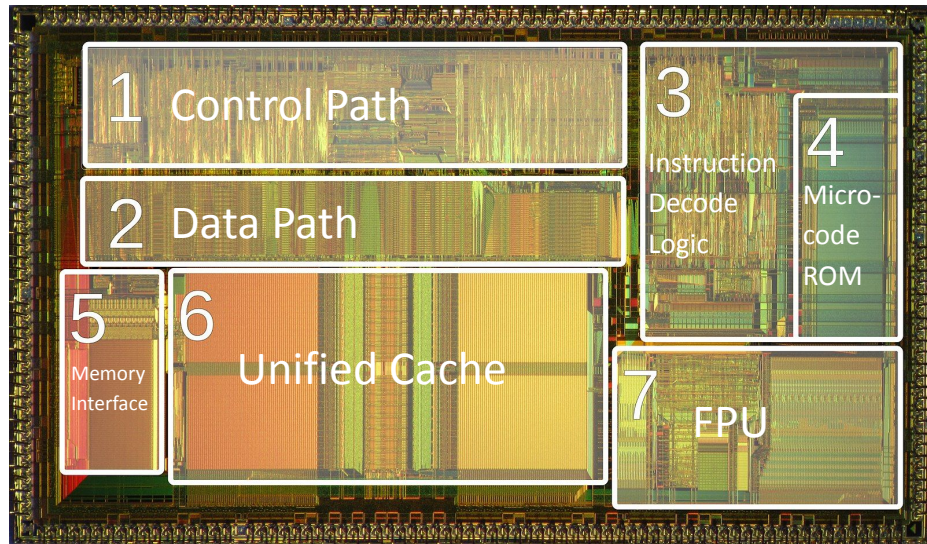
= 11 ns also!

	Hit time	Hit rate
L1 Cache	1 ns	50%
L2 Cache	10 ns	75%
L3 Cache	20 ns	80%
DRAM	100 ns	100%

Caches (and More!) in Real CPUs

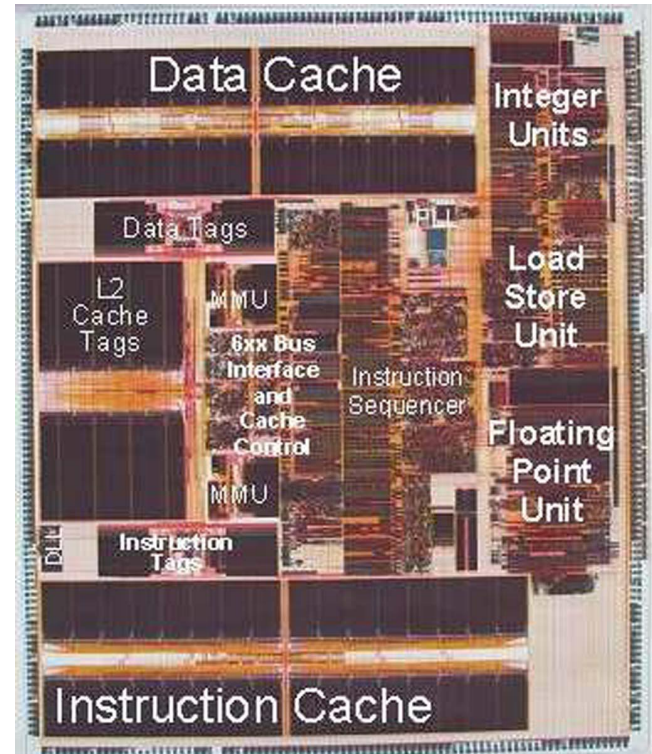
Intel 486 (1989)

- First major CPU to include on-chip cache!
- **8KiB unified L1**
- Compatible with additional external cache (acted like L2)
 - These existed before on-chip caches
- <10 cycle memory latency
 - 20-33 MHz clock frequency
- 5-stage pipelining
- Chief architect: Pat Gelsinger.



Early PowerPC (2001)

- Cache
 - **32 KiB Instruction L1 & 32 KiB Data L1 caches**
 - **External L2 cache interface** with integrated controller and cache tags, supports up to 1 MiB external L2 cache
 - Dual Memory Management Units (MMU) with Translation Lookaside Buffers (TLB)
 - TLBs are used for virtual memory (next unit)
- Pipelining
 - Superscalar (3 inst/cycle)
 - 6 execution units (2 integer and 1 double precision IEEE floating point)

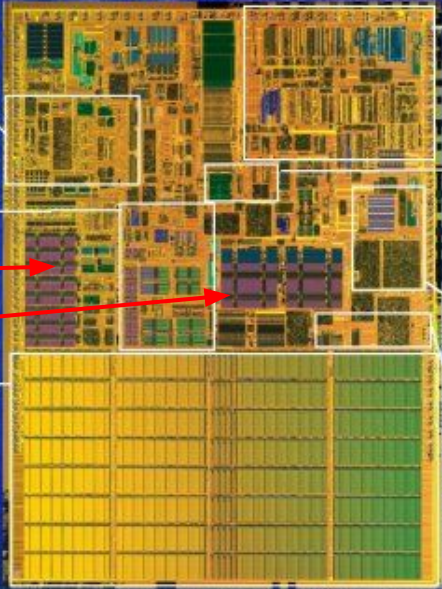


Intel Pentium M (2003)

Intel® Pentium® M Processor

New Micro Architecture

77 Million Transistors



Micro-Ops Fusion – fuses operations together to enable faster execution of instructions at lower power

Advanced Branch Prediction – fewer re-dos for increased performance

32KiB L1 I\$ →

32KiB L1 D\$ →

1MB Power Optimized L2 Cache – enables higher CPU performance

Streaming SIMD Extensions II – compatible with Pentium® 4 Processor optimized software

Dedicated Stack Management – faster instruction at lower power levels


Enhanced Intel® SpeedStep® Technology – Multiple voltages & frequency operating points

400 MHz Power Optimized System Bus – faster system bus to enhance performance at lower power levels

intel®

centrino
MOBILE TECHNOLOGY

Intel Core i7 (2010)



CPU-Z - ID : 9h78gq

CPU		Caches	Mainboard	Memory	SPD	Graphics	Bench	About
Processor								
Name	Intel Core i7 Extreme 980X							
Code Name	Guilford		Max TDP		130 W			
Package	Socket 1366 LGA							
Technology	32 nm	Core Voltage	1.208 V					
Specification Intel(R) Core(TM) i7 CPU X 980 @ 3.33GHz								
Family	6	Model	C	Stepping	2			
Ext. Family	6	Ext. Model	2C	Revision	B1			
Instructions	MMX, SSE, SSE2, SSE3, SSE3, SSE4.1, SSE4.2, EM64T, VT-x, AES							
Clocks (Core #0)								
Core Speed	1603.49 MHz							
Multiplier	x 12.0 (12 - 27)							
Bus Speed	133.63 MHz							
Rated FSB	3206.99 MHz							
Cache								
L1 Data	6 x 32 KBytes		8-way					
L1 Inst.	6 x 32 KBytes		4-way					
Level 2	6 x 256 KBytes		8-way					
Level 3	12 MBytes		16-way					
Selection Processor #1								
Cores 6 Threads 12								

CPU-Z Ver. 1.86.0.x64 Tools Validate Close

CPU-Z



CPU		Caches	Mainboard	Memory	SPD	Graphics	About
L1 D-Cache							
Size	32 KBytes		x 6				
Descriptor	8-way set associative, 64-byte line size						
L1 I-Cache							
Size	32 KBytes		x 6				
Descriptor	4-way set associative, 64-byte line size						
L2 Cache							
Size	256 KBytes		x 6				
Descriptor	8-way set associative, 64-byte line size						
L3 Cache							
Size	12 MBytes						
Descriptor	16-way set associative, 64-byte line size						
Size							
Descriptor							

CPU-Z Version 1.54 OVERCLOCKERS.HK Validate OK

- 6 cores, 12 threads
 - Hyperthreading
- Per core:
 - **32 KiB L1 I\$, 32KiB L1 D\$**
 - **256 KiB L2**
- **12 MiB shared L3 cache**

AMD Zen 3 (2020)

- “Chiplet” architecture enables multiple dies to be fabricated and integrated on one CPU
- Each chiplet has its own **32 MiB L3 cache**
 - Each core still has its own L1 and L2 caches
- “3D stacking” enables **up to 96 MiB L3 cache!**



Cache Coherence

Caching with multithreading (1/2)

- Each core has its own L1 cache, and sometimes its own L2 cache
- Most caches are write-back, so it doesn't update main memory until the block gets evicted
- If multiple threads are running with shared memory, how do we guarantee that the correct version of data is being used?

Caching with multithreading (2/2)

- Let's look at the potentially problematic scenarios:
- Simultaneous reads:
 - Should be allowed; as long as no one changes data, things are safe
- Simultaneous writes:
 - Once two different cores have conflicting data in their caches, it's difficult to return to a reasonable state
 - Example: git merge conflicts
 - Therefore, we should not allow these to happen
 - Only one thread should be able to write to a block at a time!
- Simultaneous read and write:
 - Get the read the new value after writing

MSI Protocol (1/2)

- Recall: In our cache, we keep track of whether a line is valid or not, and whether it's dirty or not
- Can we use this metadata to handle cache coherence?
- On a read:
 - Check if any other core has its dirty bit set for that block; tell that core to write back its data
- On a write:
 - If that block is valid in any other core, invalidate it (and if dirty, write back)
- Requirement: Set up a system where you can "**snoop**" on other caches to see if they have a cache block with the same tag.

MSI Protocol (2/2)

- **Invalid:** Same as valid bit off in regular cache; the block isn't in the cache
- **Shared:**
 - The block is in some other cache
 - We haven't modified this block, and we're not allowed to make modifications
 - If we need to make modifications, evict the block from everyone's cache and move to modified
 - Allows for simultaneous reads
- **Modified:** Same as valid bit on, dirty bit on in regular cache
 - The block has been read, and modified
 - Further, no other cache has this block.

MSI State Transitions

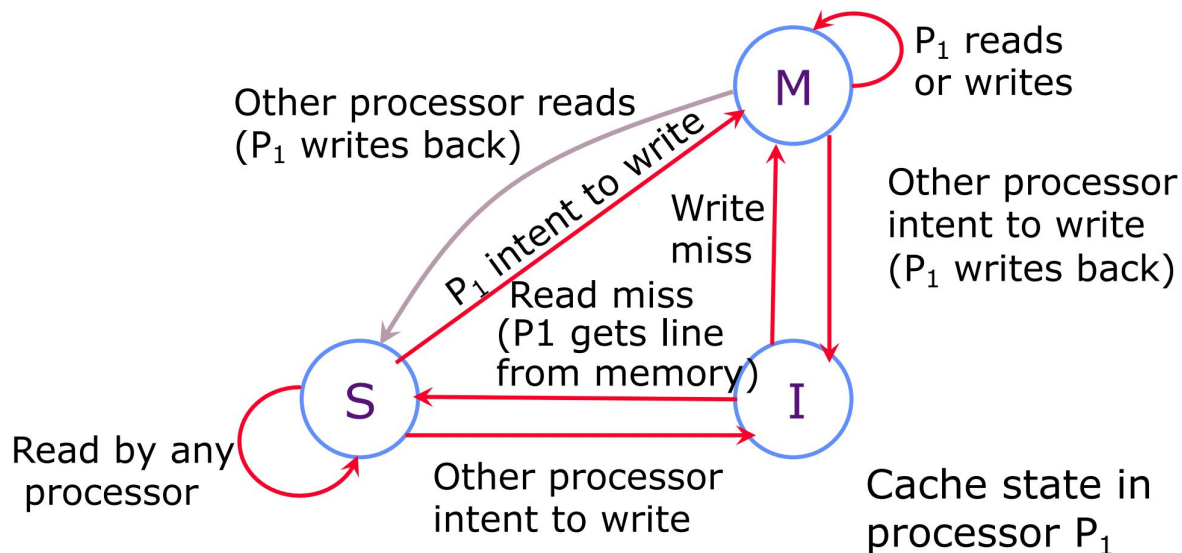
Each cache line has state bits



M: Modified

S: Shared

I: Invalid

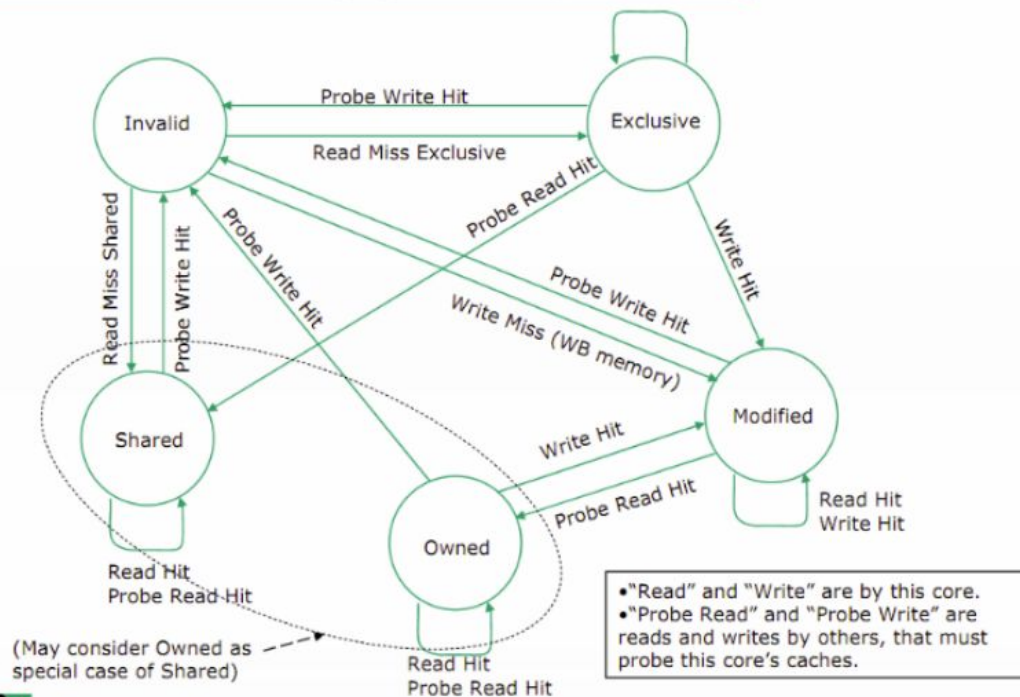


MOESI

- Adds two more states:
 - Can read then write without traffic to other cores
- **Exclusive:** Same as valid bit on, dirty bit off in regular cache
 - The block has been read, but not modified
 - Further, no other cache has this block
- **Owner:**
 - The block is in some other cache
 - If we make modifications, it's our responsibility to tell all the other caches in shared state about these changes
 - Allows for writing while other threads read the same data.

MOESI State Transitions

Cache Coherency (MOESI protocol)




Coherence Misses

- **Coherence misses**: misses caused by blocks invalidated due to other cores
- Since only one thread is allowed to have a dirty block at a time, two threads writing to the same block cause a lot of coherence misses
 - This is why the interleaved pragma omp was much slower than the blocked pragma omp
 - Another form of **thrashing**
- Therefore, when writing multithreaded code, we want to have each thread work on a separate block as much as possible
 - But also close enough so that the shared L3 cache gets hits.

False Sharing

- Coherence misses due to interleaving are an example of **false sharing**
- The entire block is invalidated and must be reloaded, even though technically no data is shared.



<https://www.codeproject.com/Articles/85356/Avoiding-and-Identifying-False-Sharing-Among-Threa>

In conclusion...

- We've discussed memory caching in detail. Caching in general shows up over and over in computer systems
 - Filesystem cache, Web page cache, Game databases / tablebases, Software memoization, Others?
- Big idea: if something is expensive but we want to do it repeatedly, do it once and cache the result.
- Cache design choices:
 - Block size
 - Size of cache: speed v. capacity
 - Associativity (direct-mapped v. set v. fully associative)
 - Replacement policy (LRU, MRU, FIFO, ...)
 - Write Policy (write through v. write back)
 - Extra:
 - Multilevel cache
 - Cache coherence protocol
- Understanding caches allows us to make correct hardware and software design choices, depending on programs, technology, budget, etc.
- Next week: upending everything you (think you) know about memory :)